

Lightning Fast and Space Efficient Inequality Joins

Zuhair Khayyat, William Lucia, Megna Singh, Mourad Ouzzani, Paolo Papotti, Jorge Quiané, Nan Tang, Panos Kalnis

Motivation

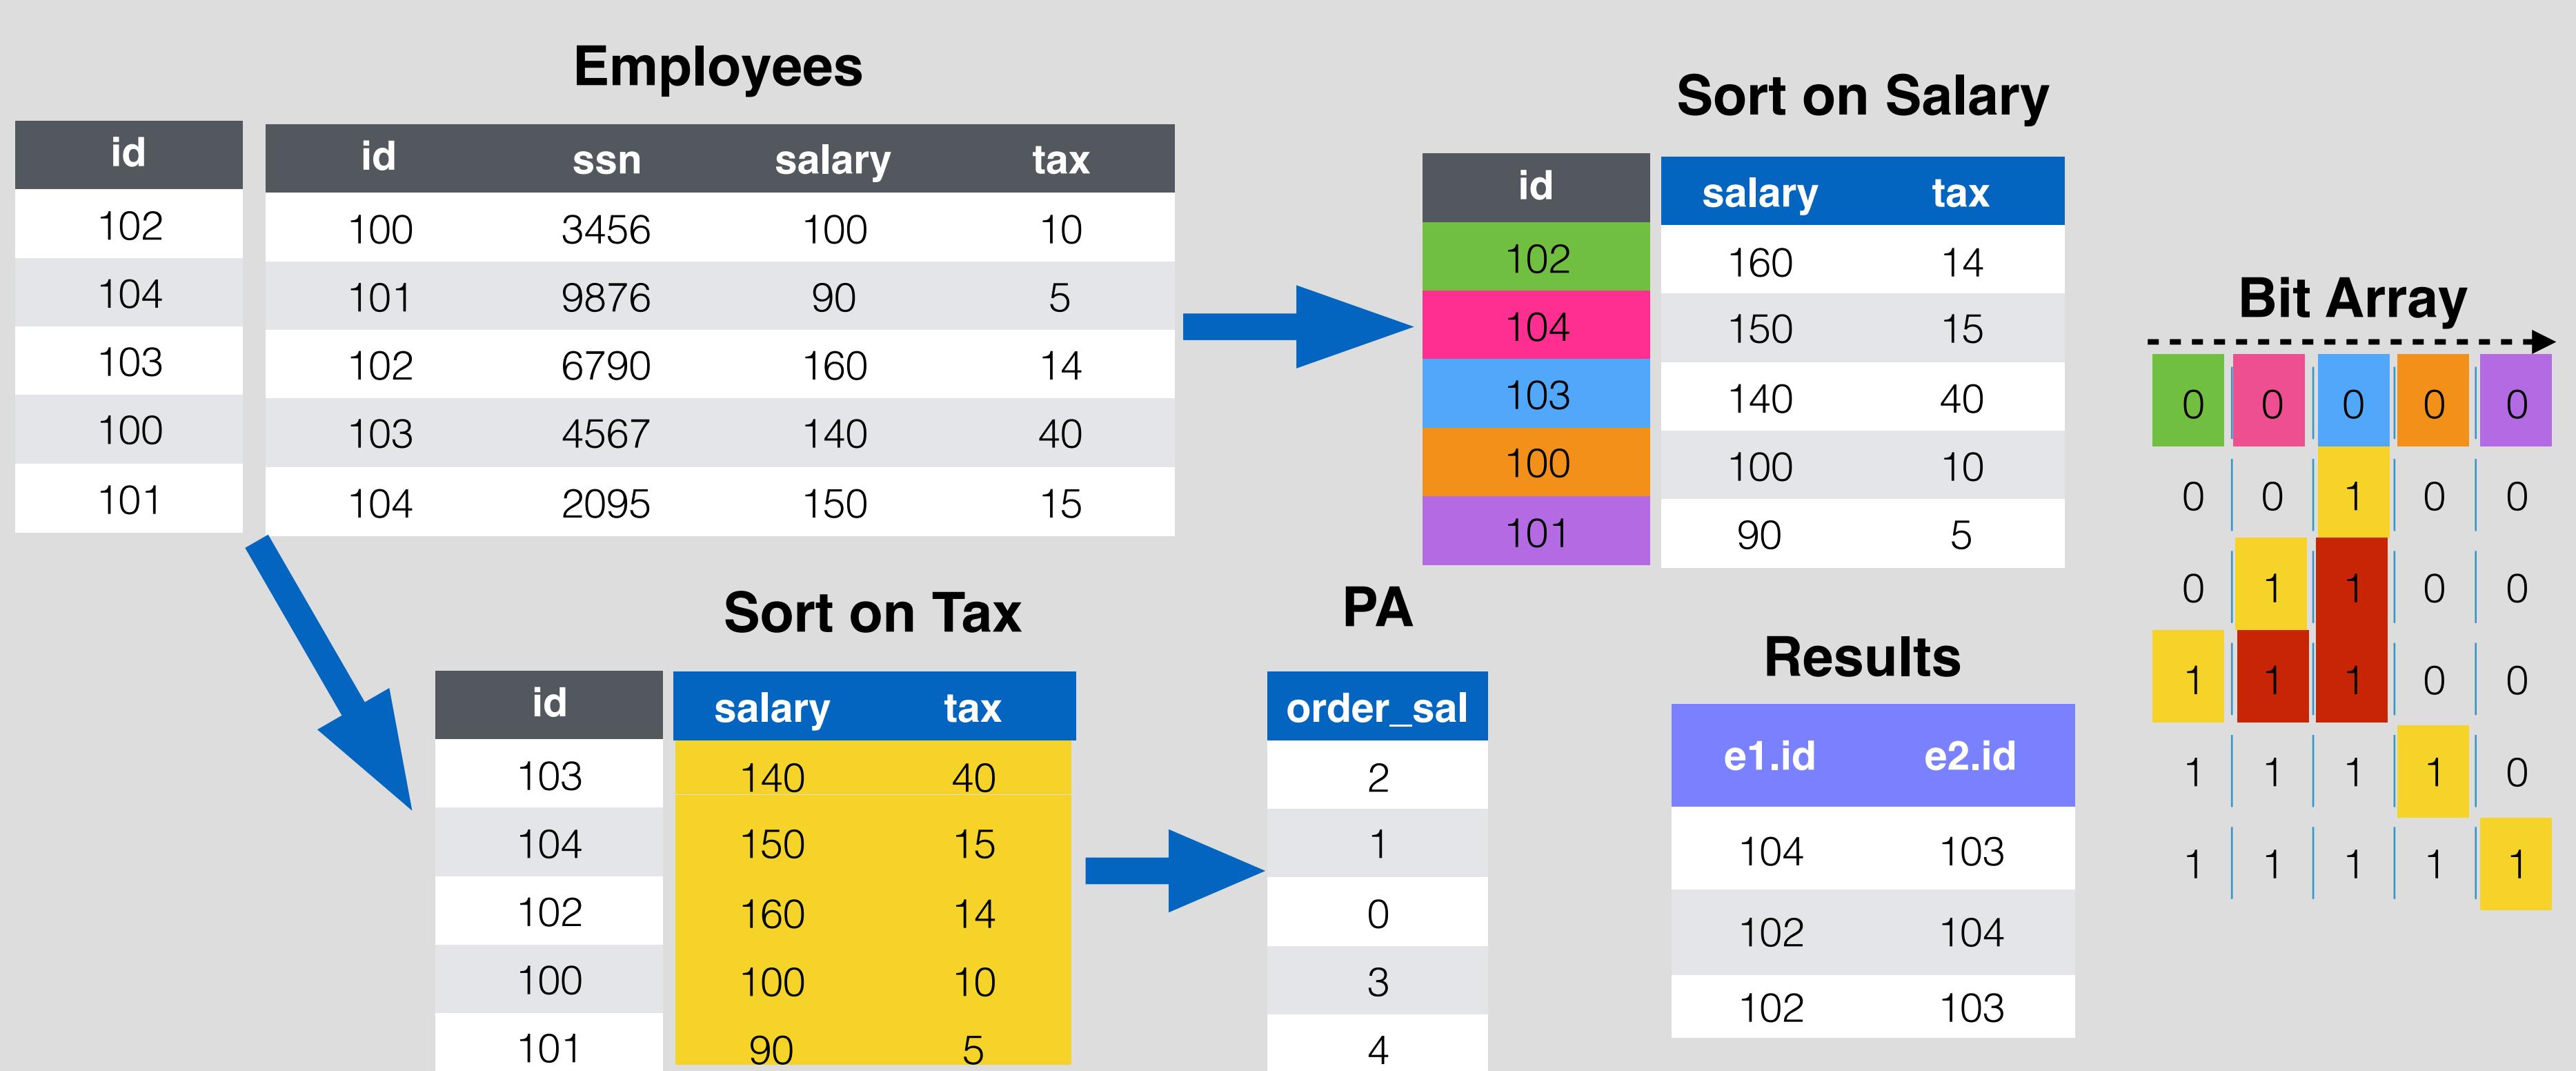
```
SELECT e1.id, e2.id
FROM Employees e1, e2
WHERE e1.salary > e2.salary
AND e1.tax < e2.tax;
```

- Important on many applications: temporal and spatial databases, data cleaning, social networks, sensor networks,..
- Inefficient solutions: cross product + selection or expensive indexes

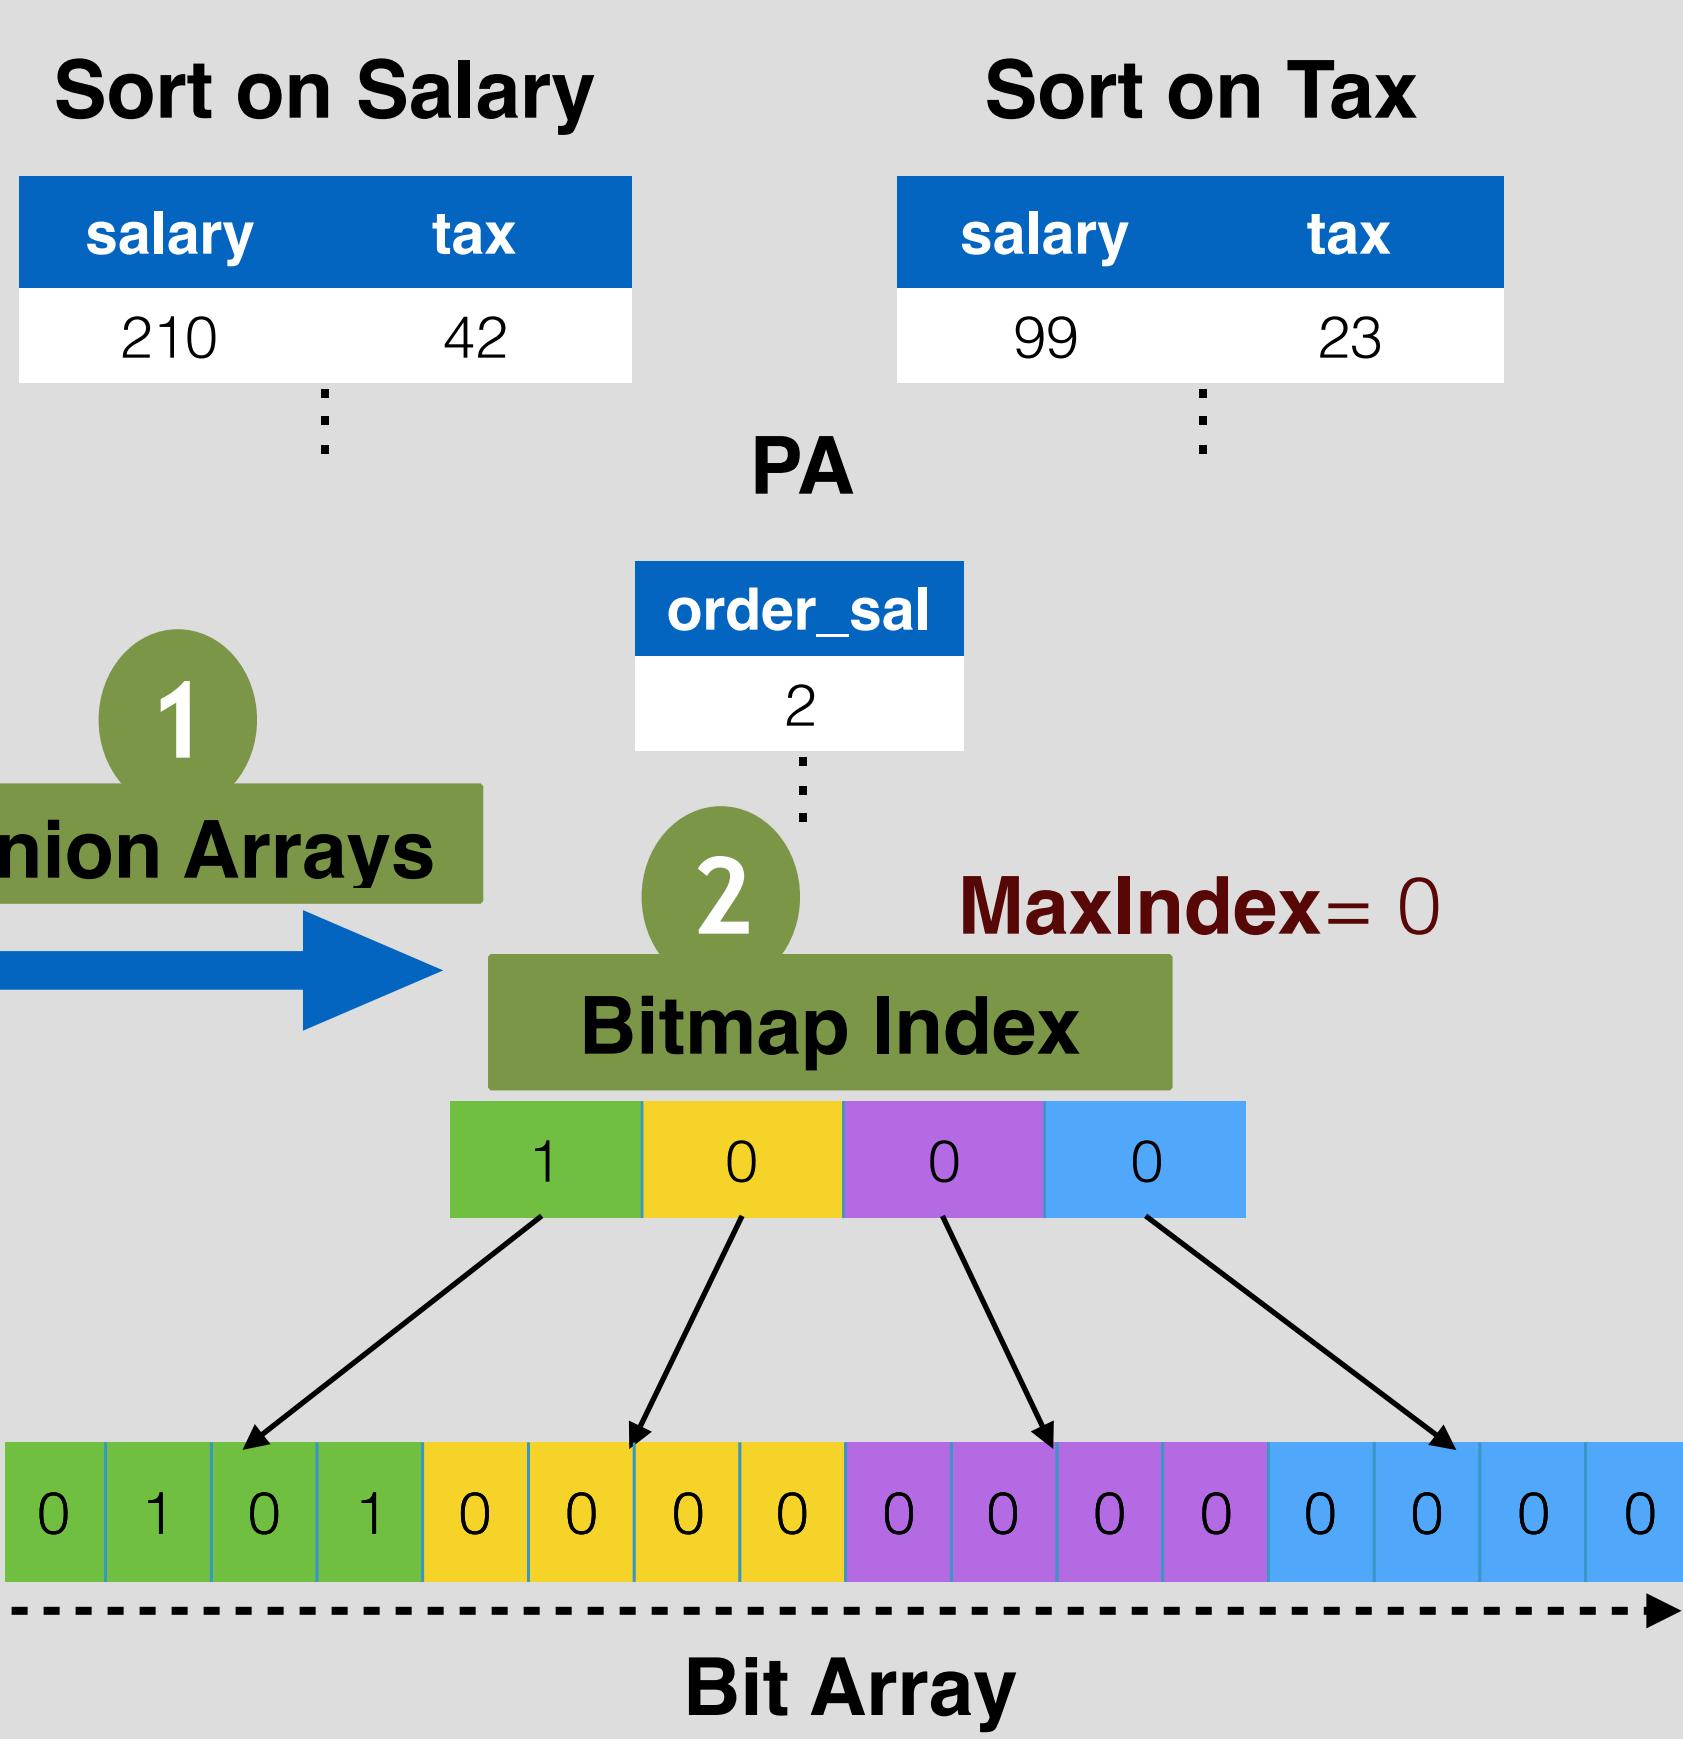
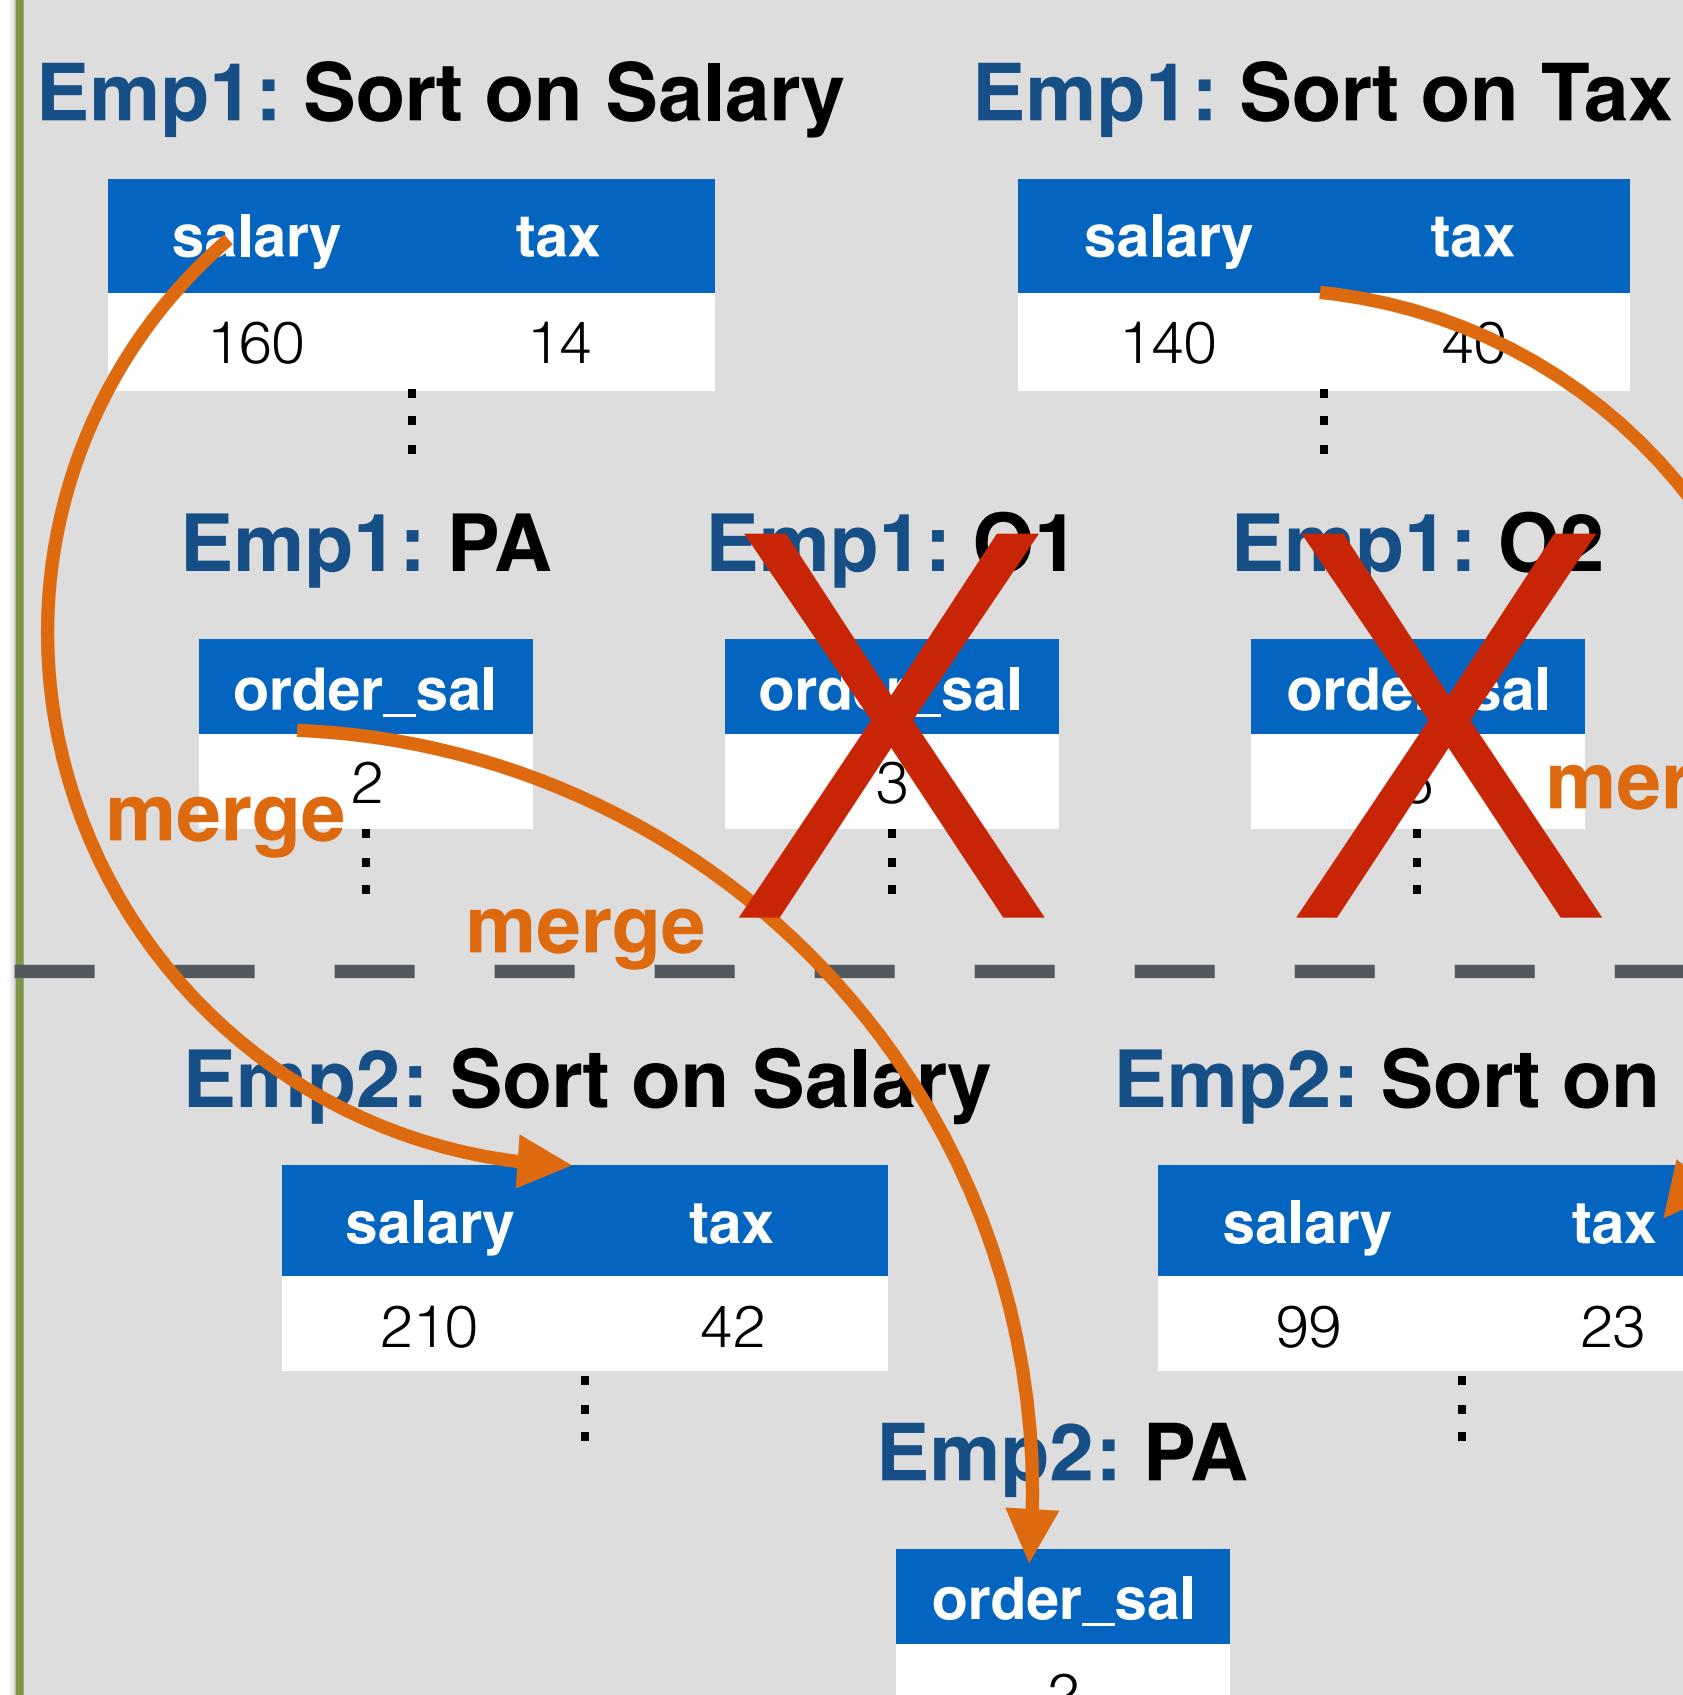
IEJoin

Algorithm 2: IESELFJOIN

```
input : query Q with 2 join predicates  $t_1.X \text{ op}_1 t_2.X$  and  $t_1.Y \text{ op}_2 t_2.Y$ , table T of size n
output: a list of tuple pairs  $(t_i, t_j)$ 
1 let  $L_1$  (resp.  $L_2$ ) be the array of column X (resp. Y)
2 if  $(\text{op}_1 \in \{\leq, \geq\})$  sort  $L_1$  in ascending order
3 else if  $(\text{op}_1 \in \{<, \leq\})$  sort  $L_2$  in descending order
4 if  $(\text{op}_2 \in \{\geq, \leq\})$  sort  $L_2$  in ascending order
5 else if  $(\text{op}_2 \in \{<, \leq\})$  sort  $L_2$  in descending order
6 compute the permutation array P of  $L_2$  w.r.t.  $L_1$ 
7 initialize bit-array B ( $|B| = n$ ), and set all bits to 0
8 initialize join_result as an empty list for tuple pairs
9 if  $(\text{op}_1 \in \{\leq, \geq\})$  and  $\text{op}_2 \in \{\leq, \geq\}$ ) eqOff = 0
10 else eqOff = 1
11 for  $(i \leftarrow 1$  to  $n$ ) do
12     pos  $\leftarrow P[i]$ 
13     B[pos]  $\leftarrow 1$ 
14     for  $(j \leftarrow \text{pos} + \text{eqOff}$  to  $n$ ) do
15         if  $B[j] = 1$  then
16             add tuples w.r.t.  $(L_1[j], L_1[P[i]])$  to join_result
17 return join_result
```



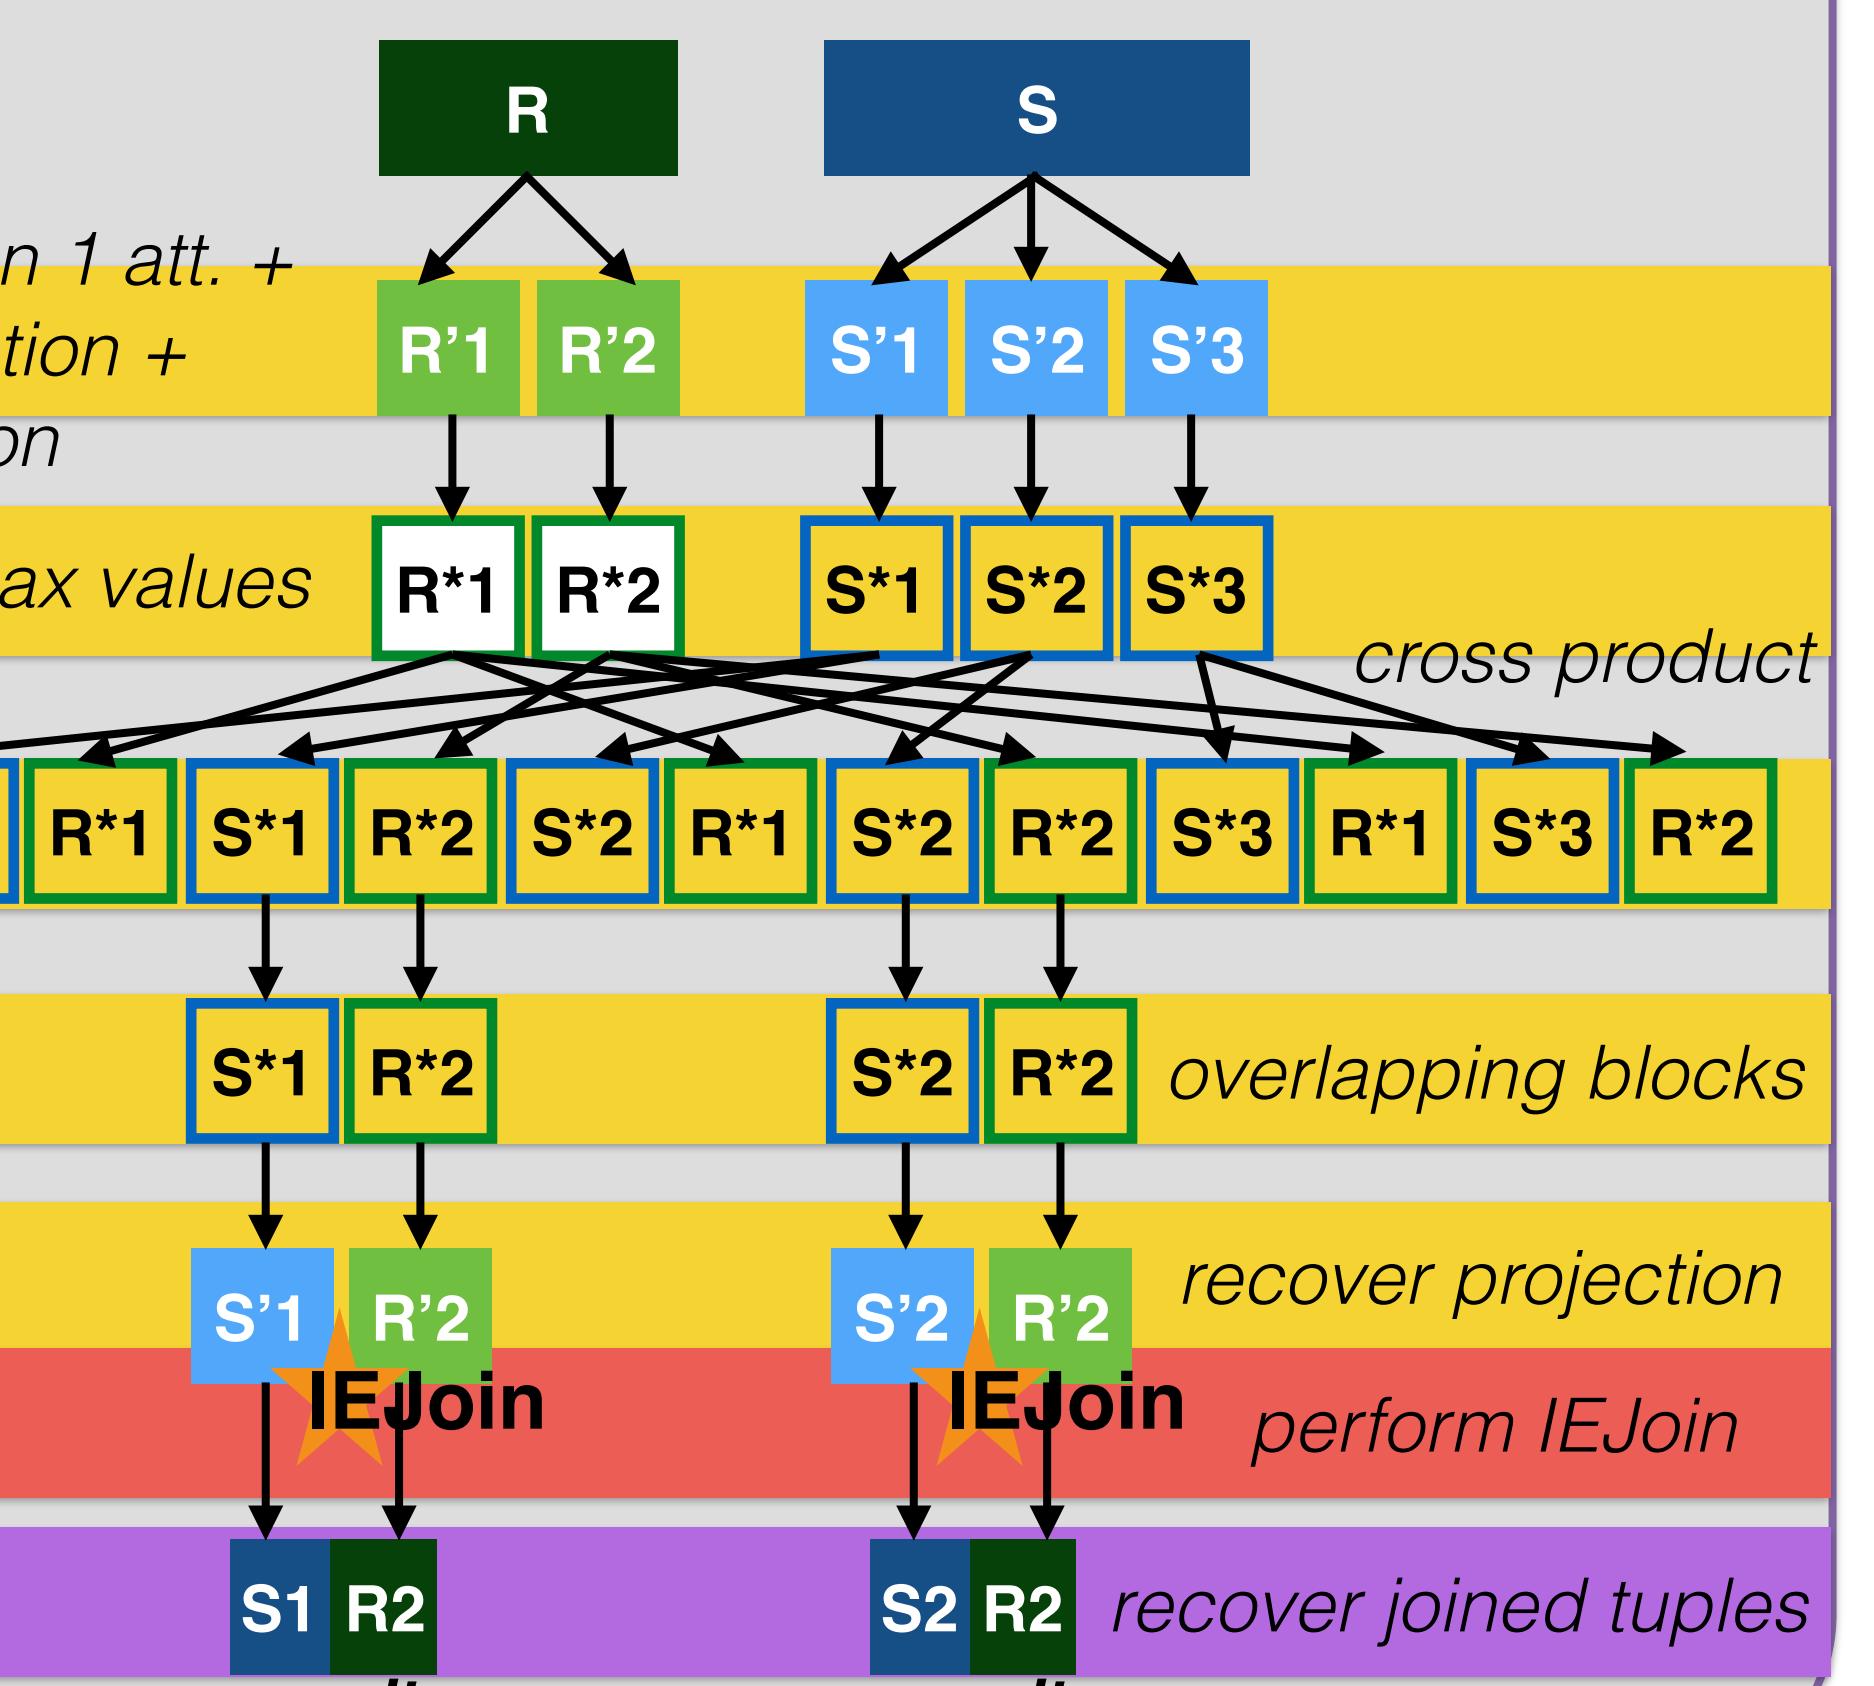
Optimizations



Scalable IEJoin

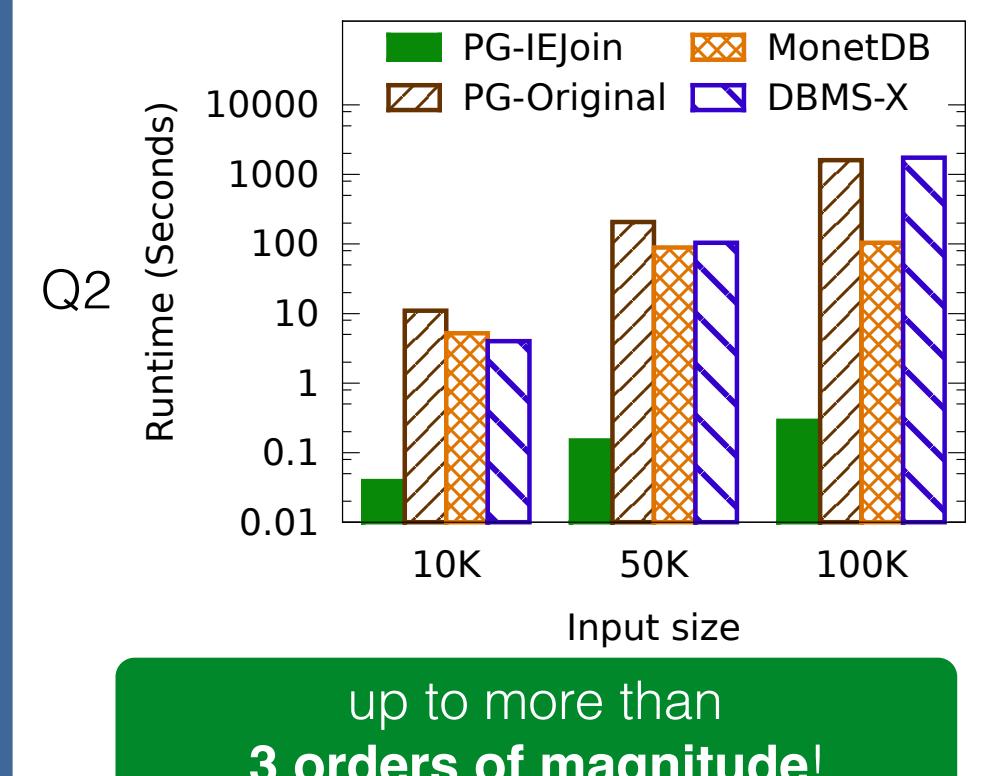
Phase 1: Pre-processing Phase 2: IEJoin

Phase 3: Post-processing

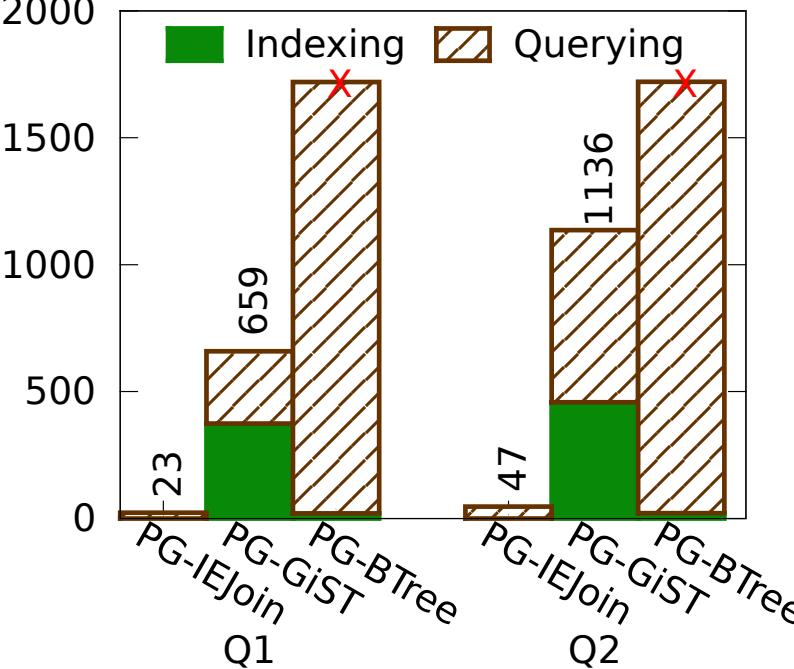


Results

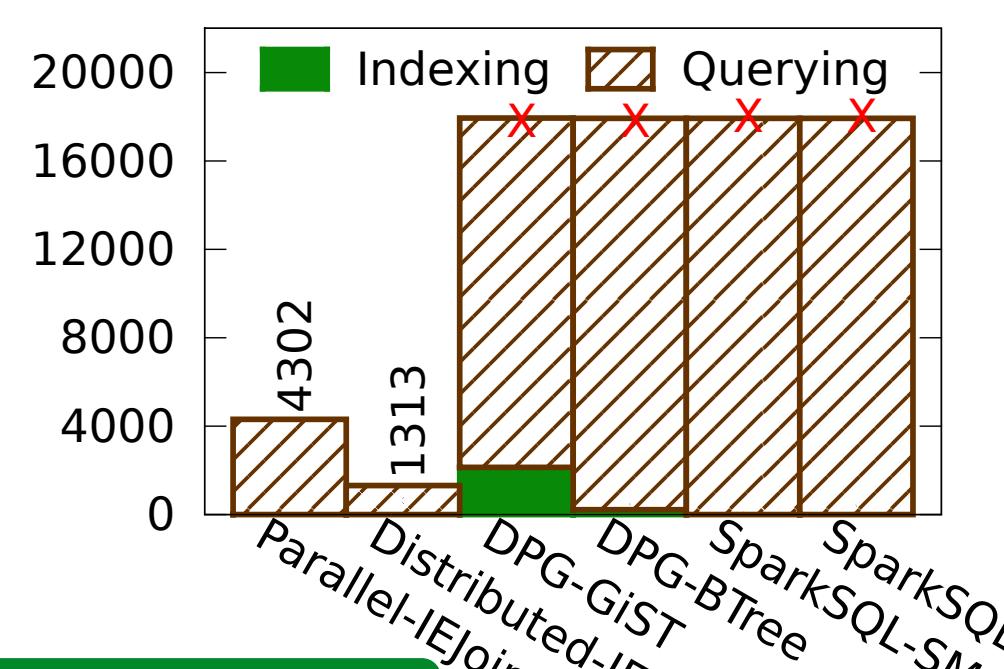
Single-Node Performance



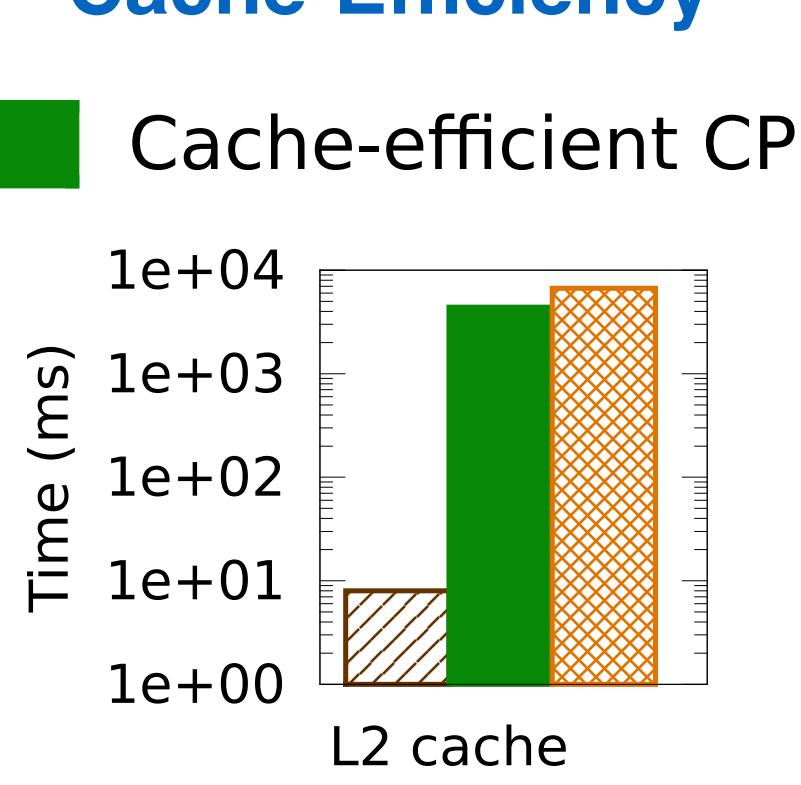
IEJoin vs Indexing



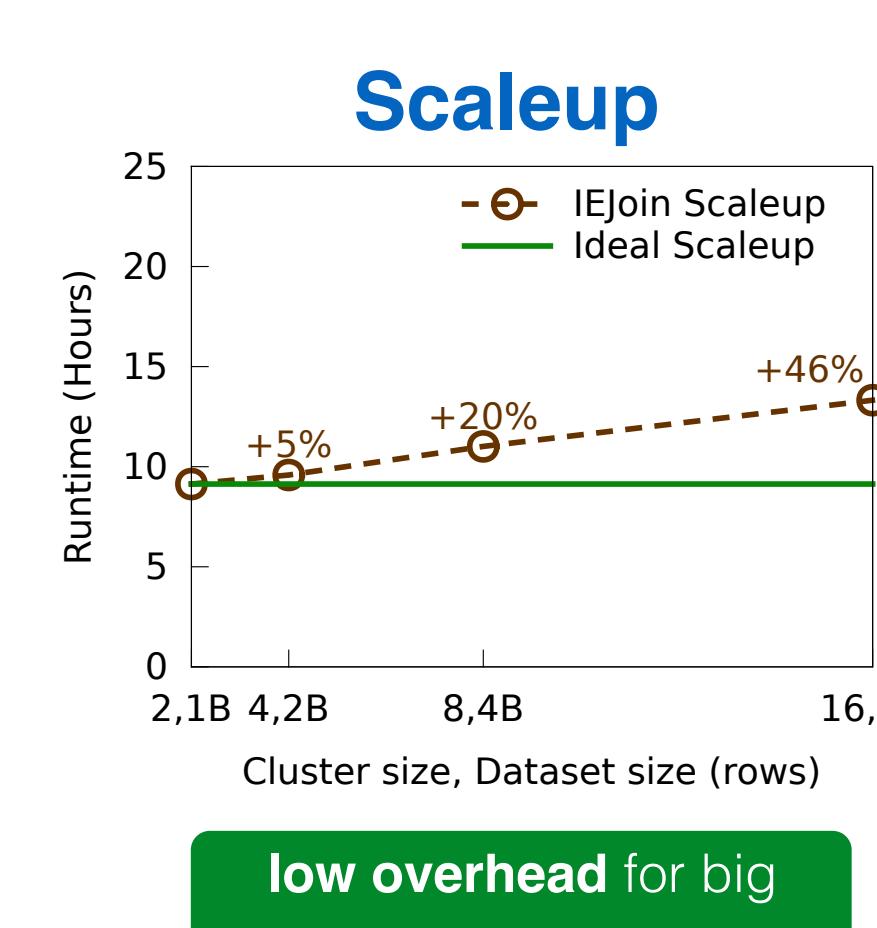
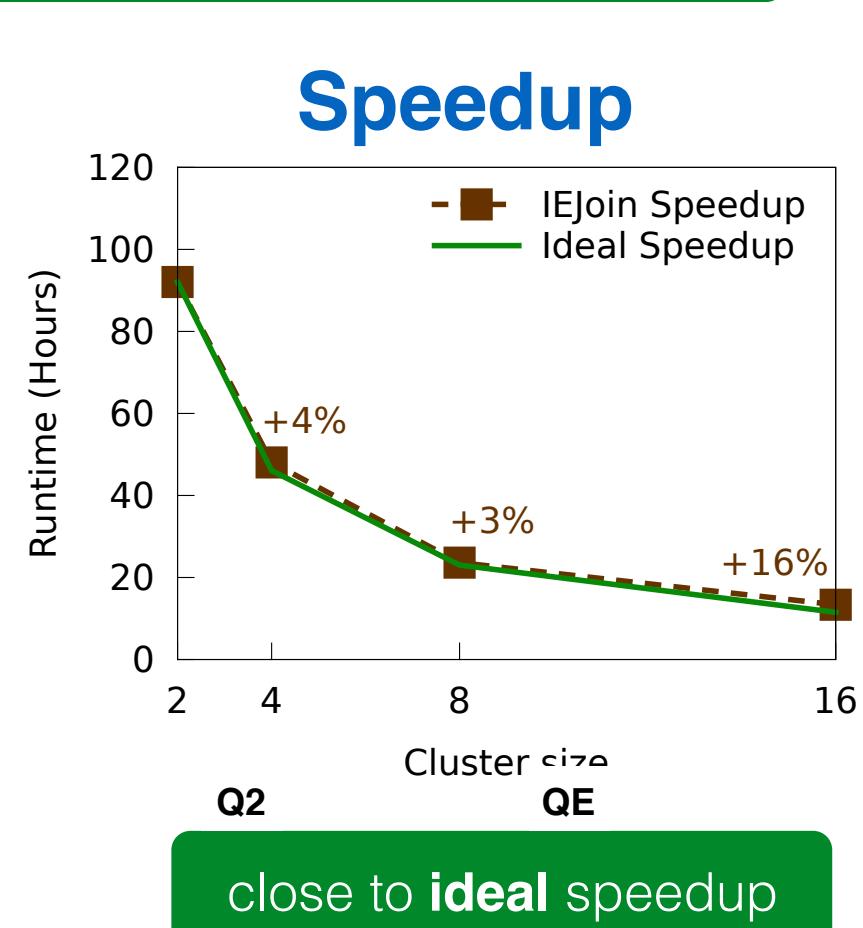
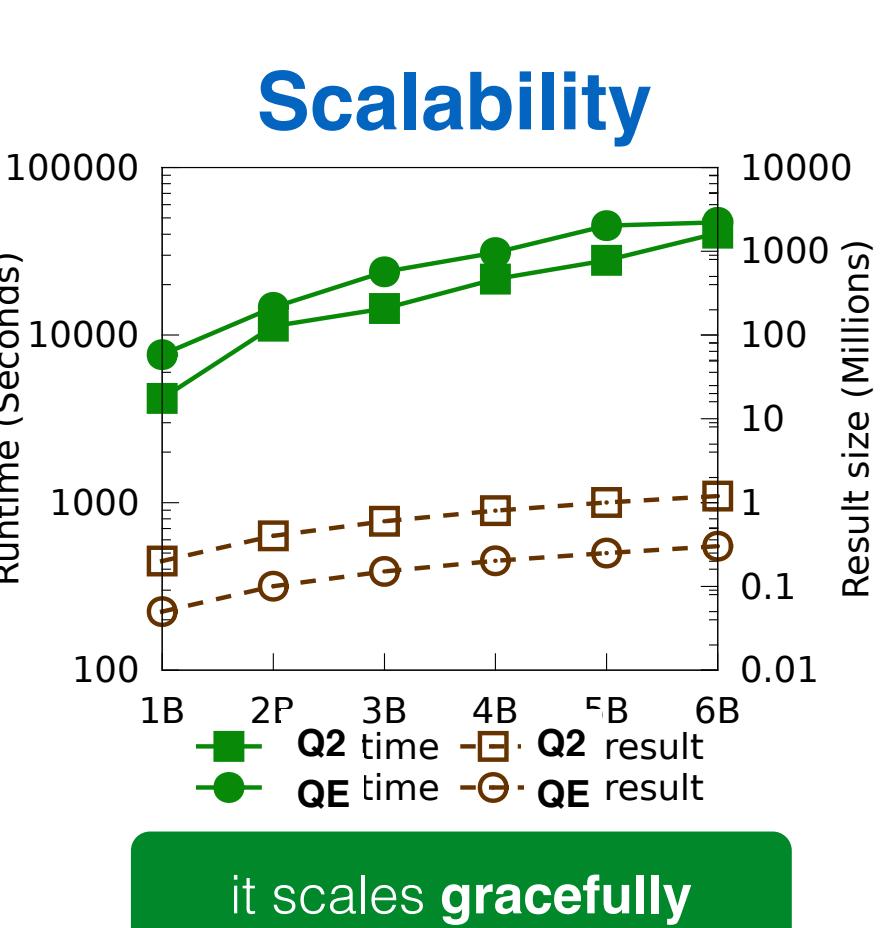
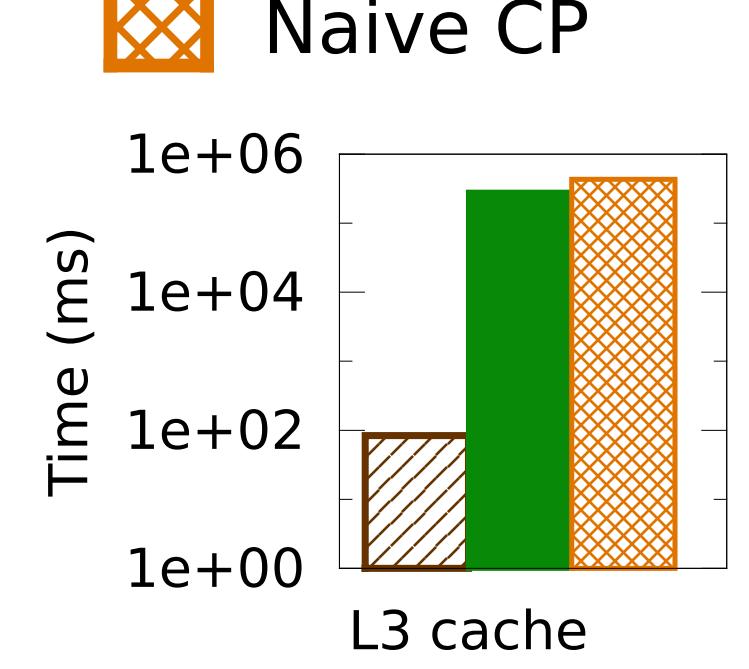
Multi-Node Performance



Cache-Efficiency



Naive CP



| Query | Input | Output | Time(secs) | Mem(GB) |
|-------|-------|--------|------------|---------|
| Q1 | 100K | 9K | 0.30 | 0.1 |
| Q1 | 200K | 1.1K | 0.5 | 0.2 |
| Q1 | 1M | 29K | 2.79 | 0.9 |
| Q1 | 10M | 3M | 27.64 | 8.8 |
| Q2 | 100K | 0.2K | 0.34 | 0.1 |
| Q2 | 200K | 0.8K | 0.65 | 0.2 |
| Q2 | 1M | 20K | 3.38 | 0.9 |
| Q2 | 10M | 2M | 59.6 | 9.7 |
| Q4 | 100K | 6M | 2.8 | 0.1 |
| Q4 | 200K | 25M | 10.6 | 0.2 |
| Q4 | 1M | 0.4M | 186 | 0.9 |
| Q4 | 10M | 50.5M | 28,928 | 8.2 |

low overhead for big

low memory footprint

Highly cache efficient

Code available at: <https://github.com/daqcri/rheem>

RHEEM