

Lightning Fast & Space Efficient Inequality Joins

Data Analytics Team
@QCRI



Once upon a time

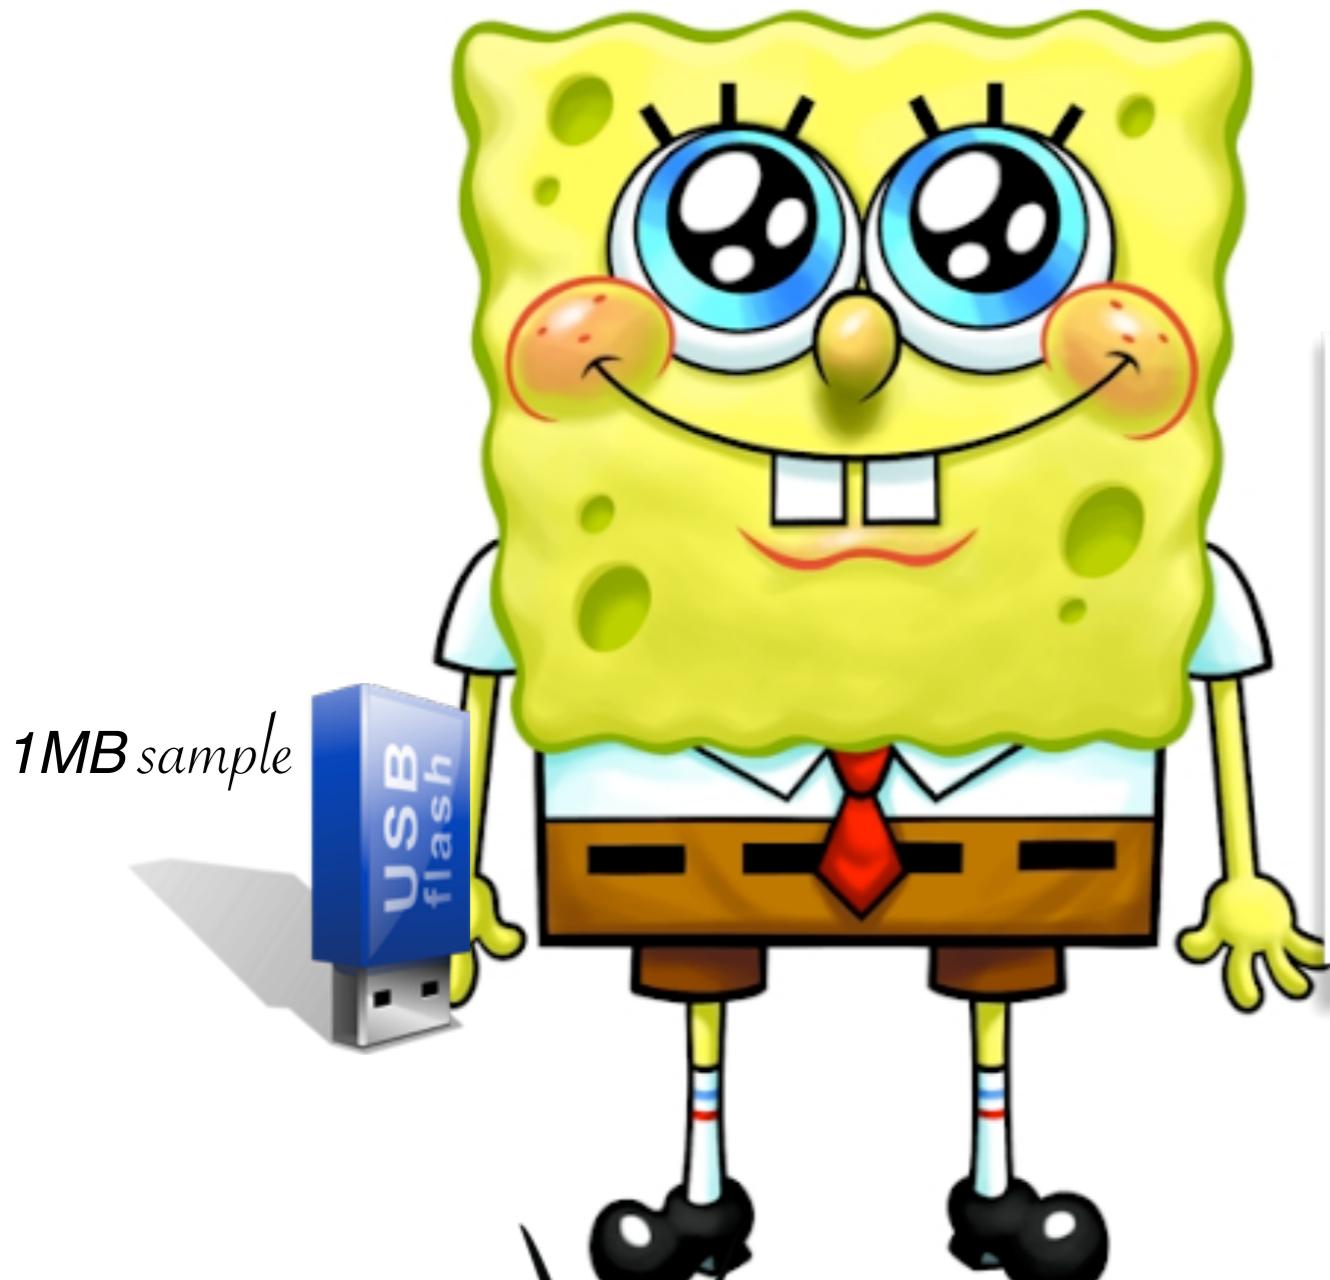
at QCRJ

Bob@QCRI



Nadeef
RMEEM

Bob@QCRI



Nadeef
RMEEM

Data quality rule

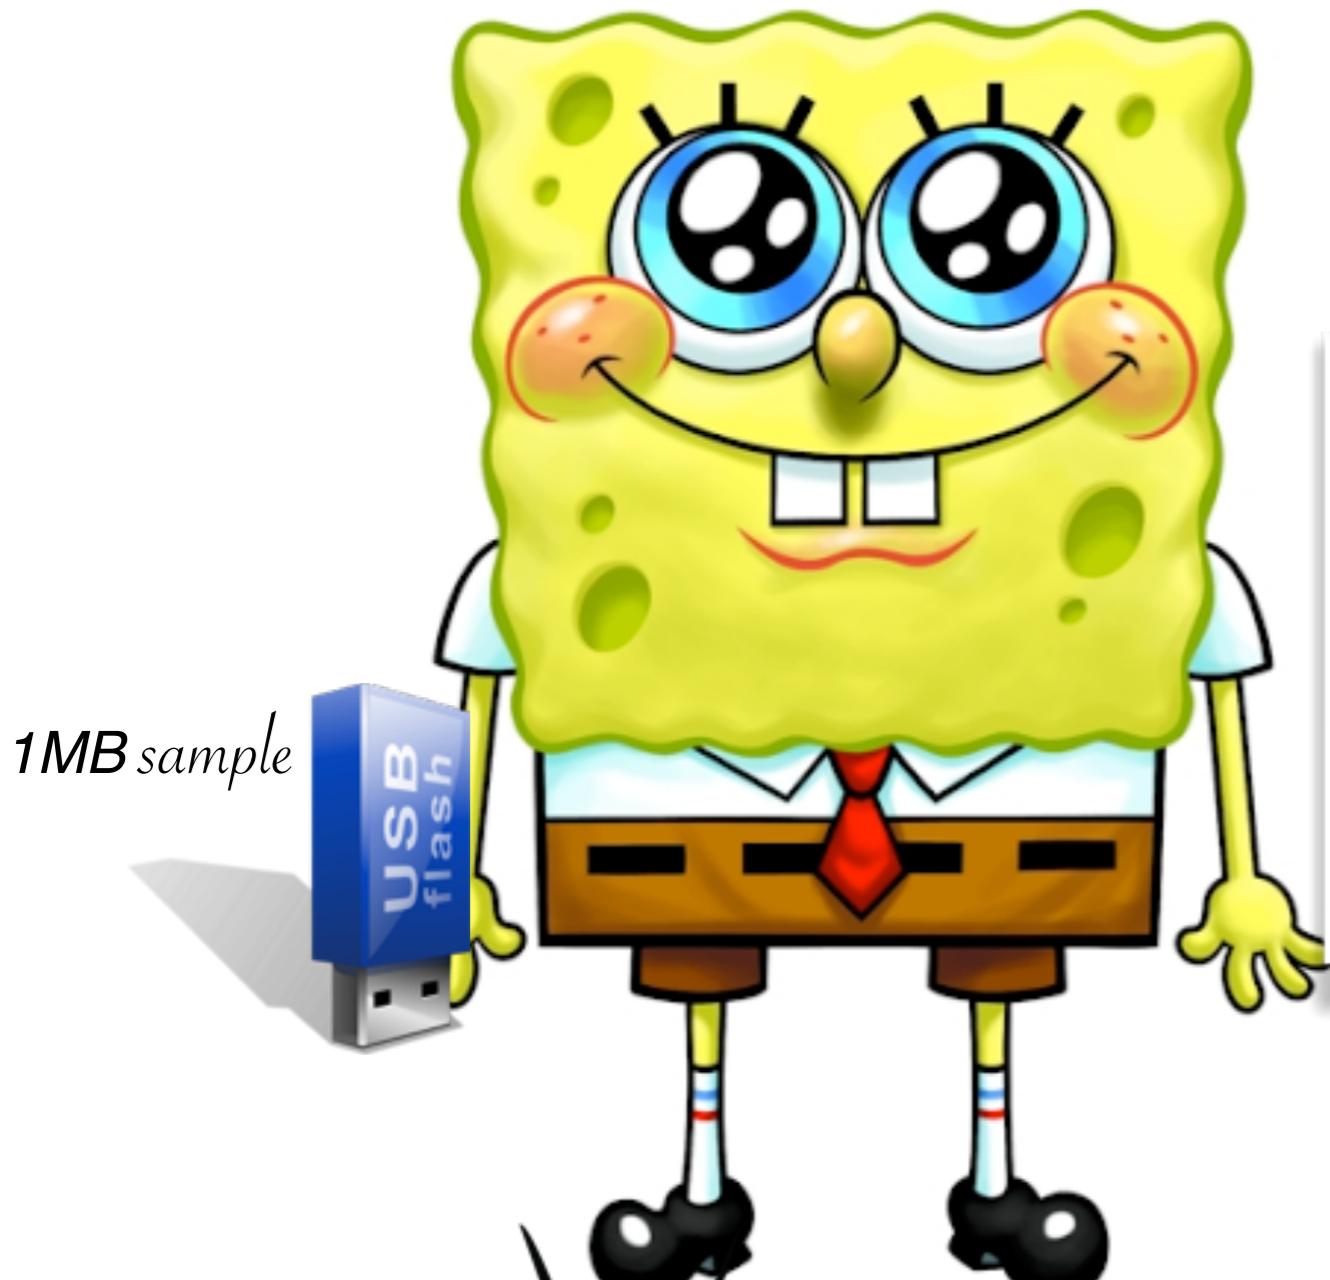
Q1:

SELECT e.id

FROM Employees e, Managers m

WHERE e.ssn = m.ssn

Bob@QCRI



Nadeef
RMEEM

Data quality rule

Q1:

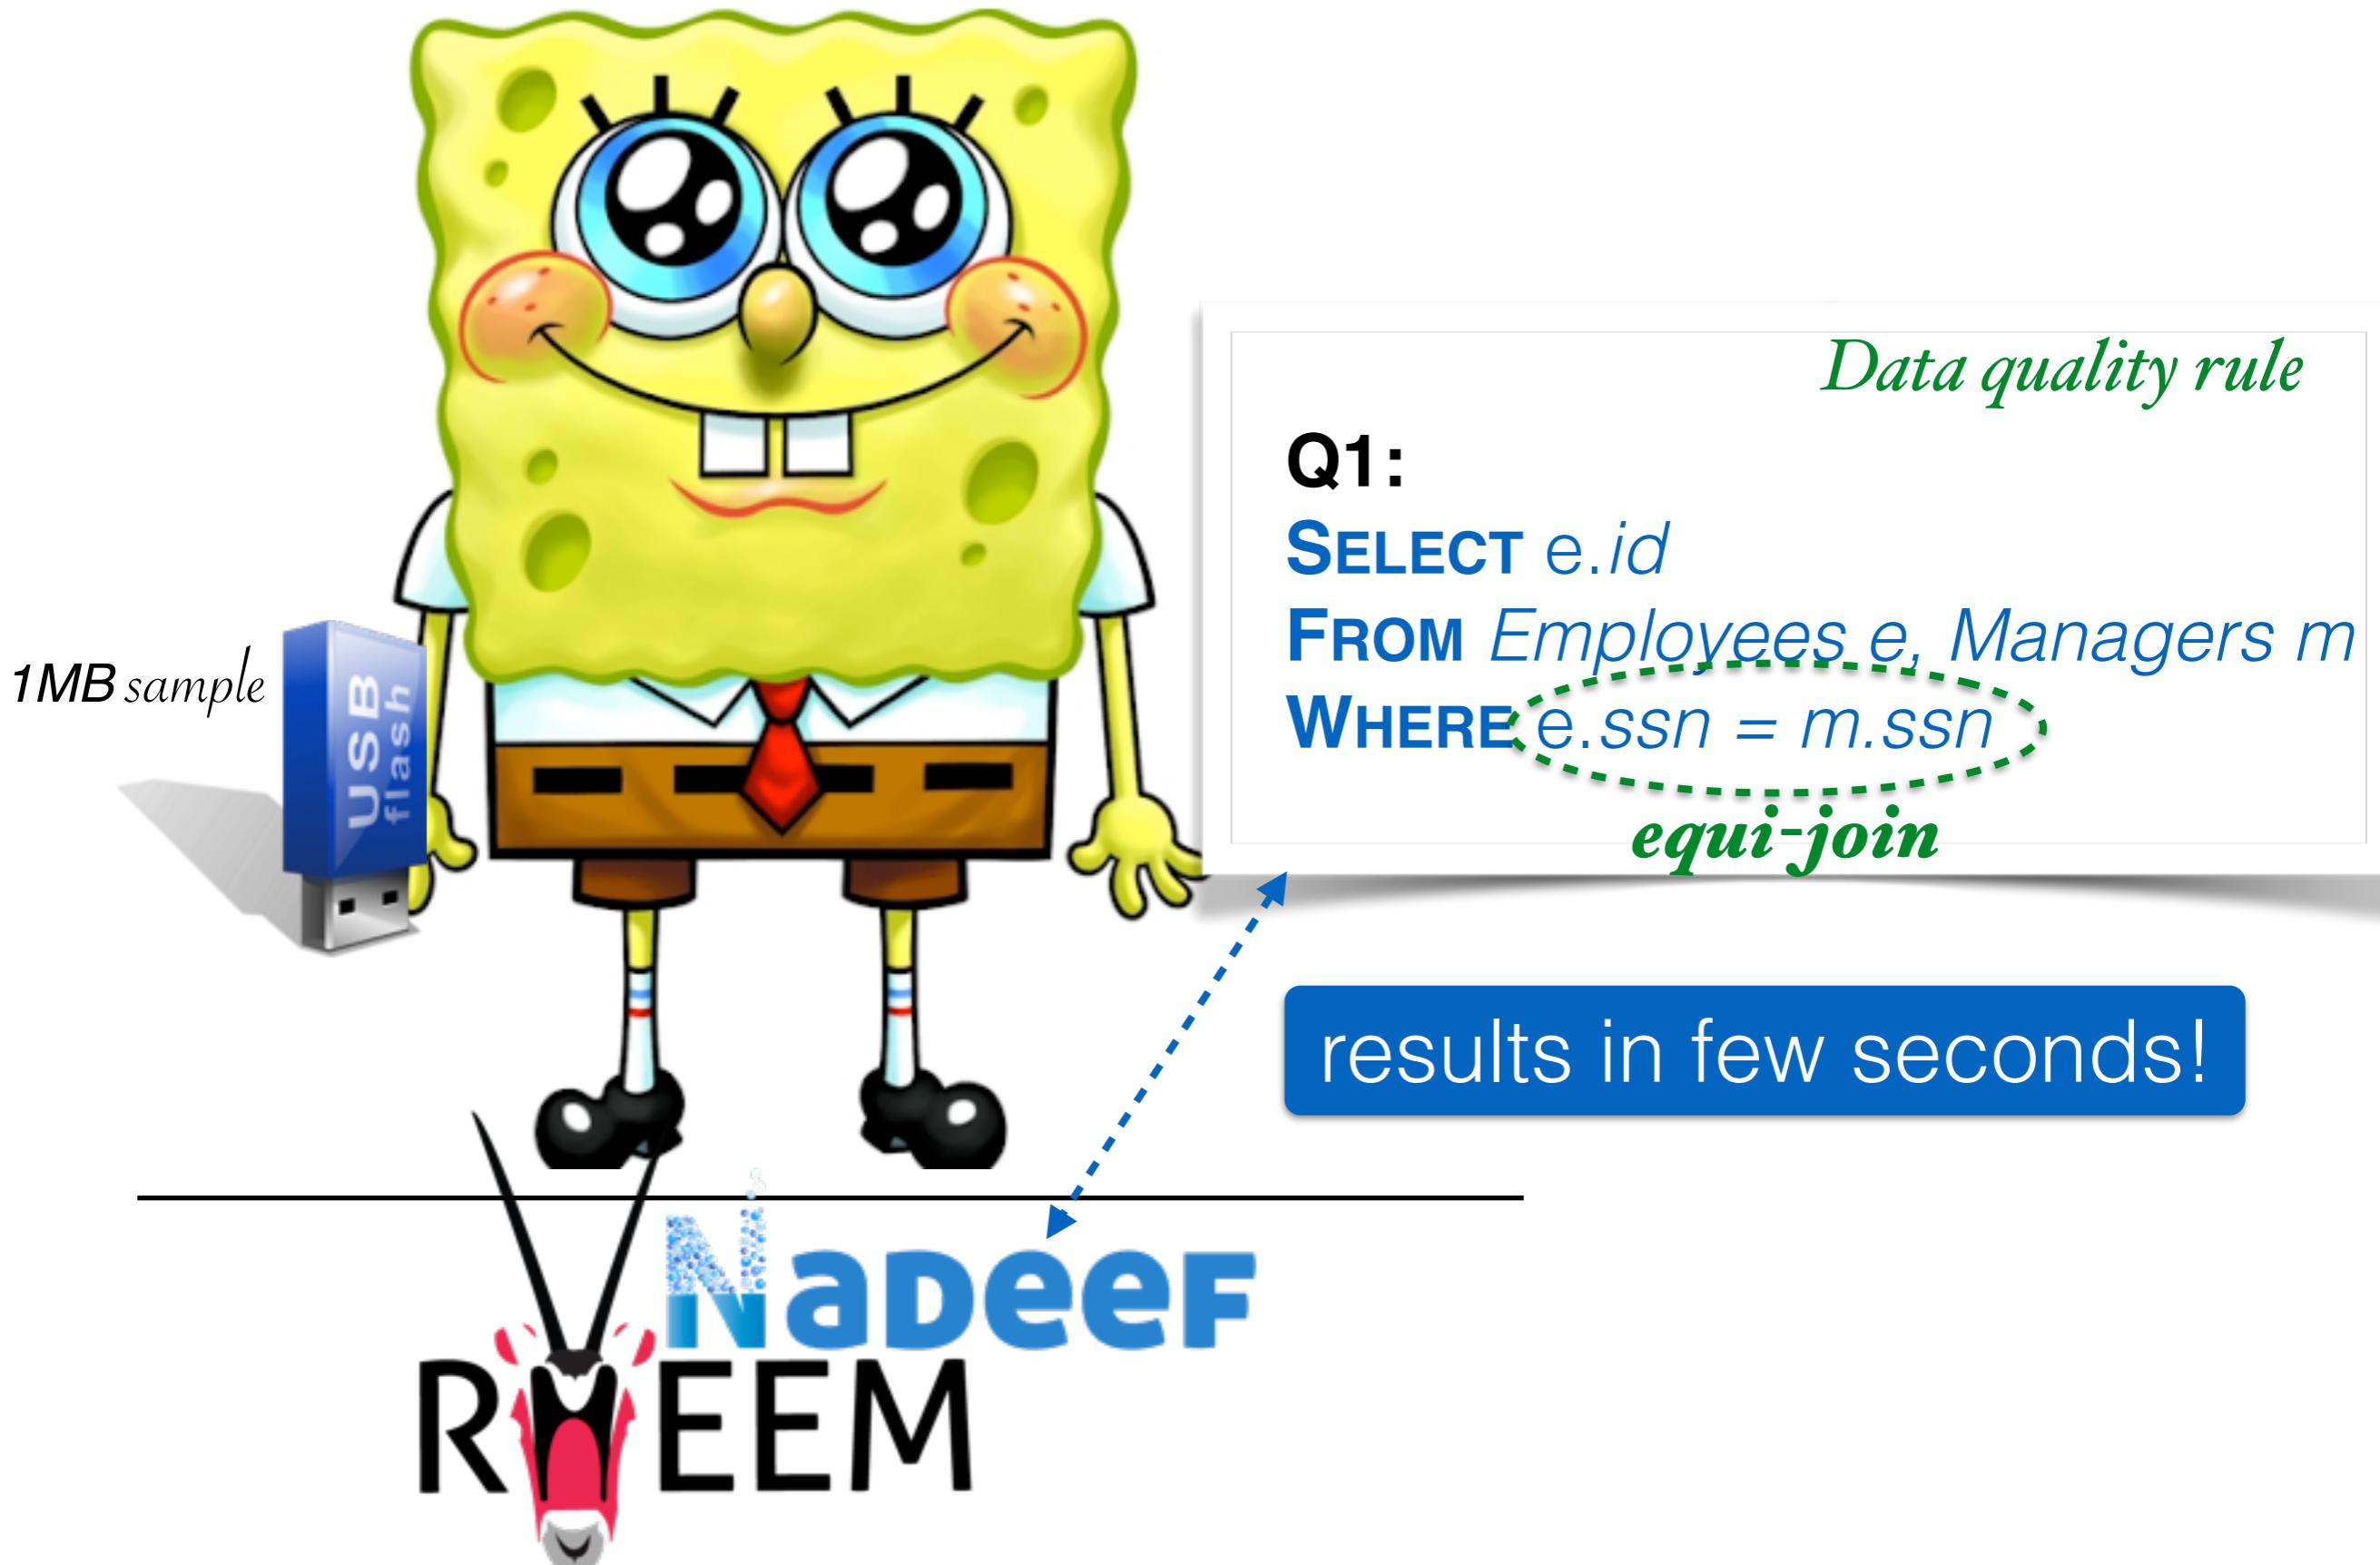
SELECT e.id

FROM Employees e, Managers m

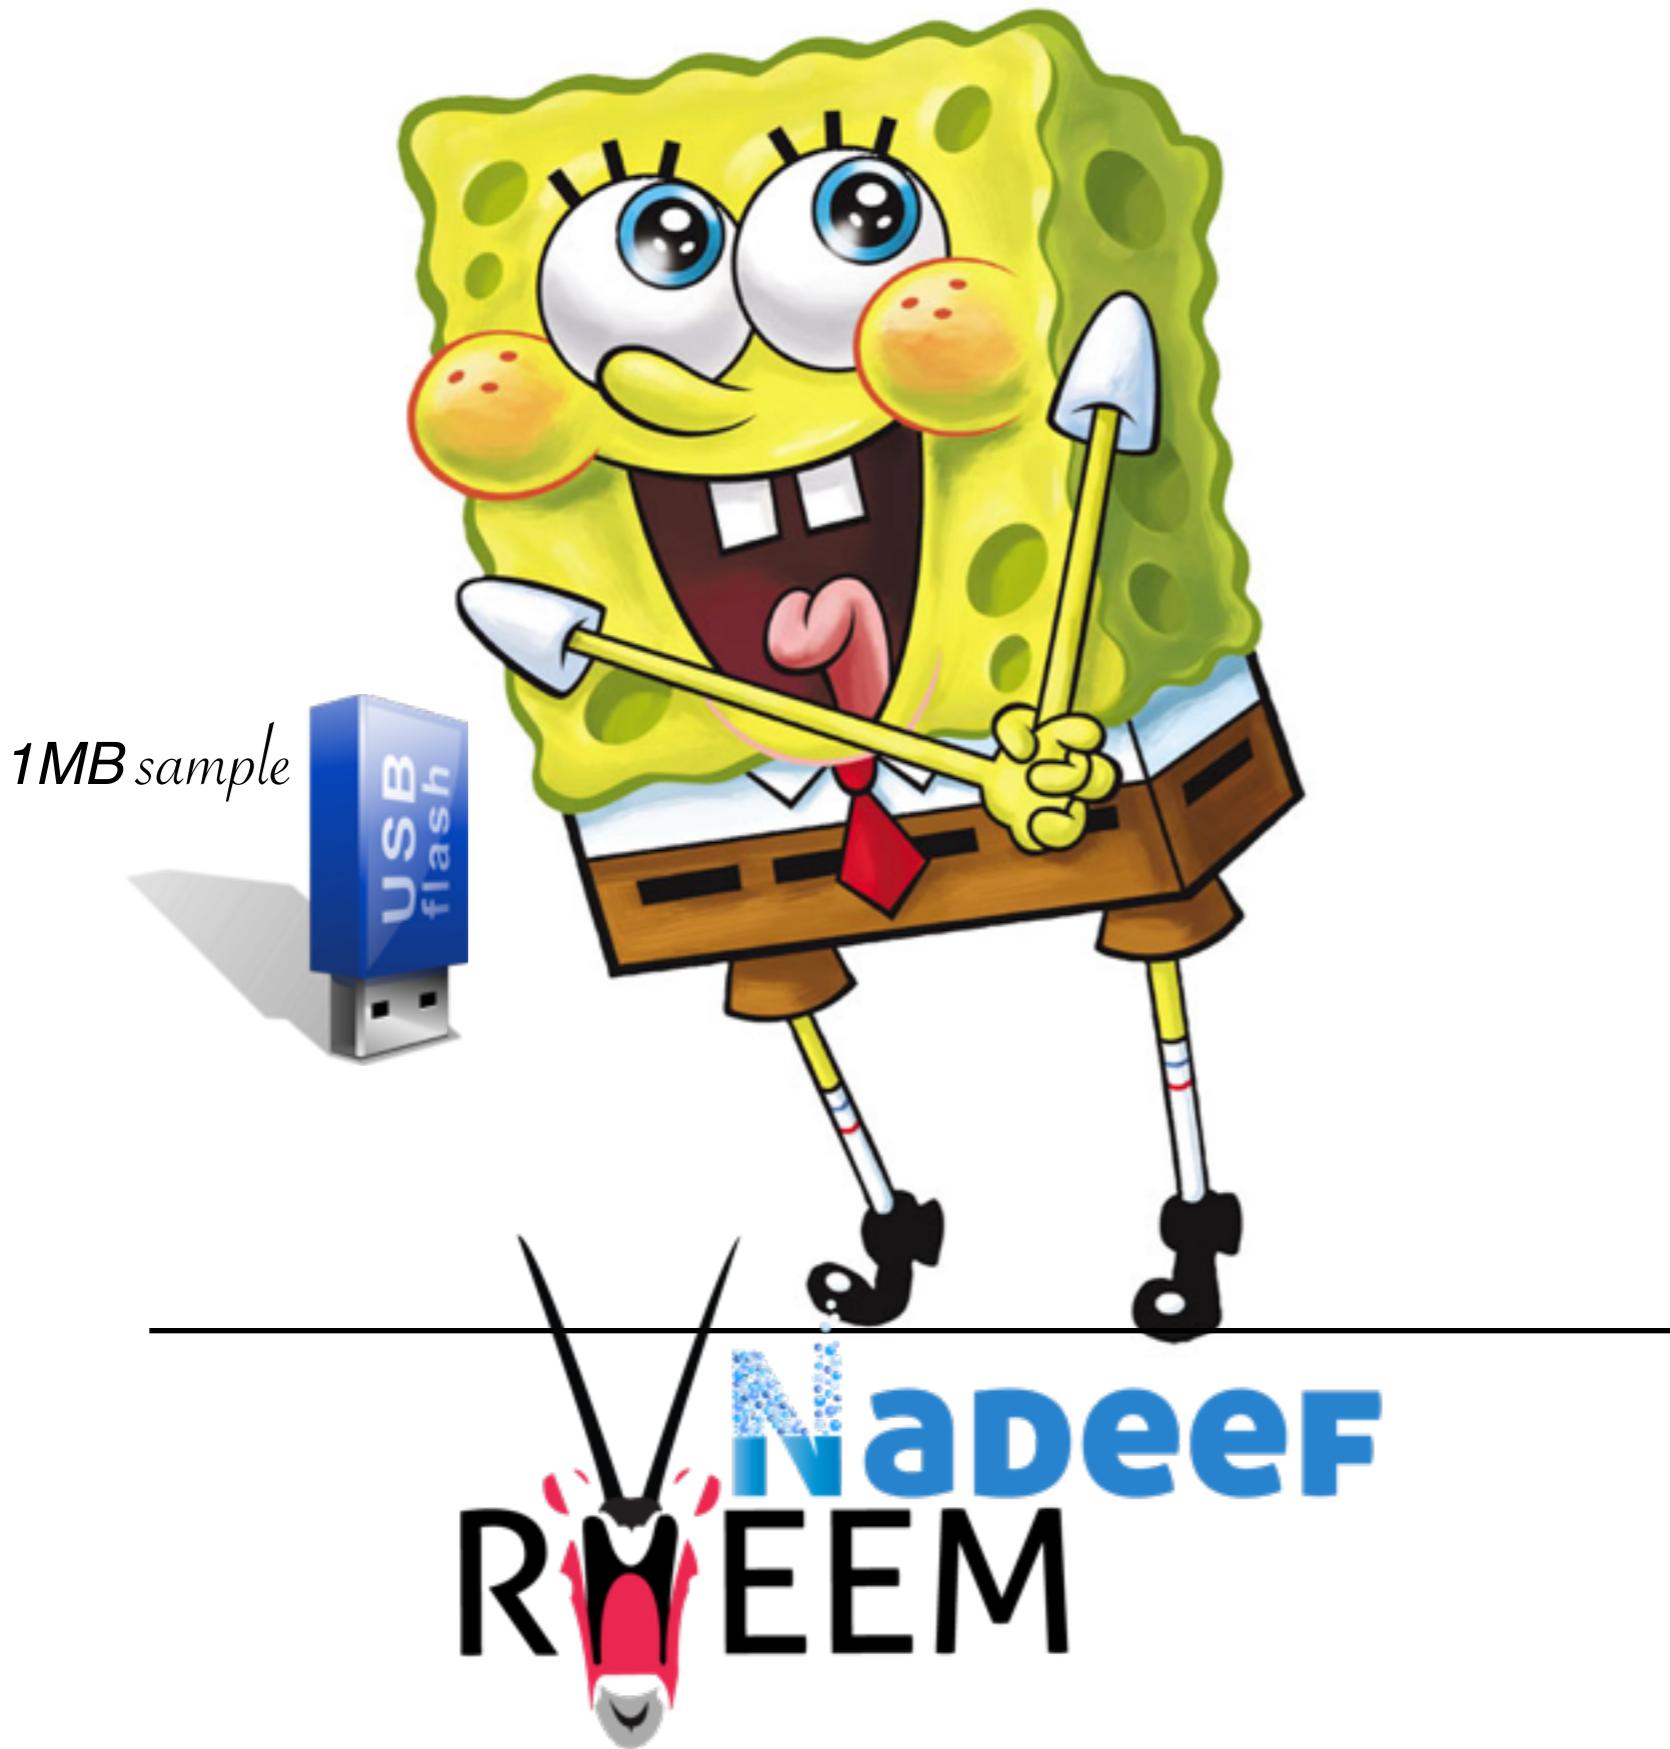
WHERE e.ssn = m.ssn

equi-join

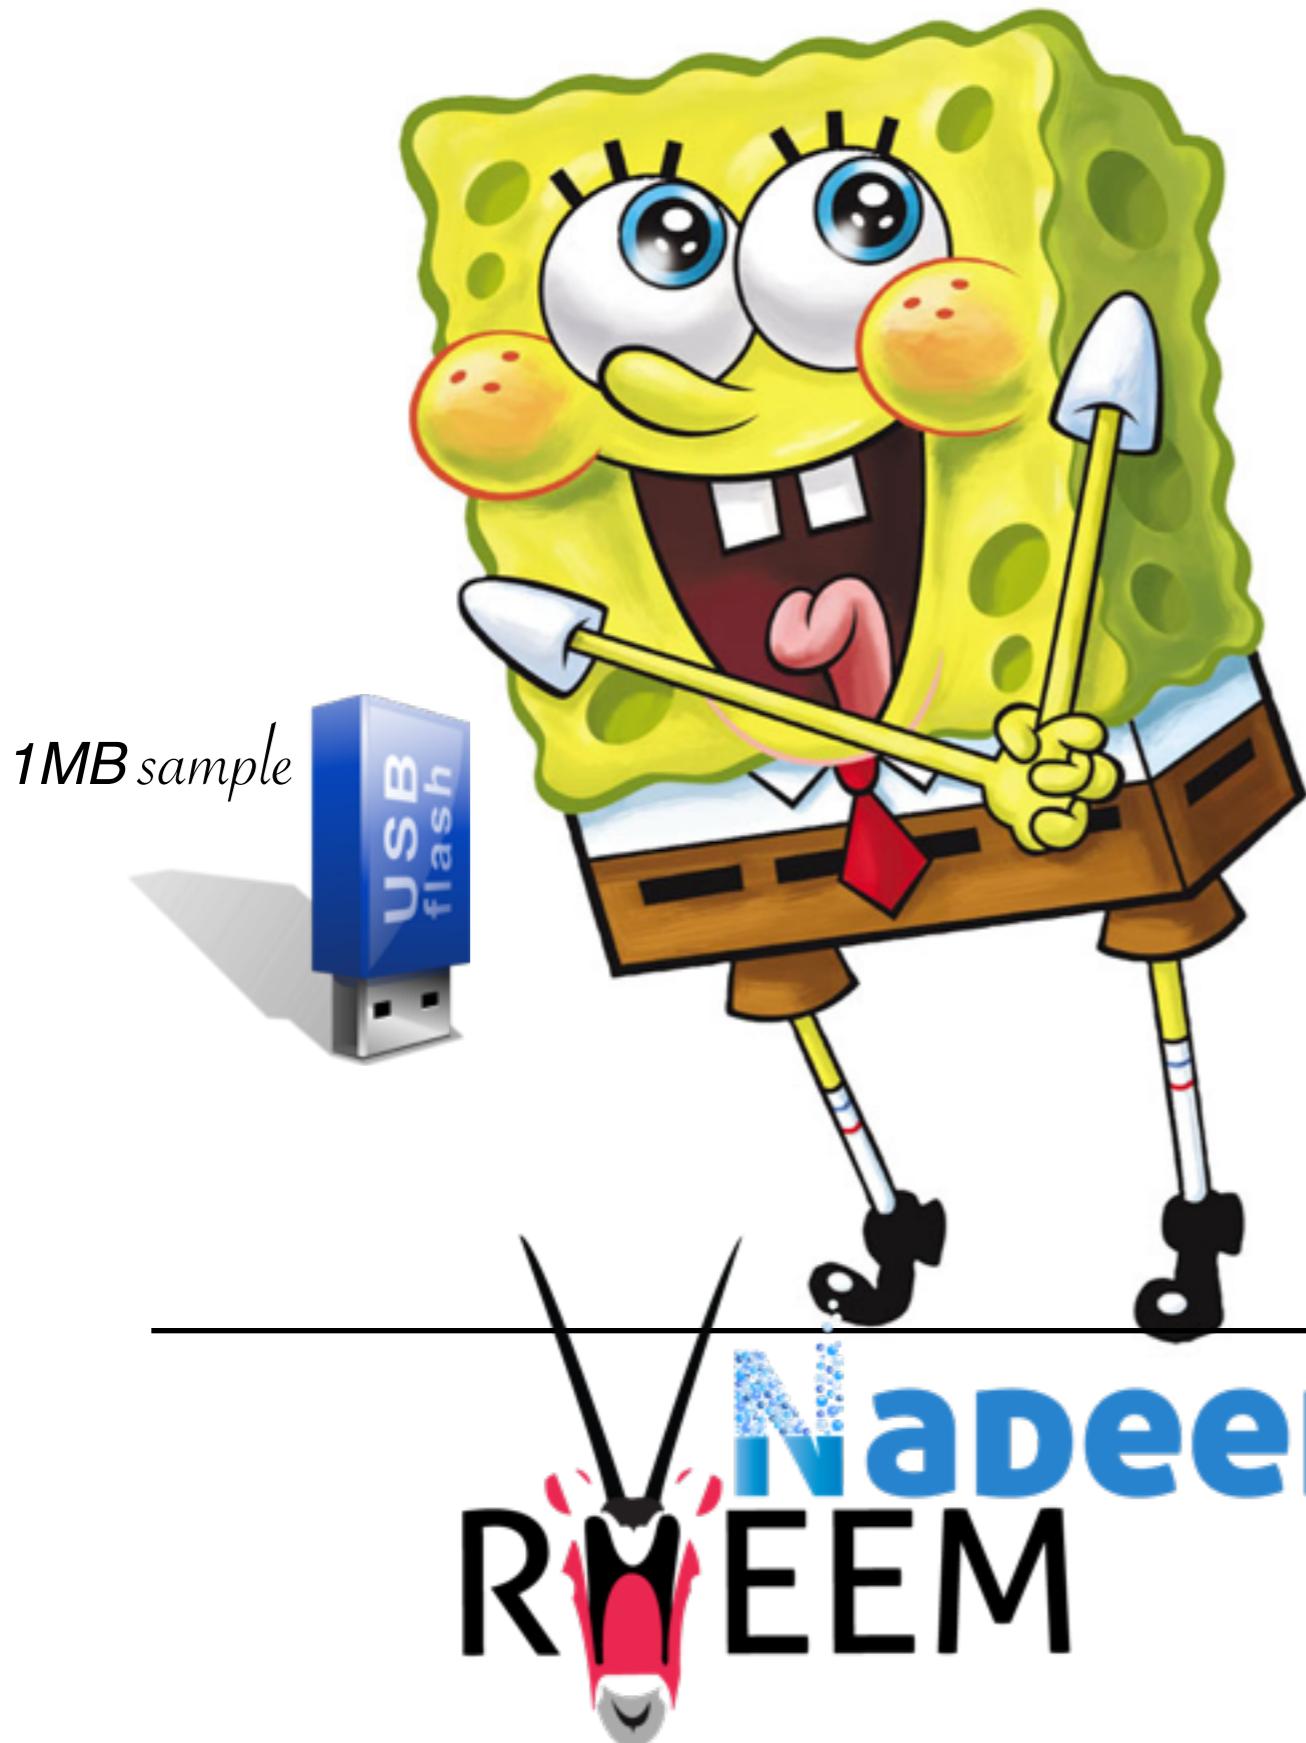
Bob@QCRI



Bob@QCRI



Bob@QCRI

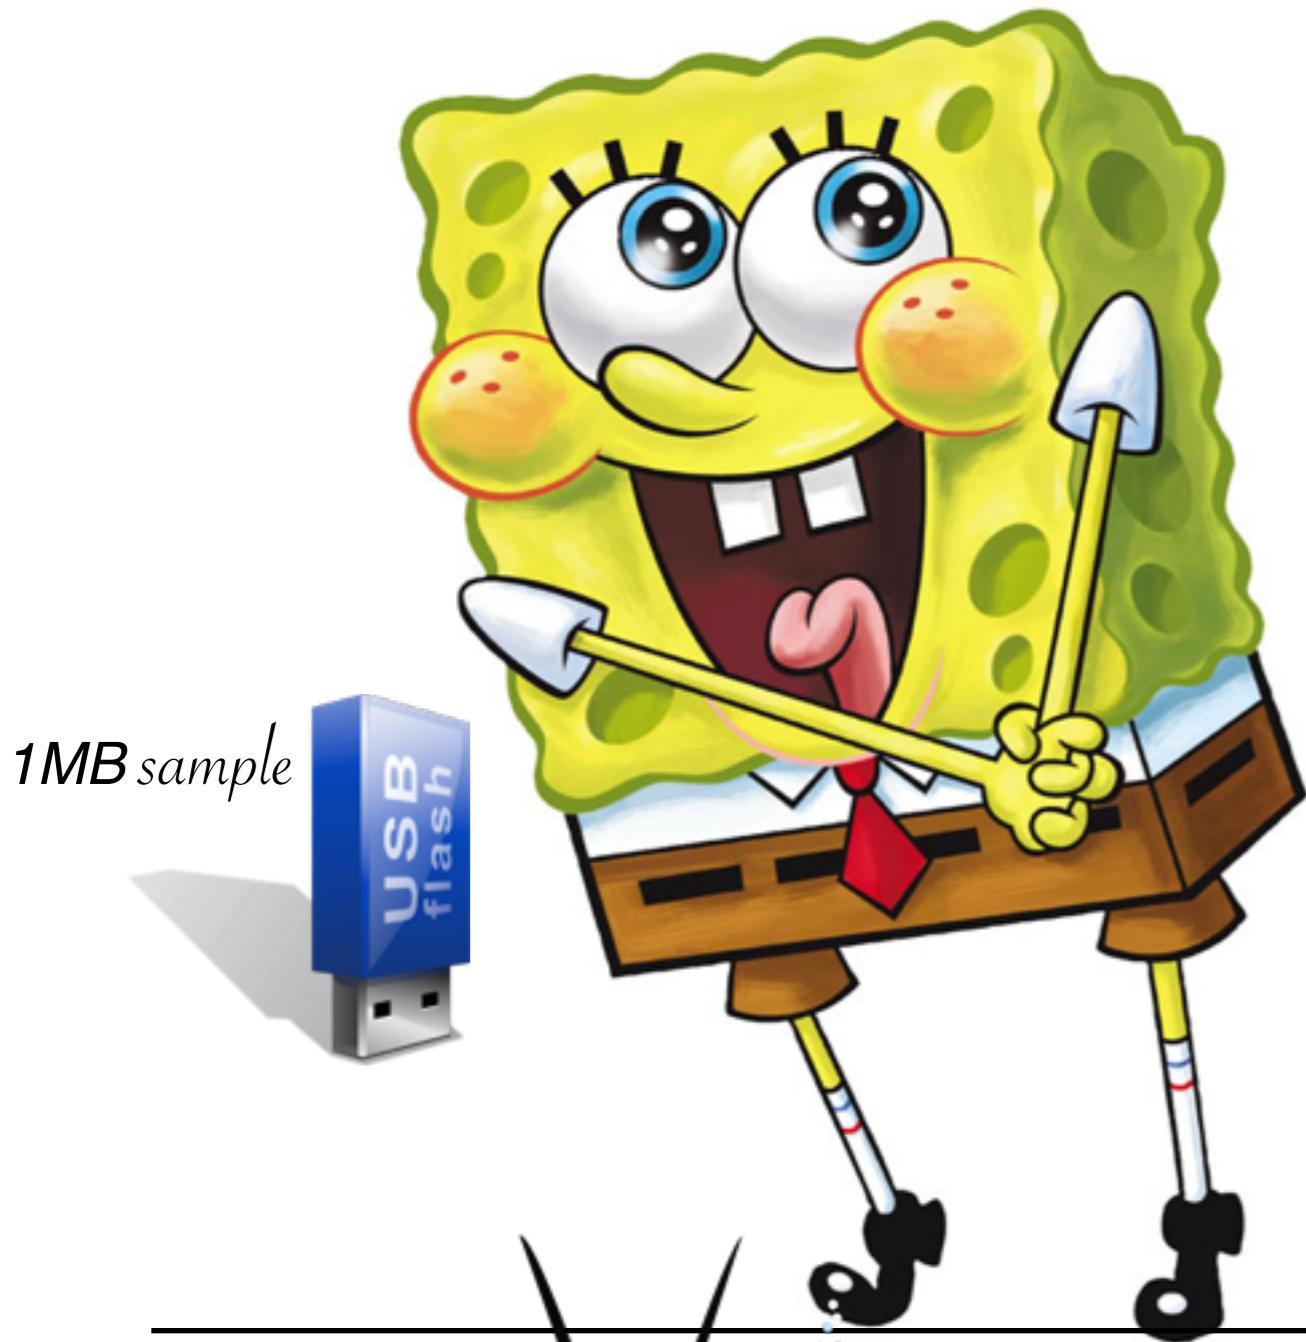


Q2:

Data quality rule

```
SELECT e1.id, e2.id  
FROM Employees e1, e2  
WHERE e1.salary > e2.salary  
AND e1.tax < e2.tax;
```

Bob@QCRI



Nadeef
RMEEM

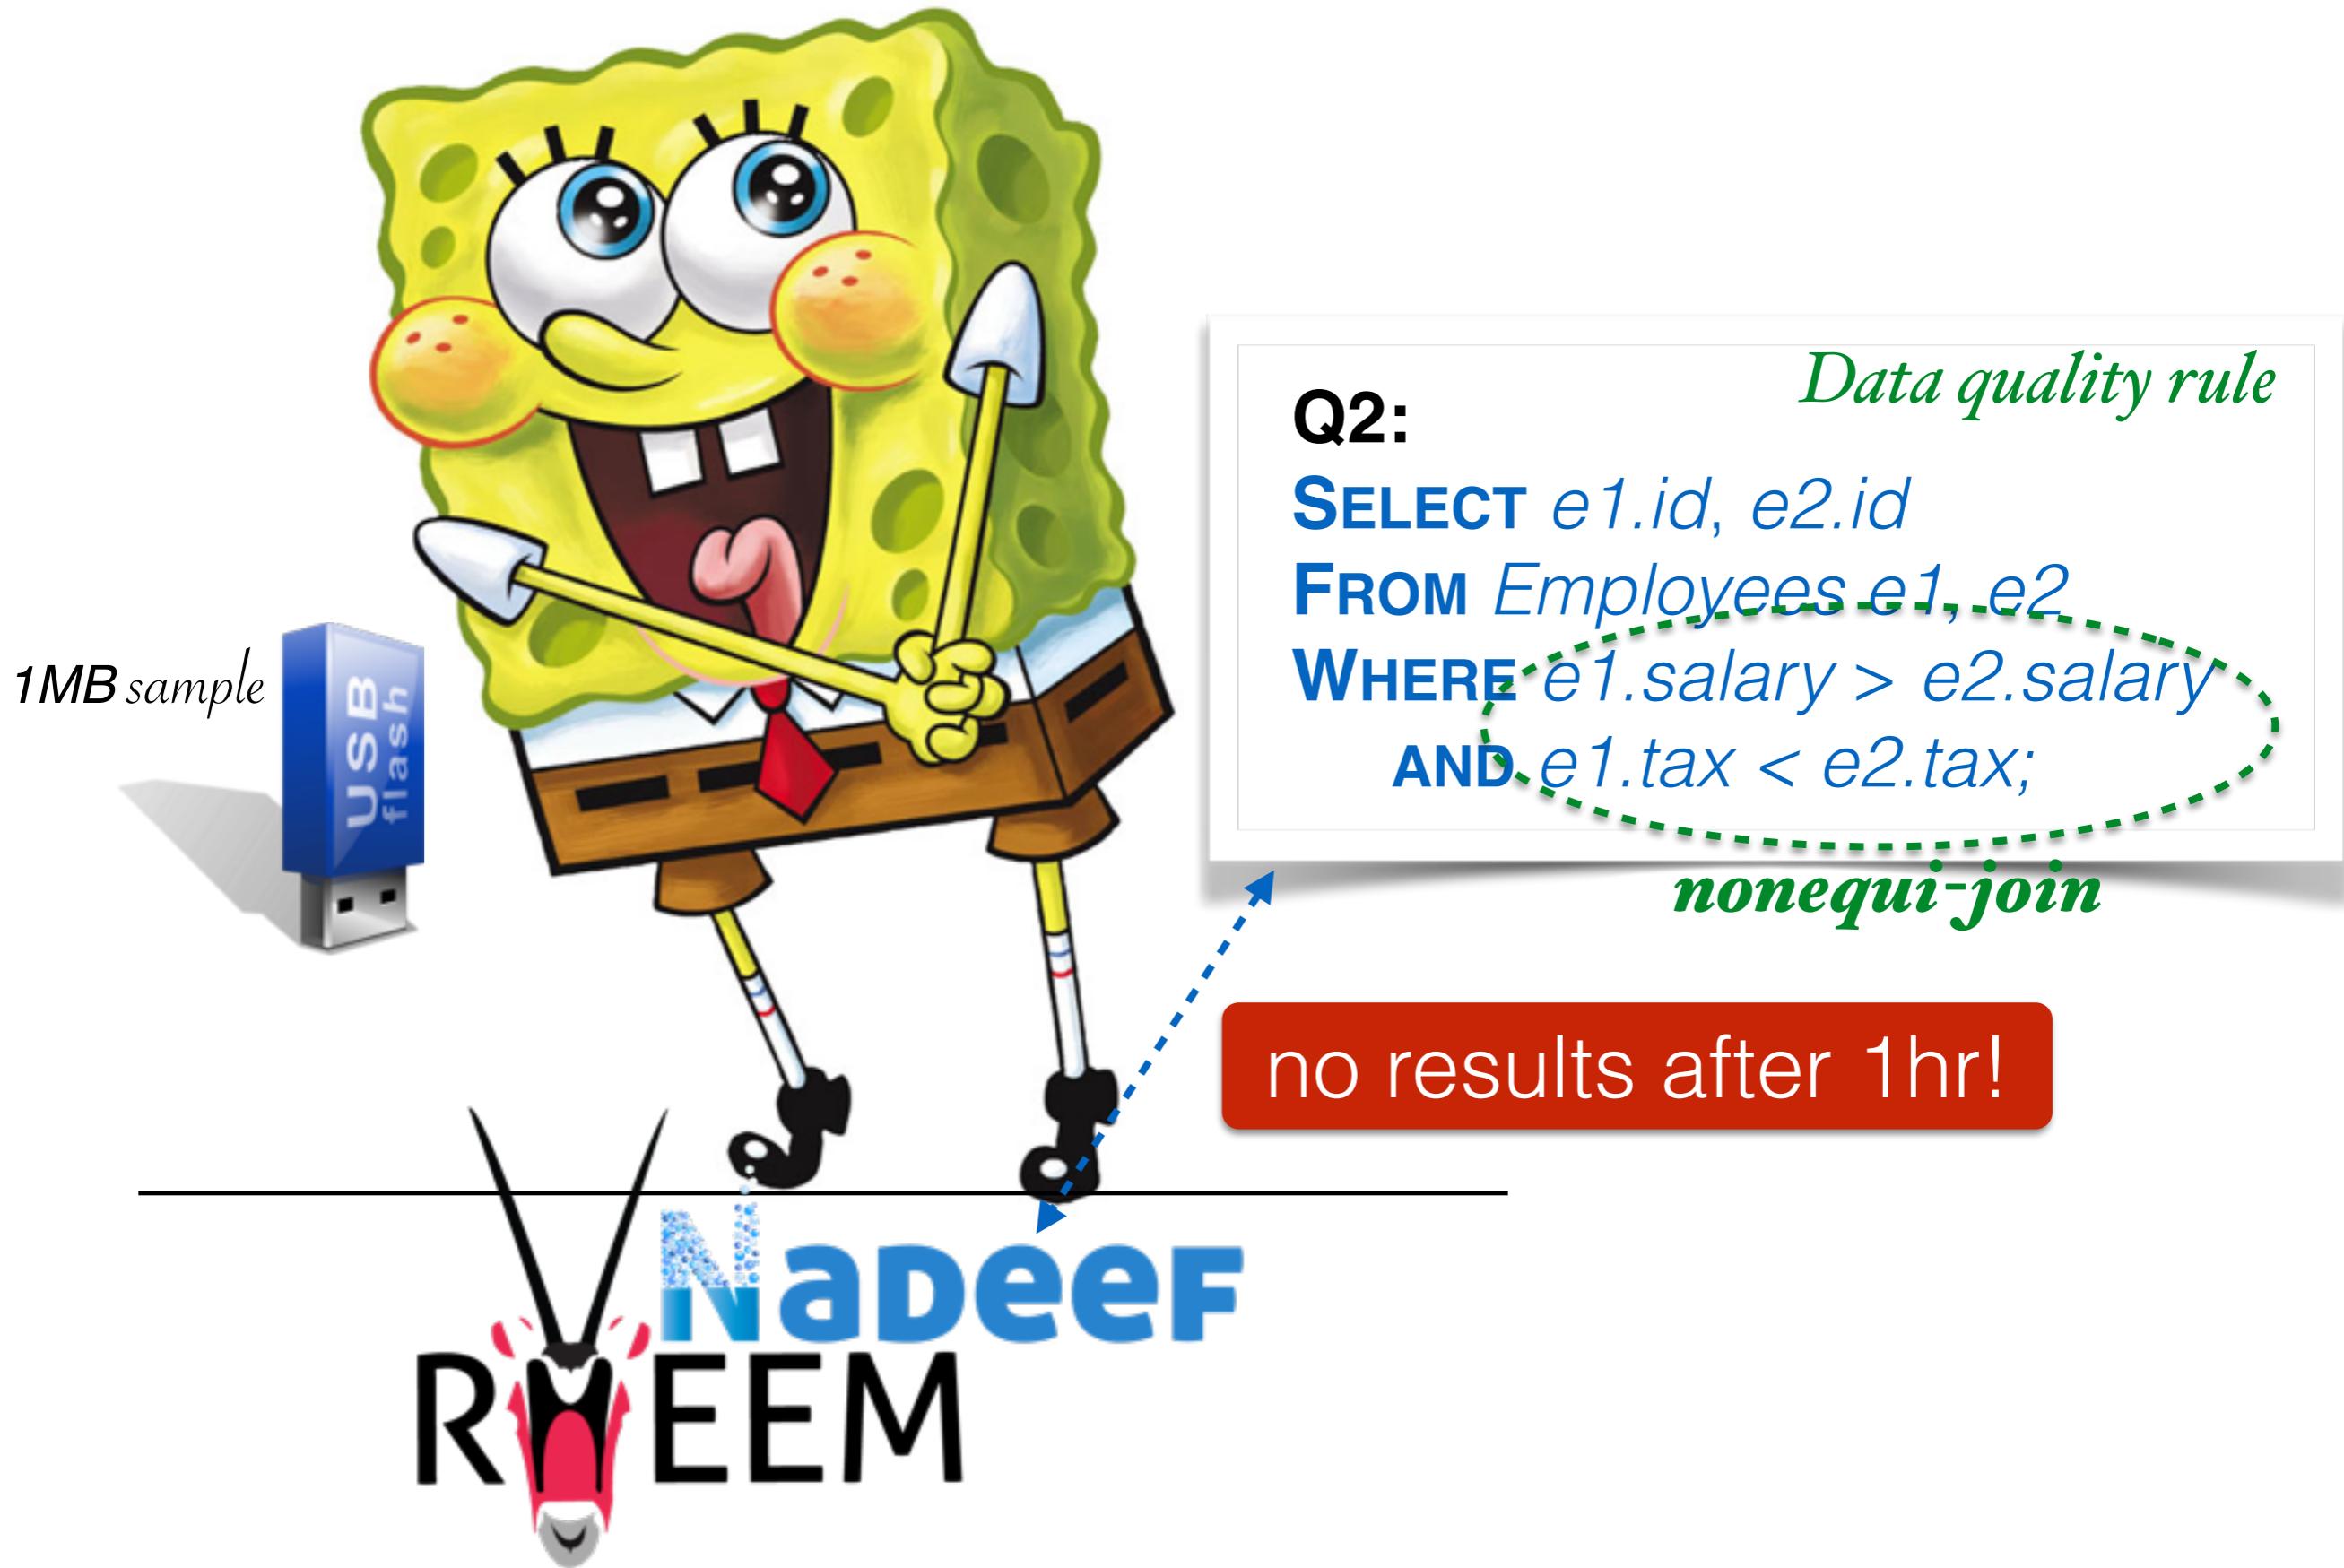
Q2:

SELECT e1.id, e2.id
FROM Employees e1, e2
WHERE e1.salary > e2.salary
AND e1.tax < e2.tax;

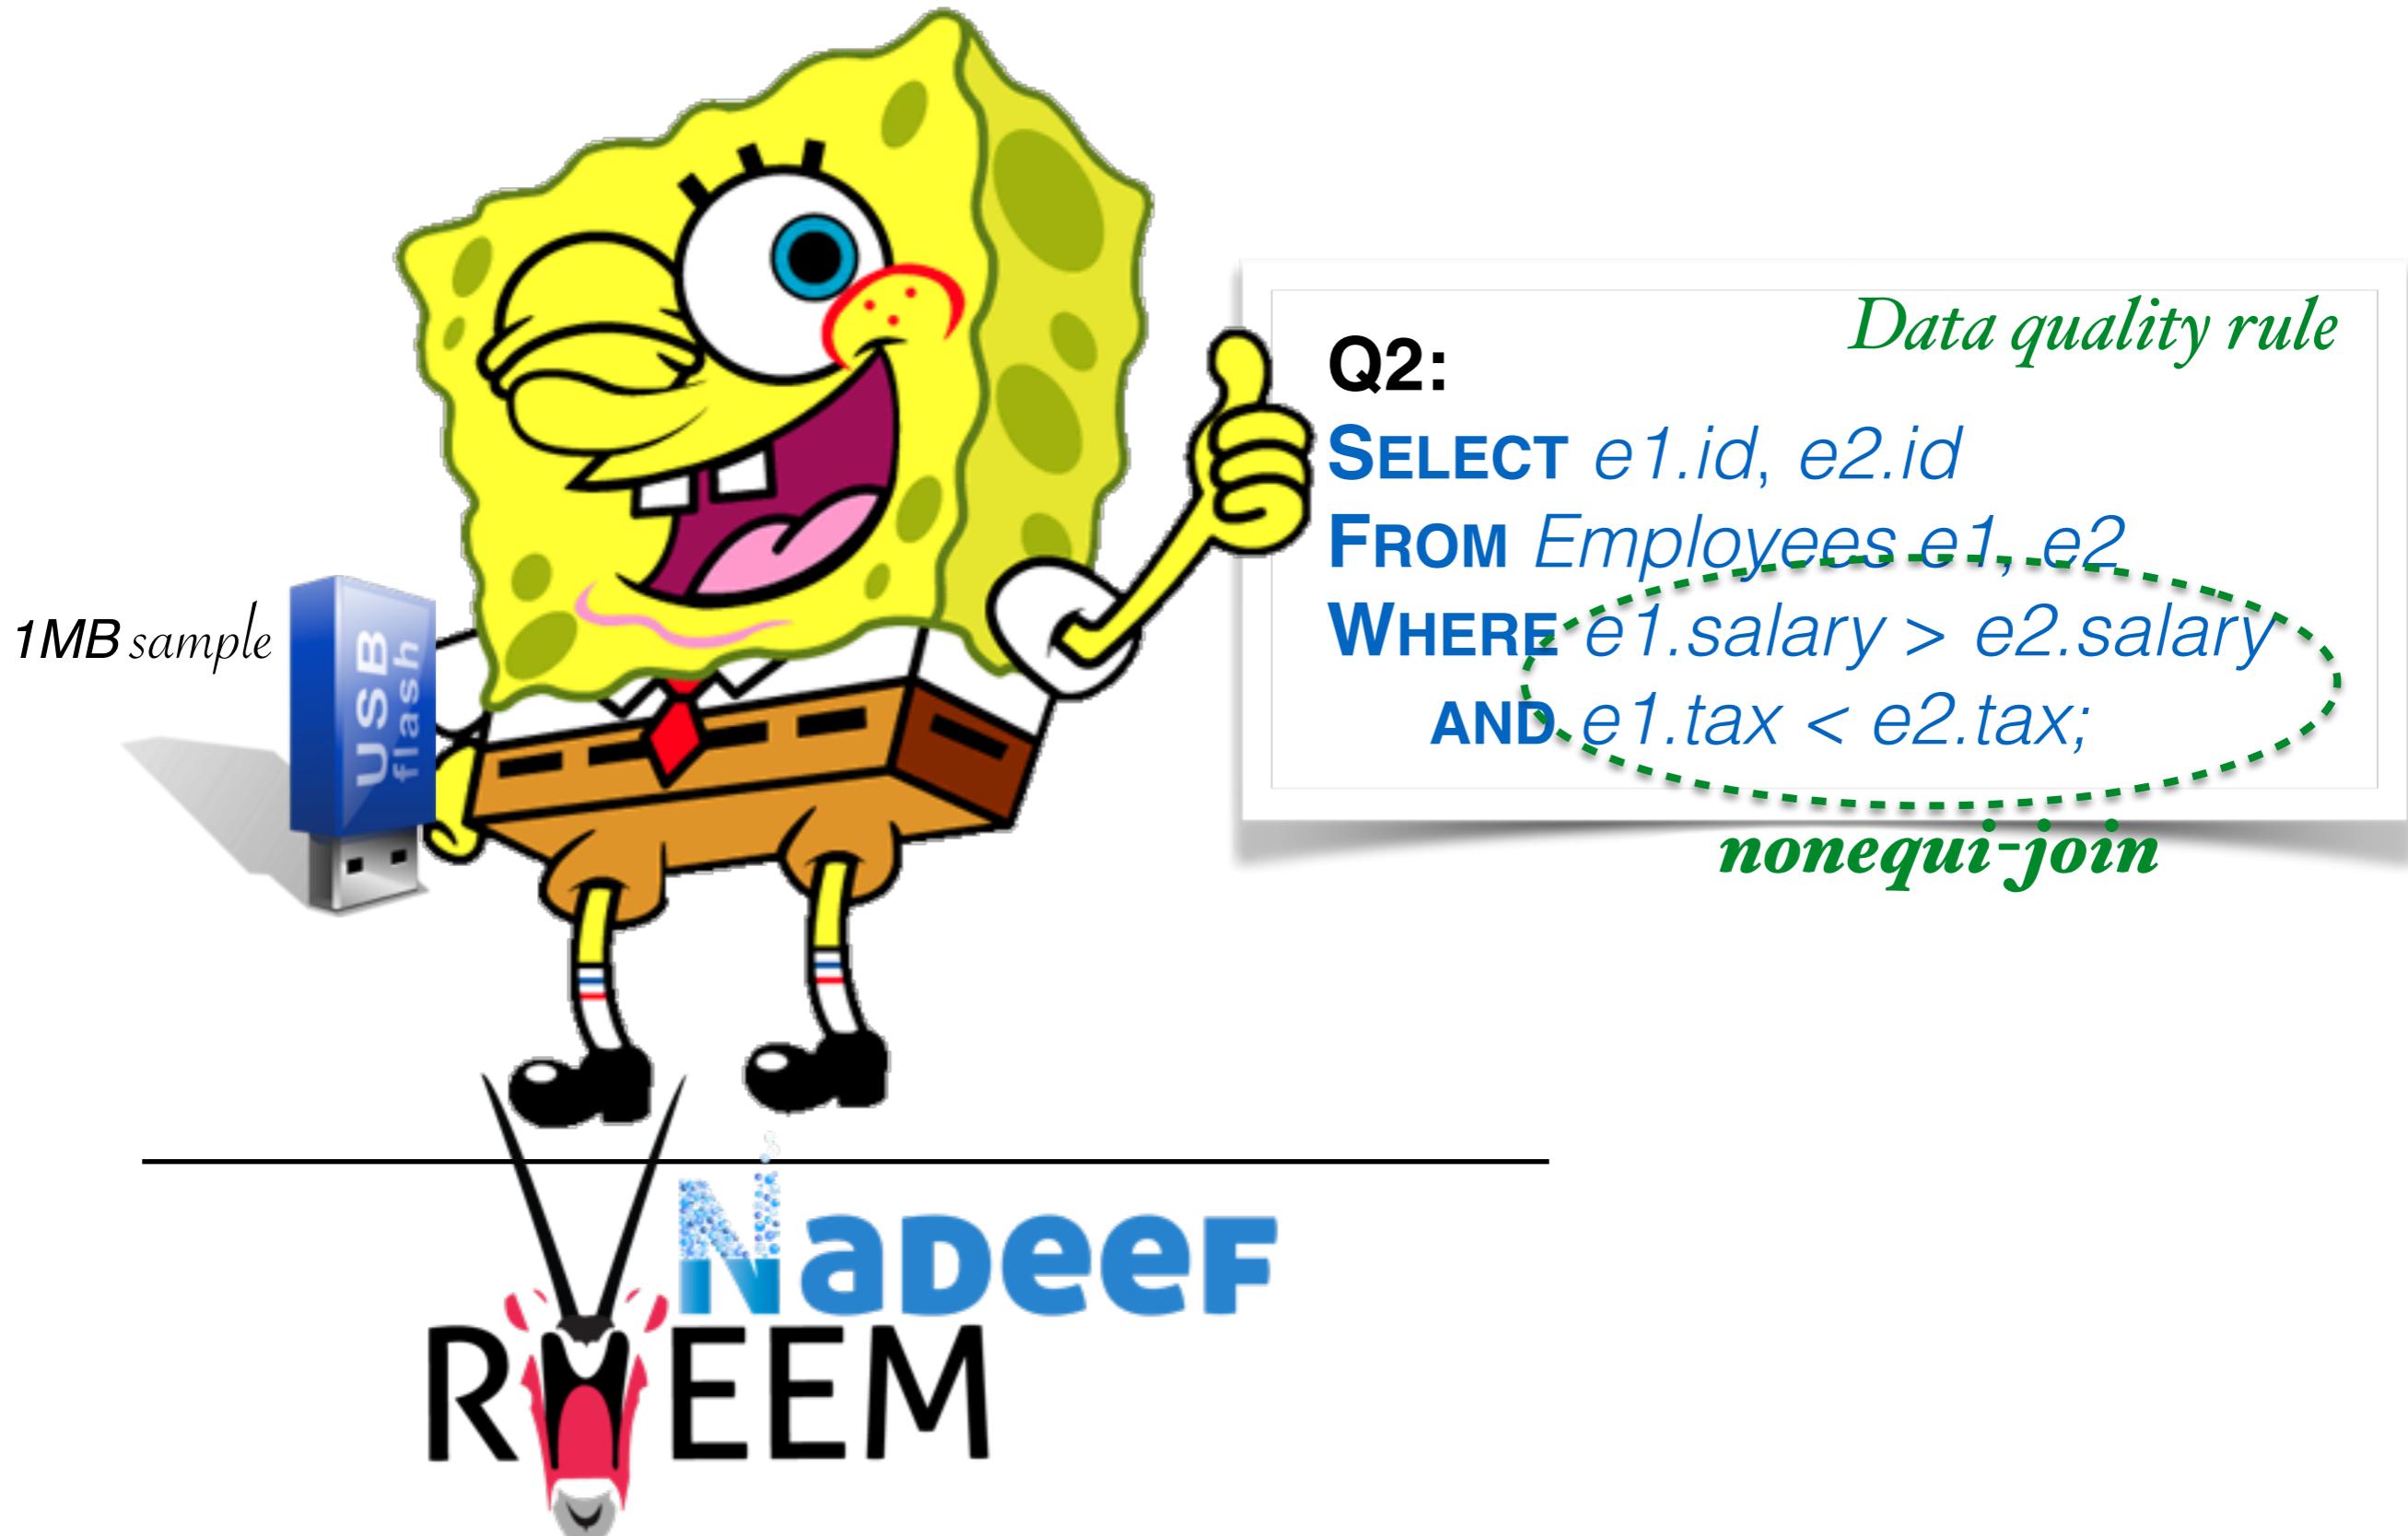
Data quality rule

nonequi-join

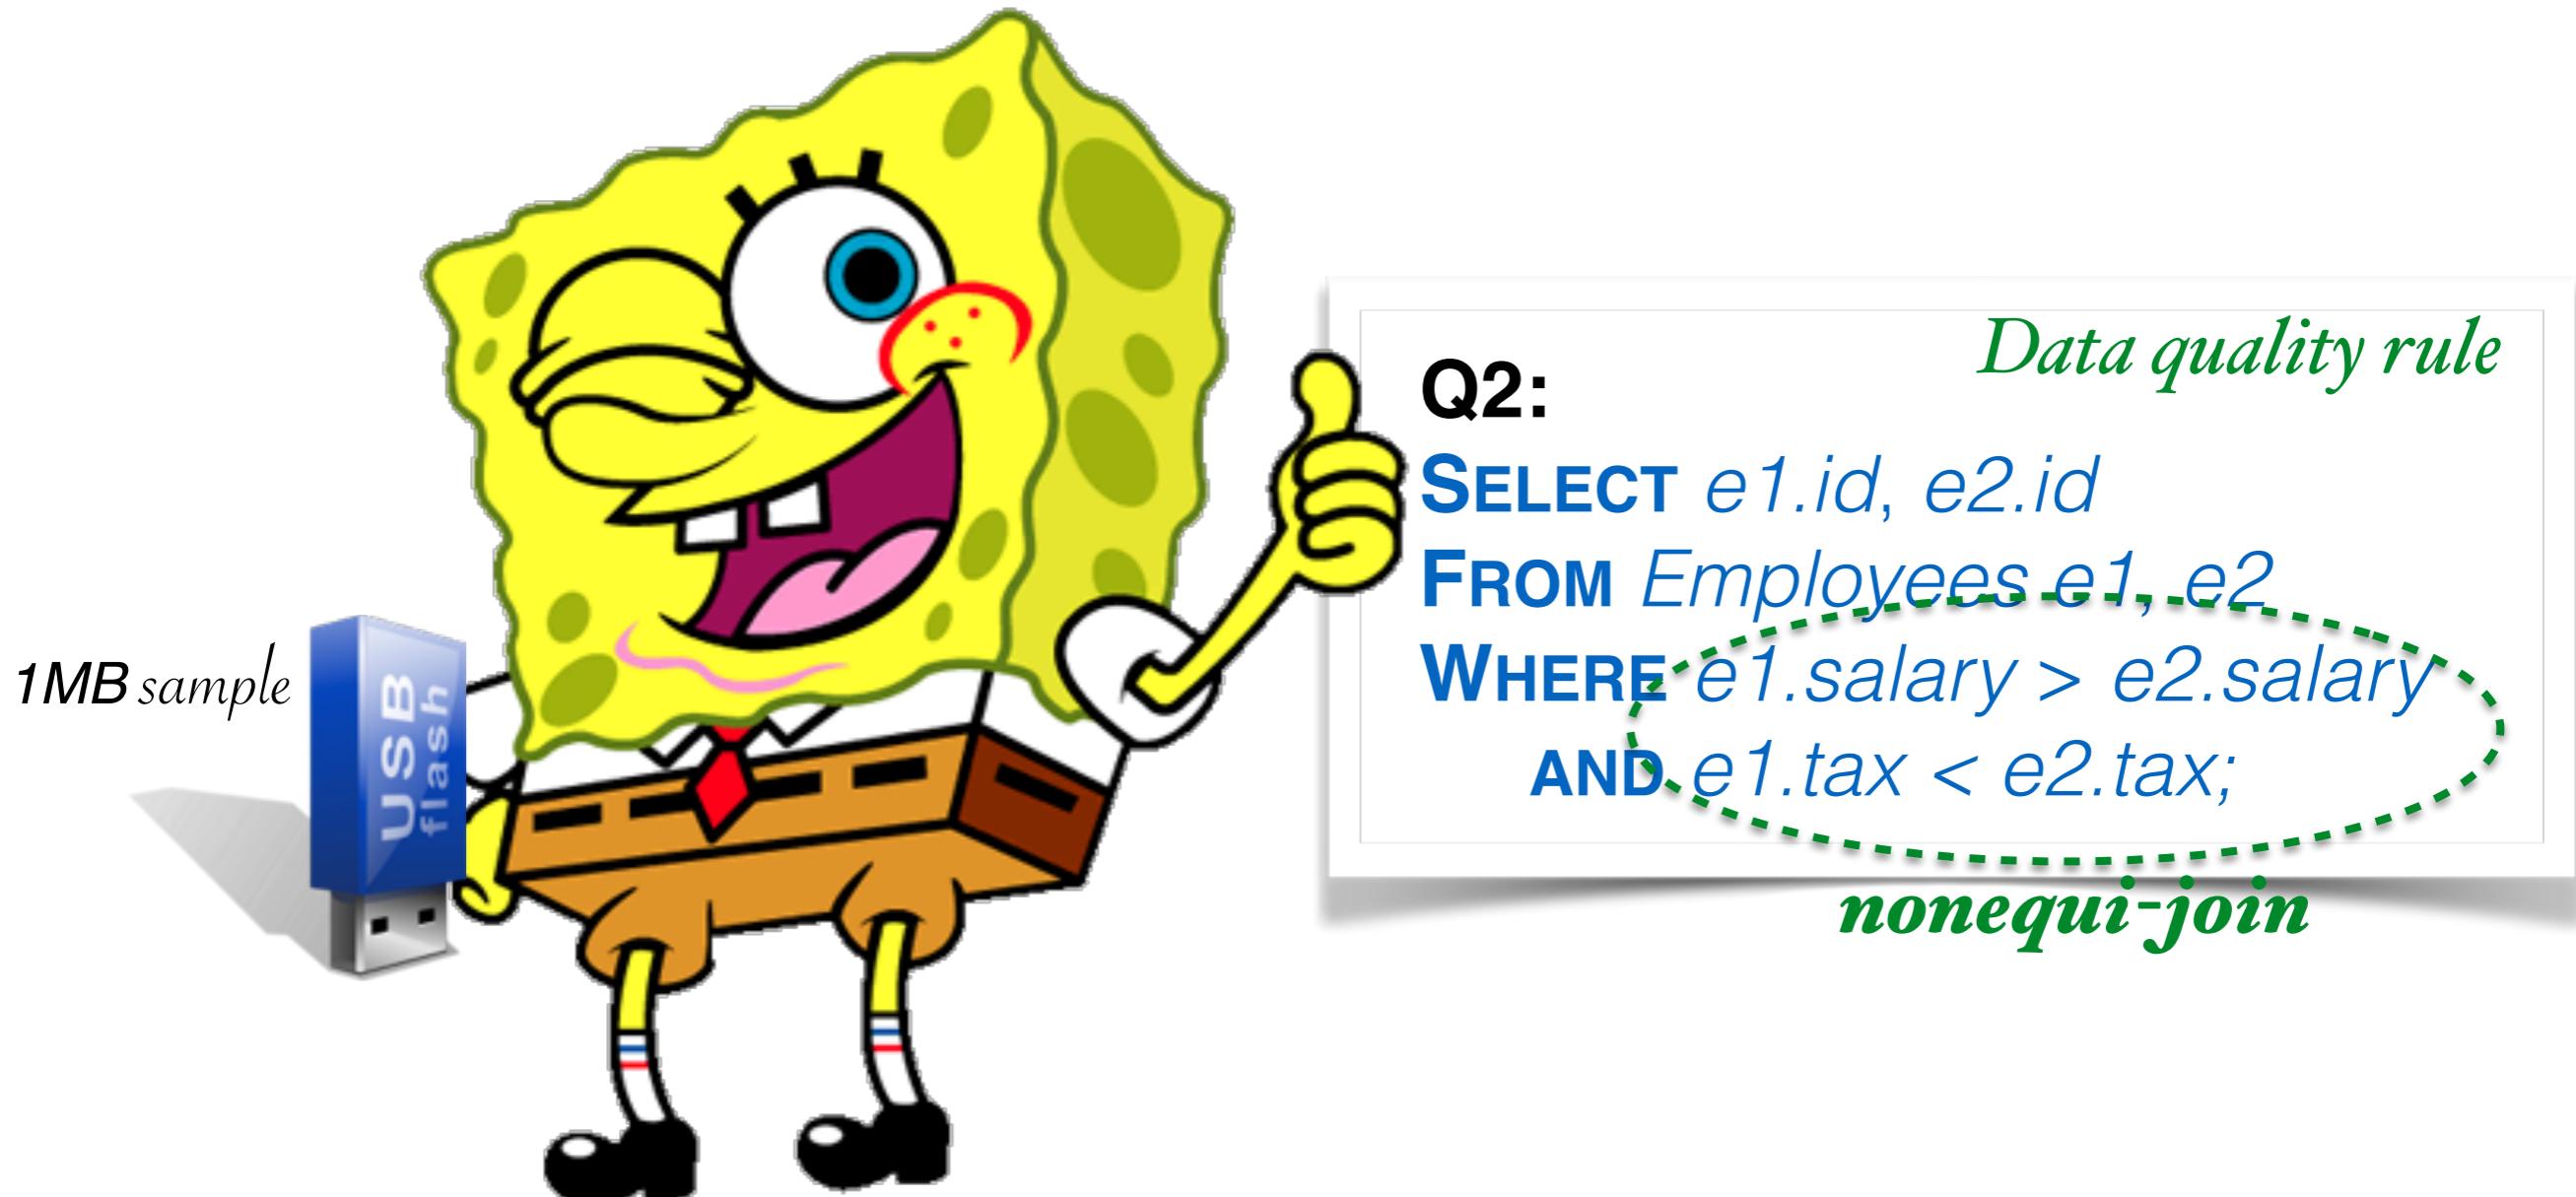
Bob@QCRI



Bob@QCRI



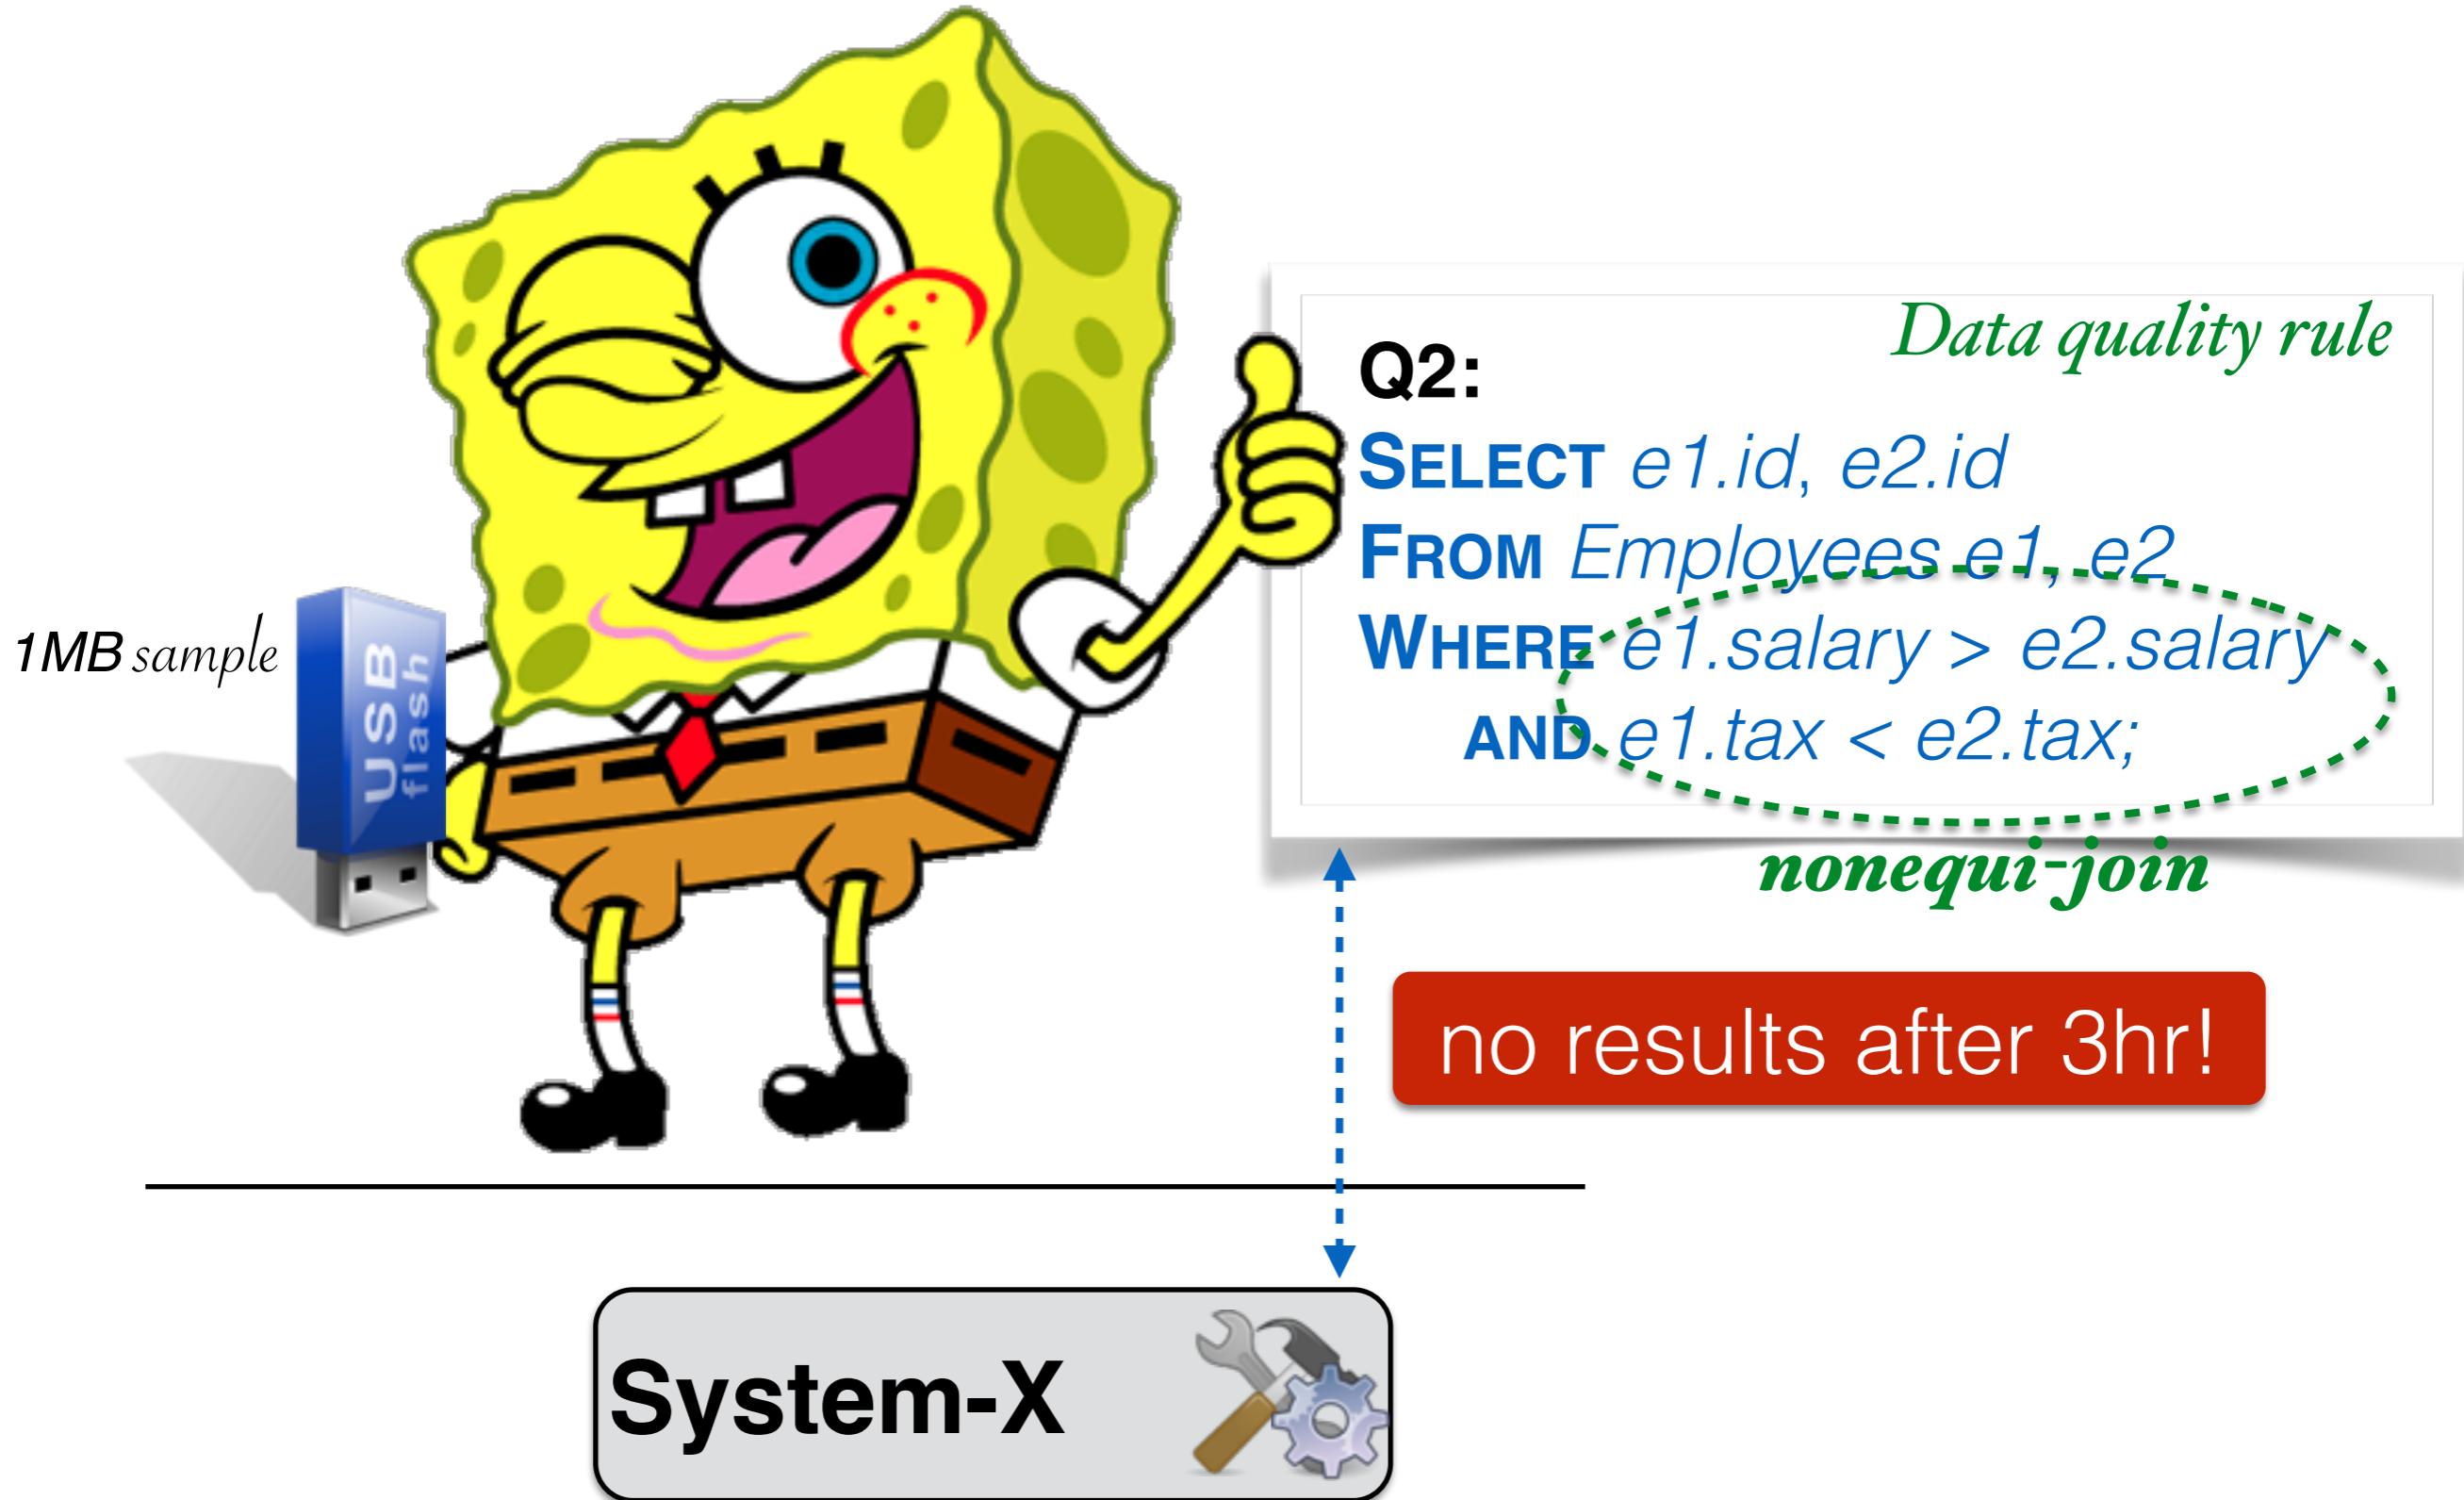
Bob@QCRI



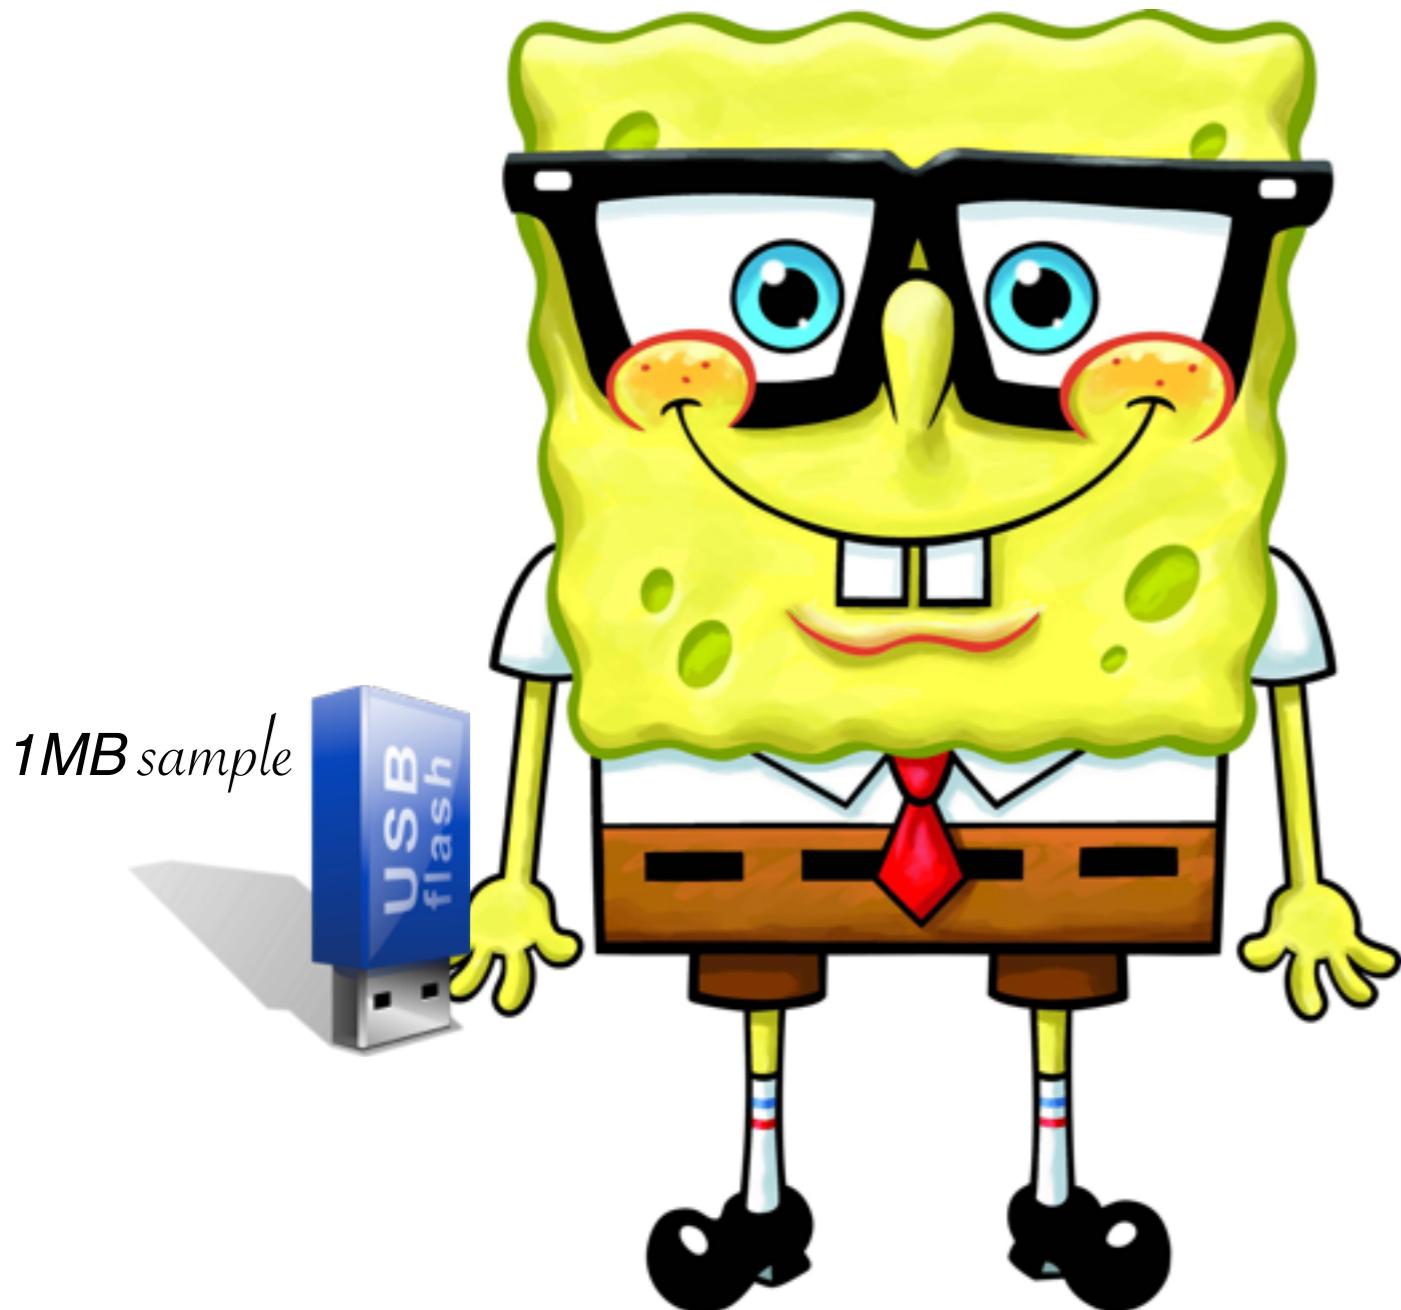
System-X



Bob@QCRI



Bob@QCRI



Q2:

SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
AND e1.tax < e2.tax;

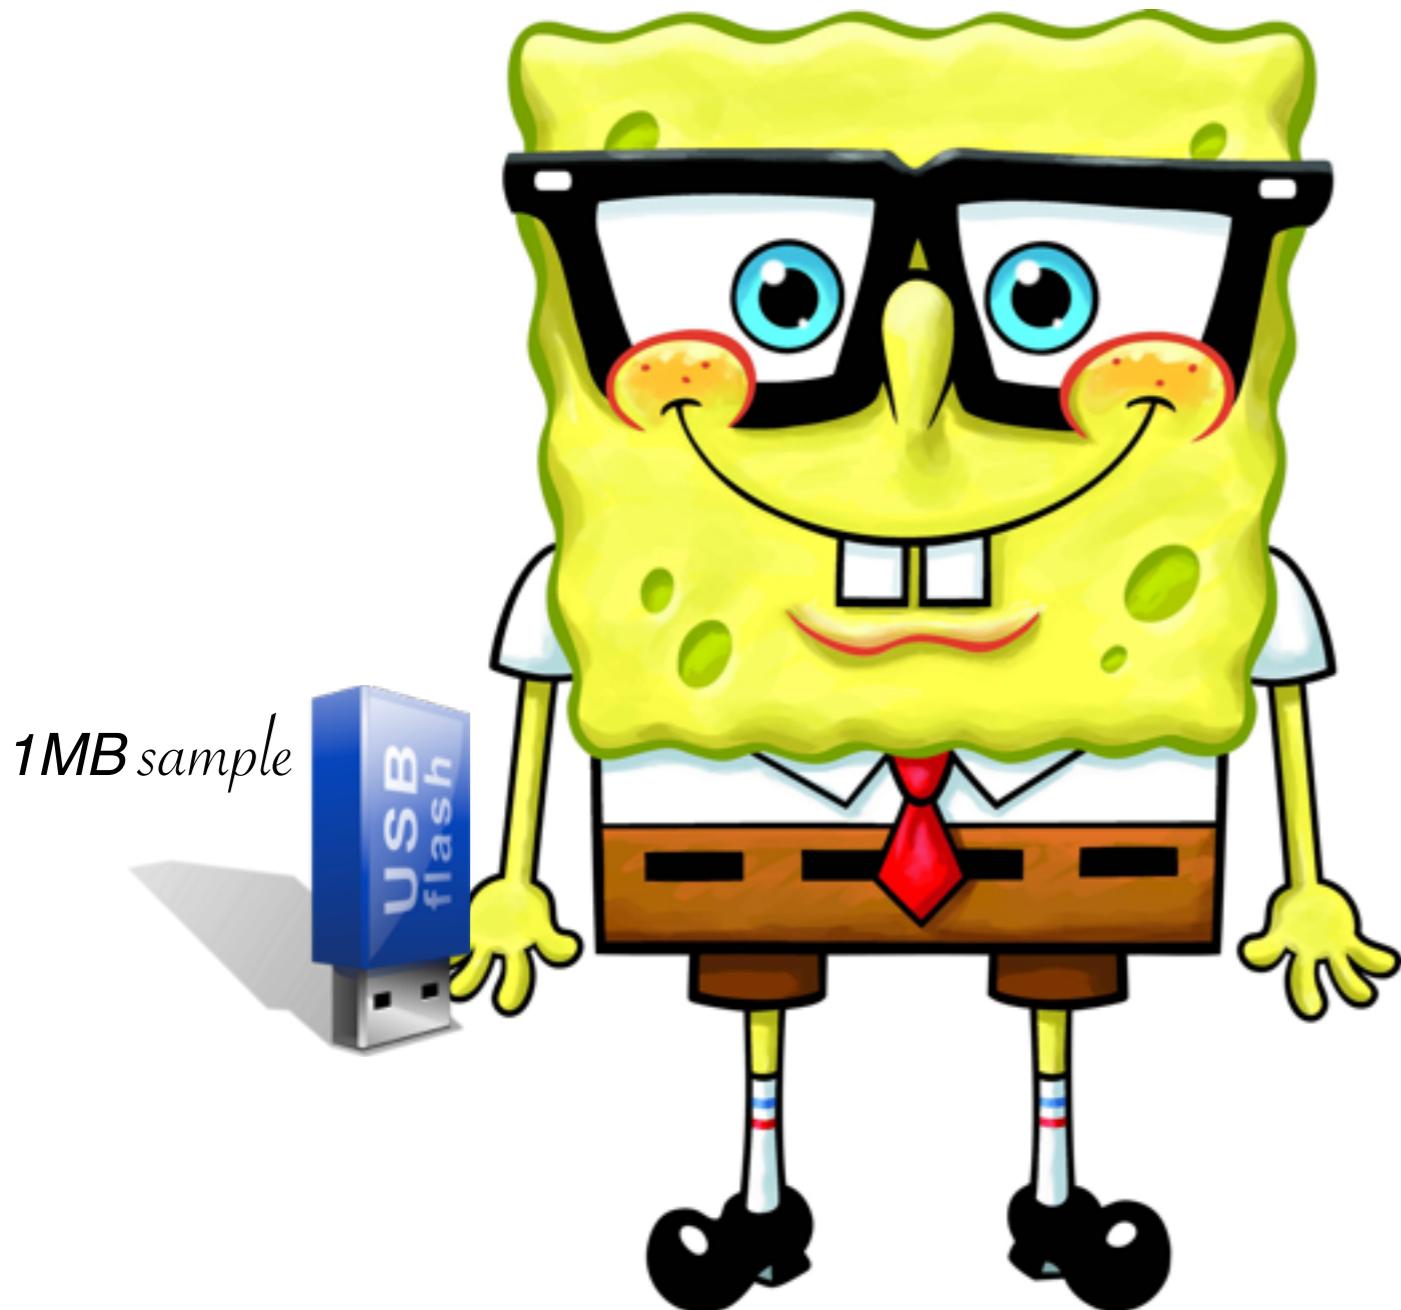
Data quality rule

nonequi-join

System-X



Bob@QCRI



Q2:

SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
AND e1.tax < e2.tax;

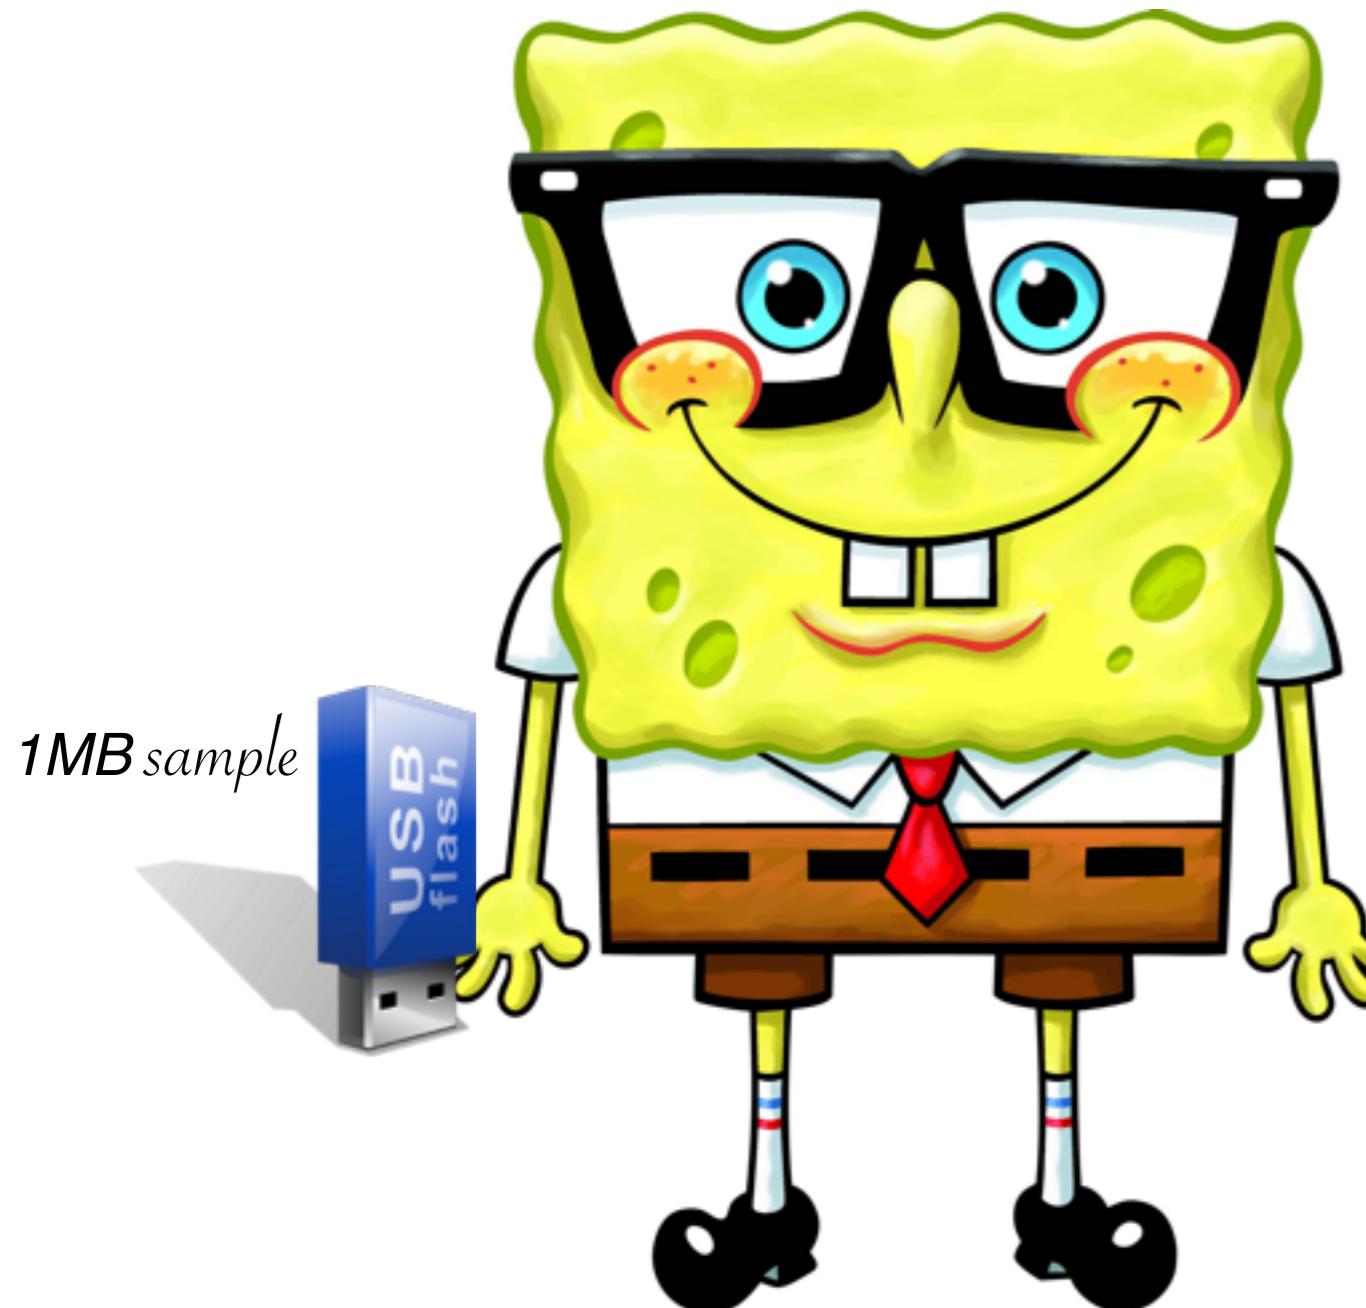
Data quality rule

nonequi-join

DBMS-X



Bob@QCRI



Q2:

SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
AND e1.tax < e2.tax;

Data quality rule

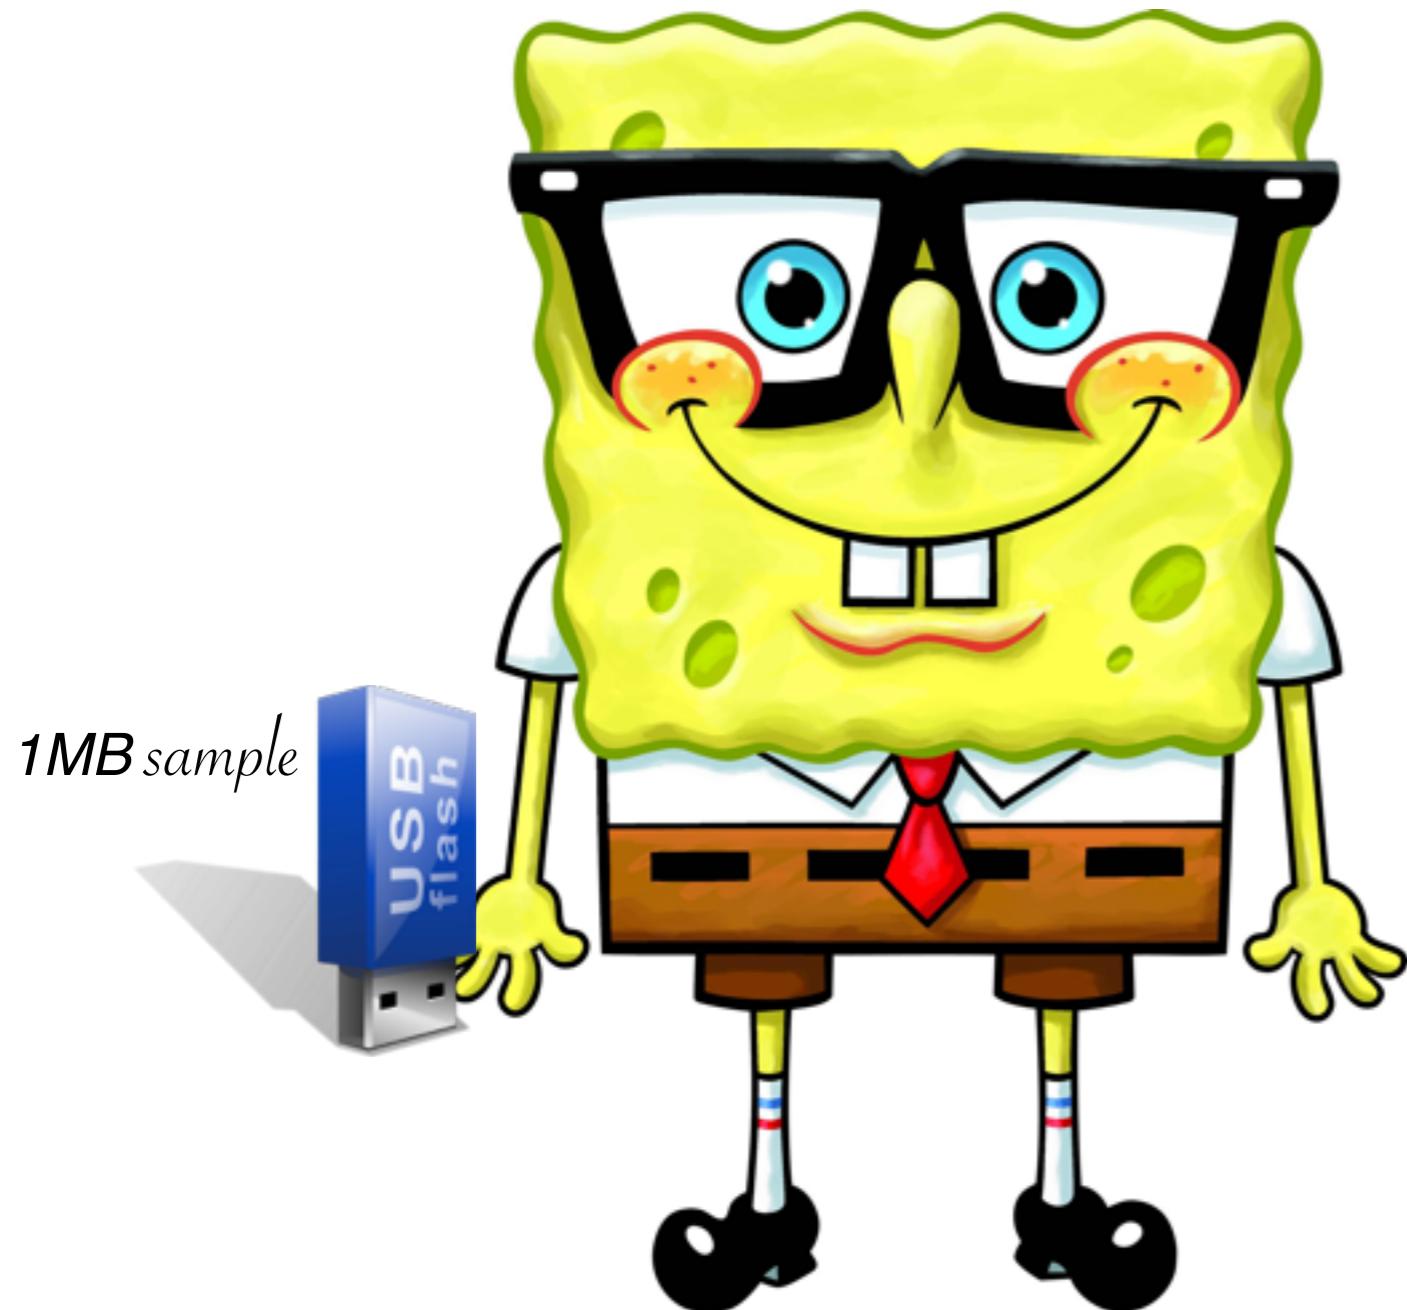
nonequi-join

results after 2hr!

DBMS-X



Bob@QCRI



Q2:

SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
AND e1.tax < e2.tax;

Data quality rule

nonequi-join

results after 2hr!

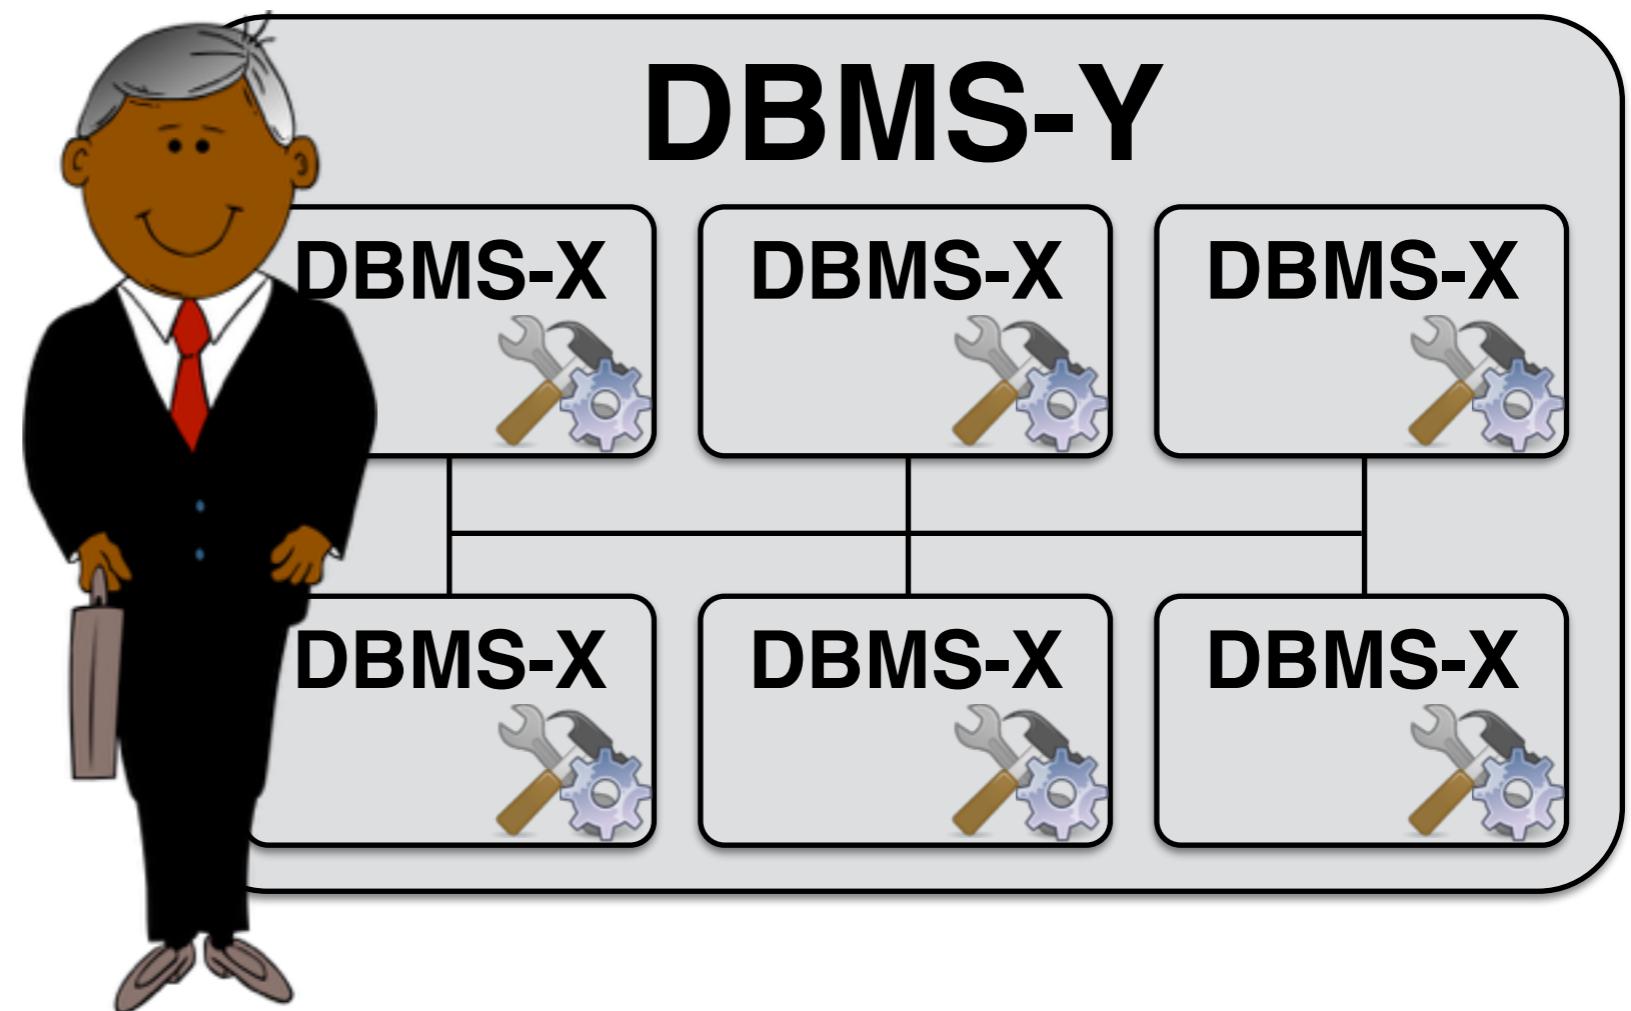
DBMS-X



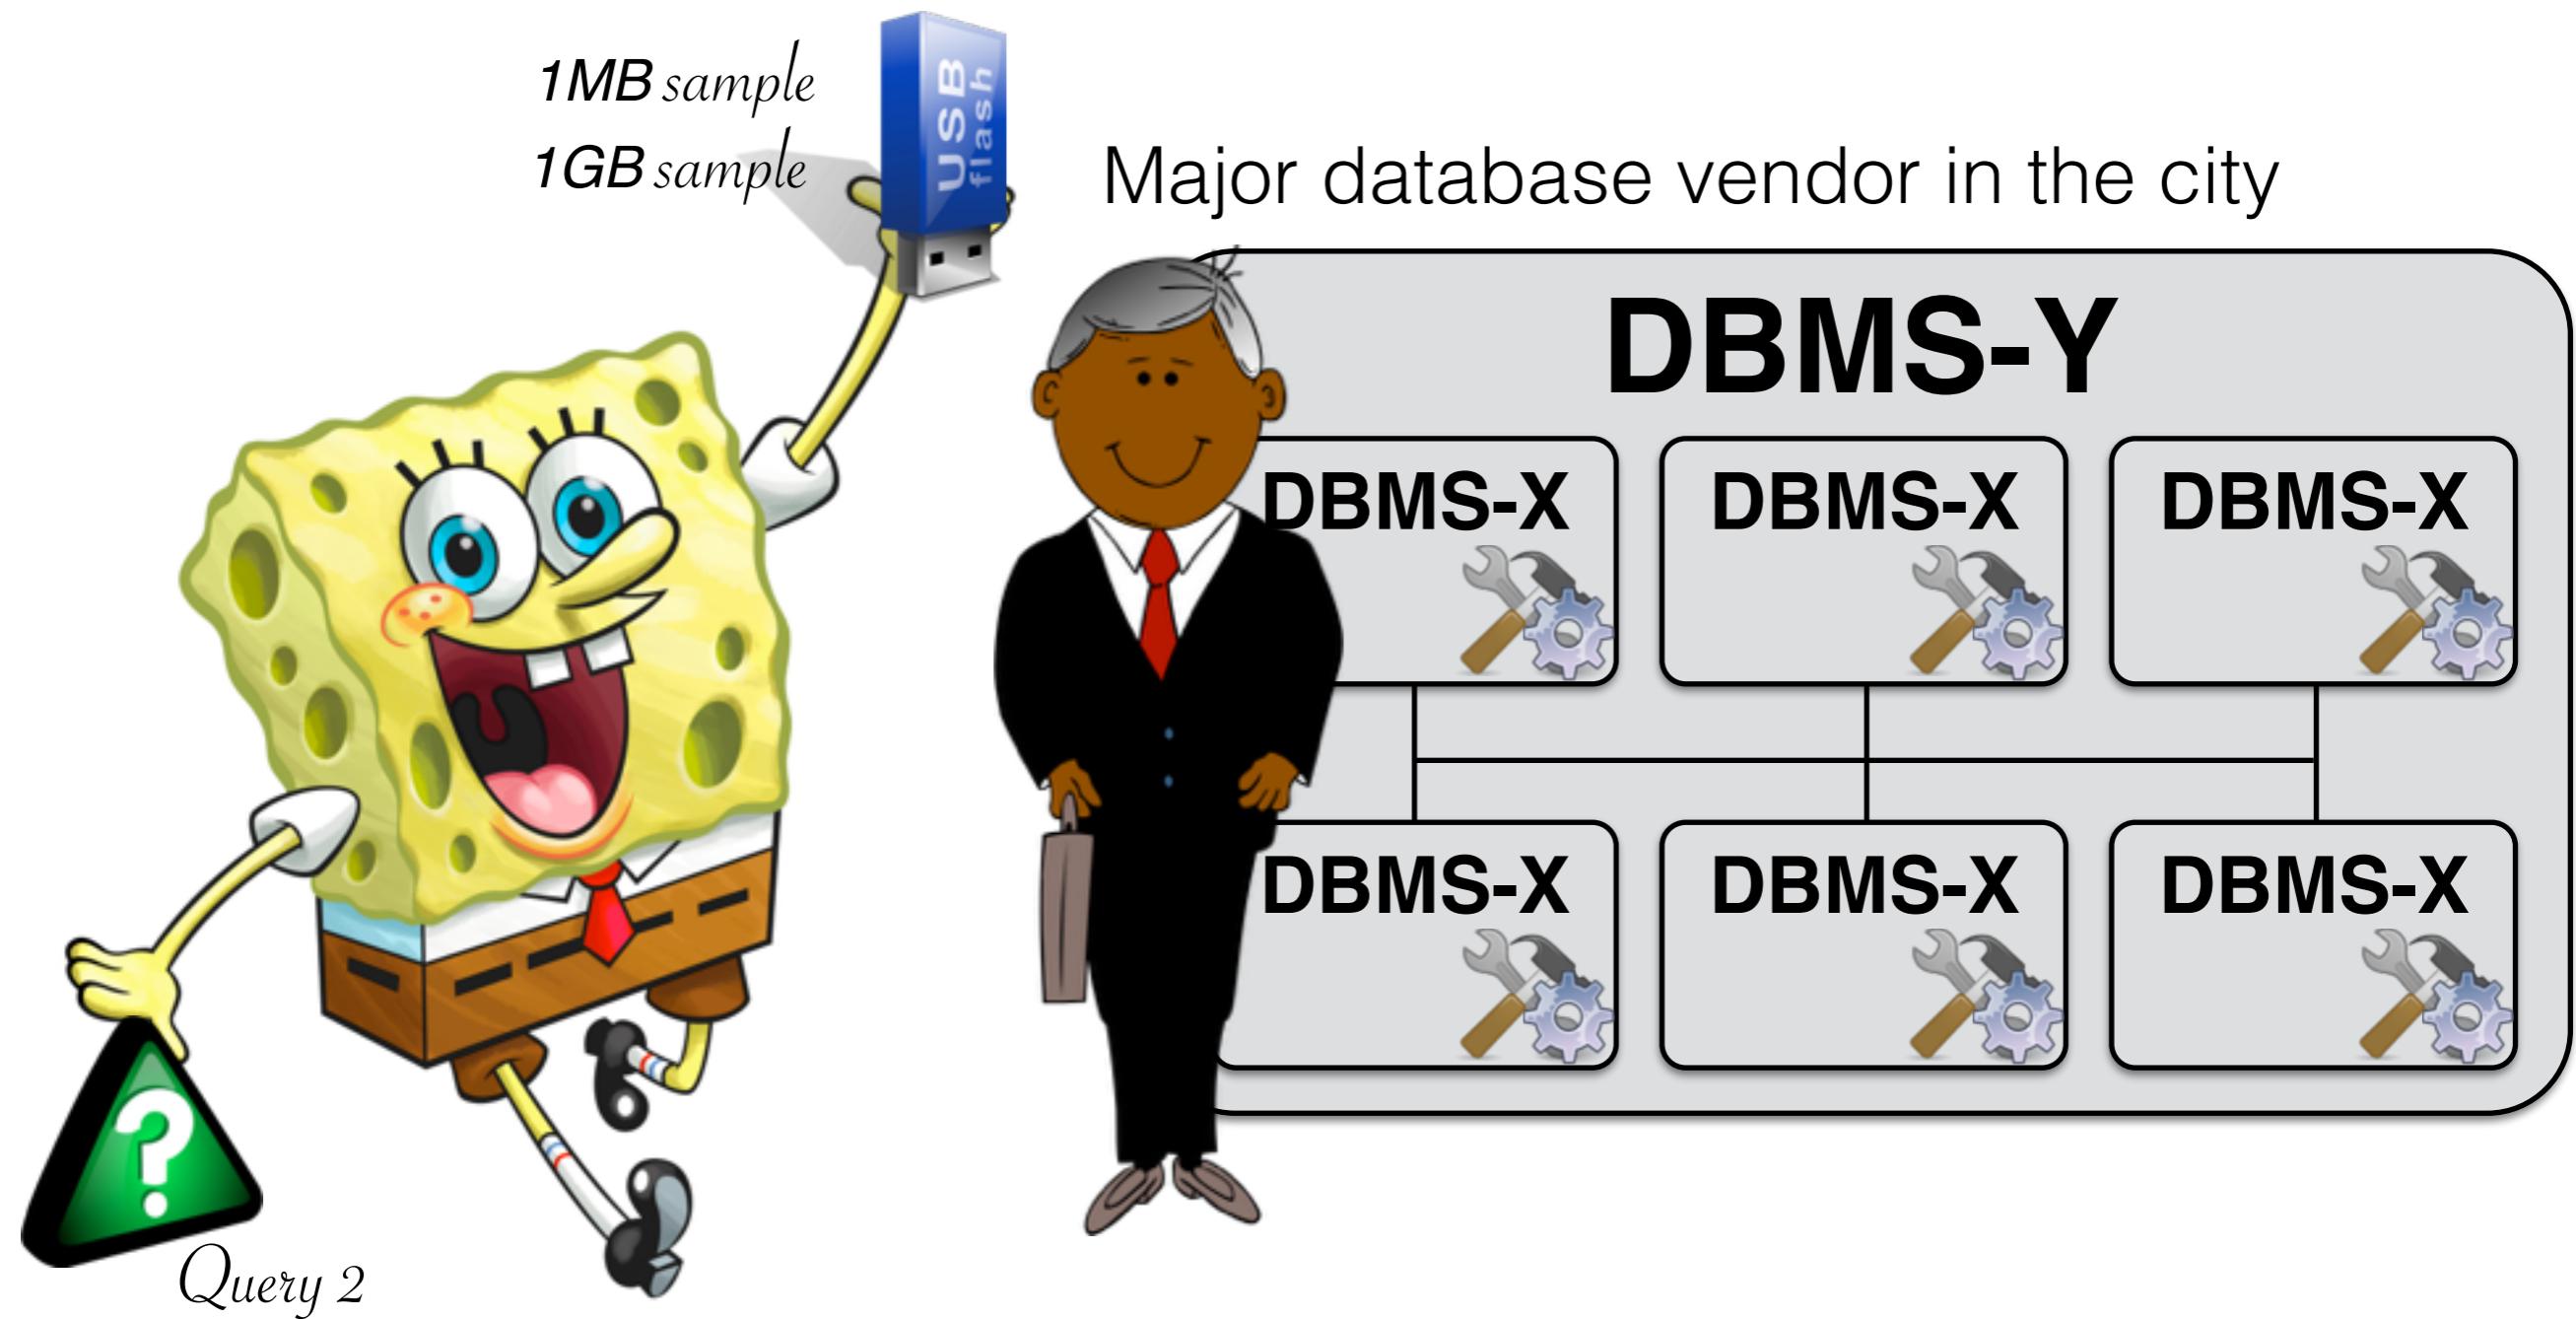
$\times + \sigma$

Bob@QCRI

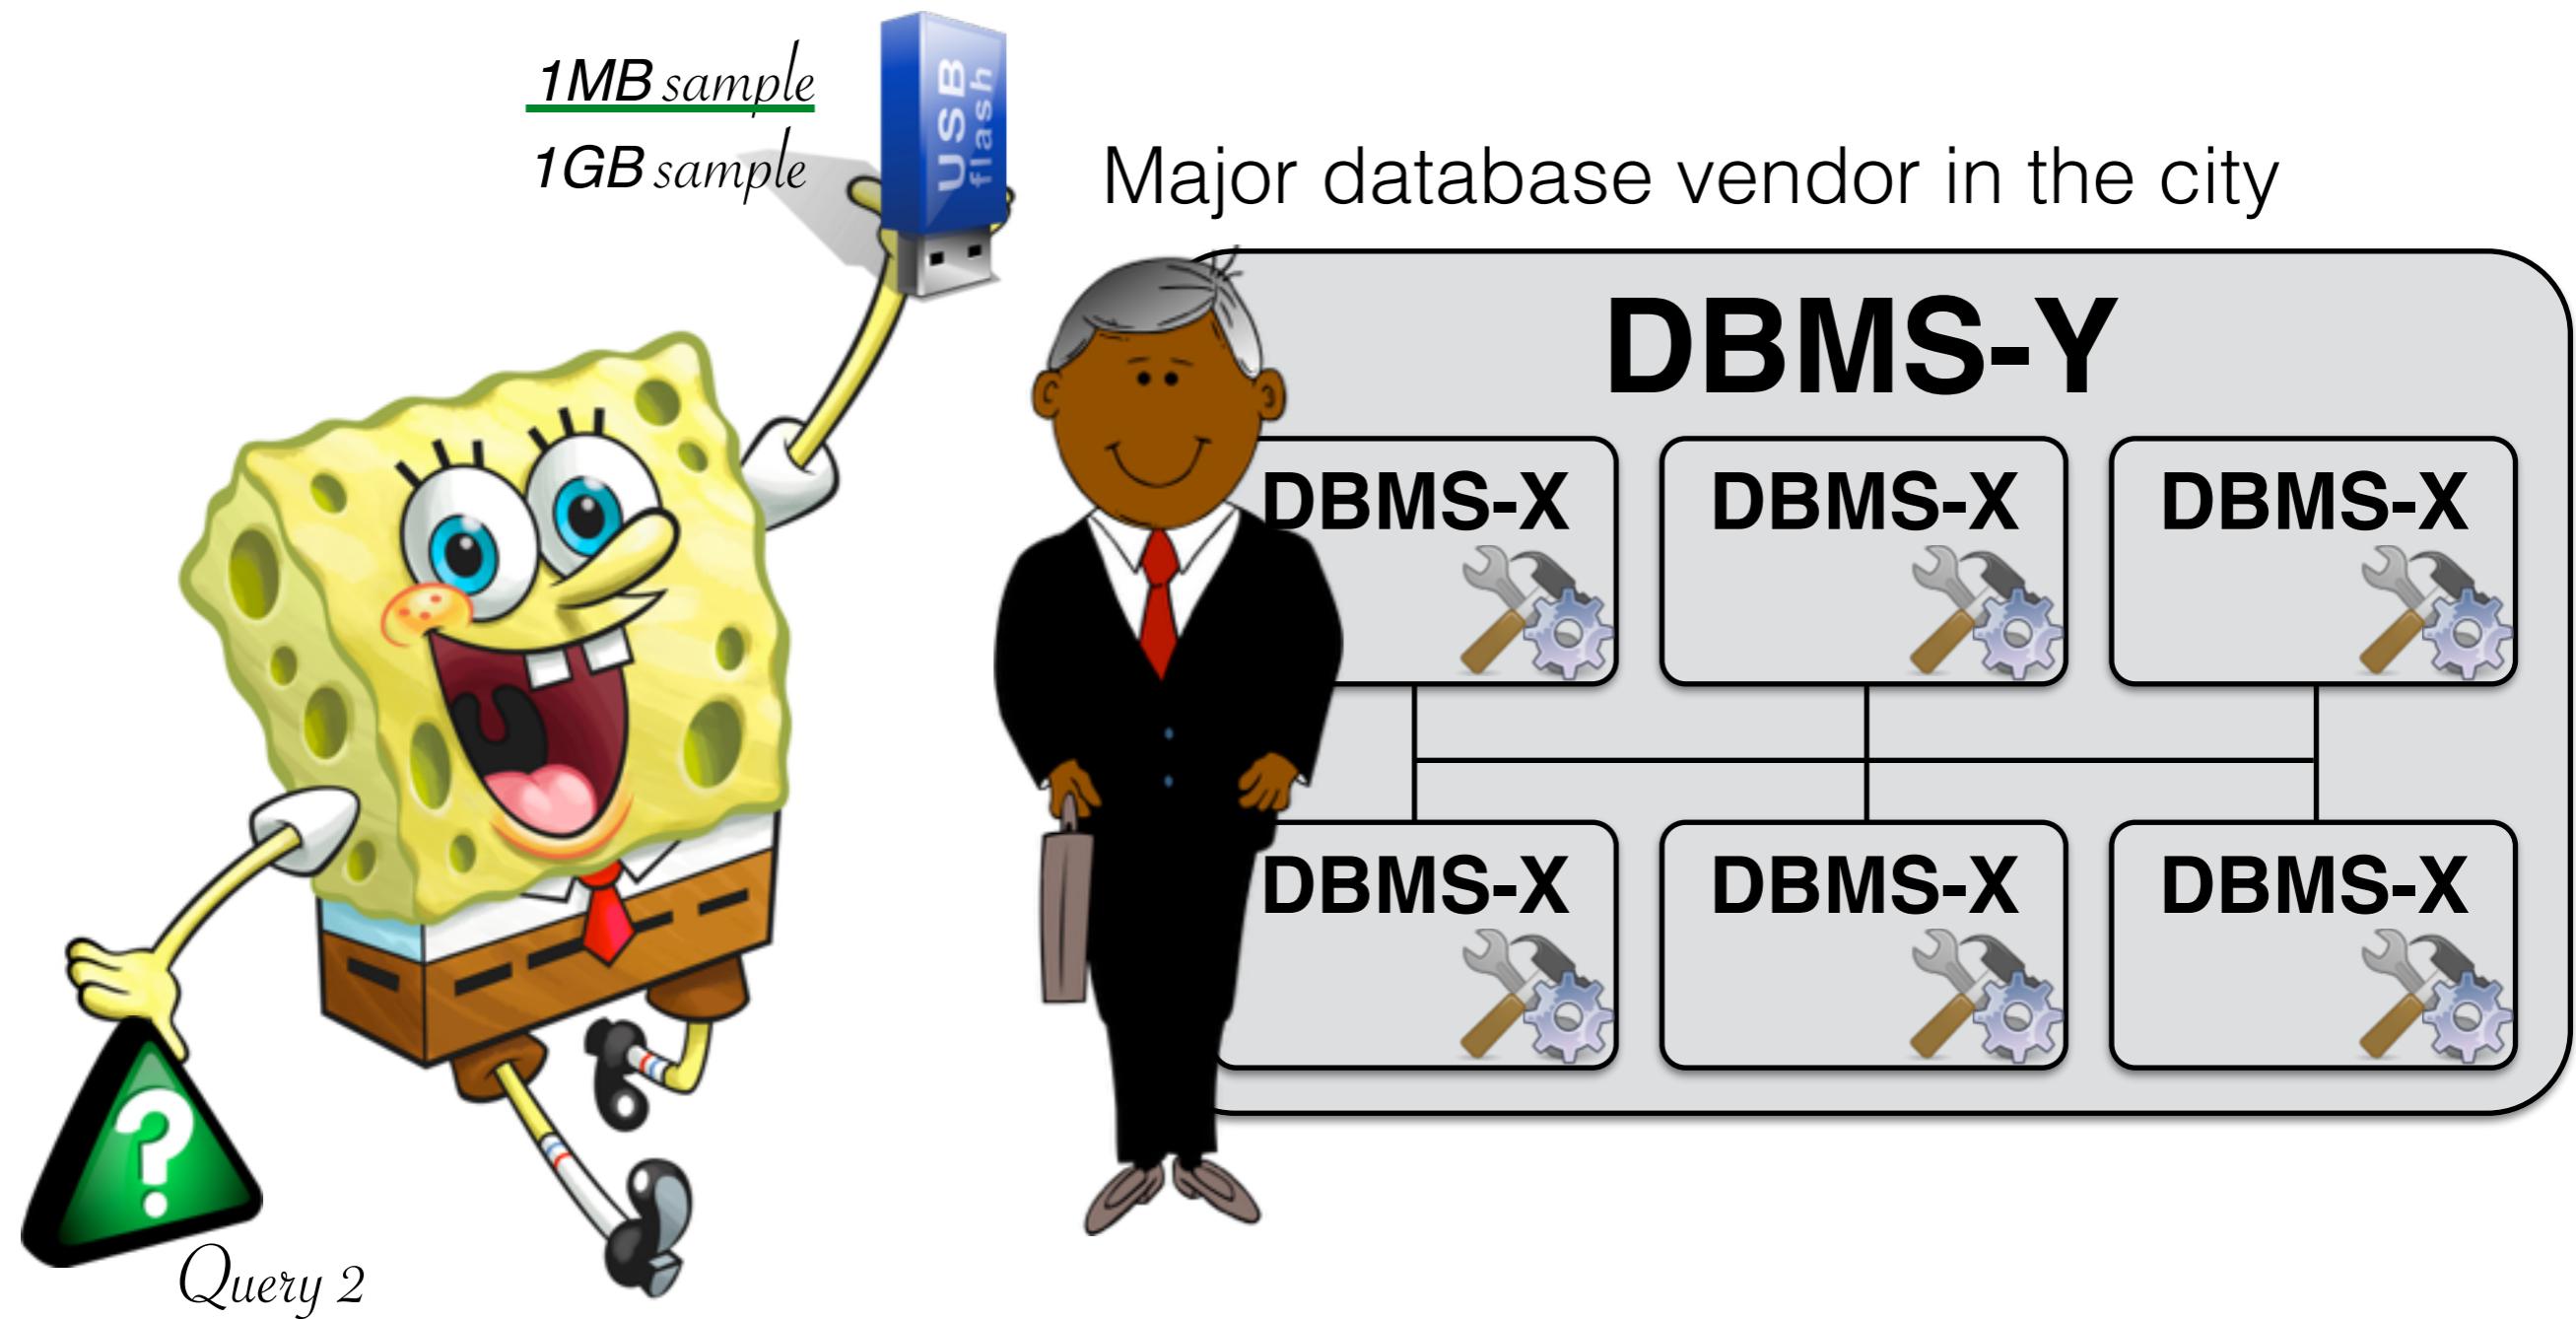
Major database vendor in the city



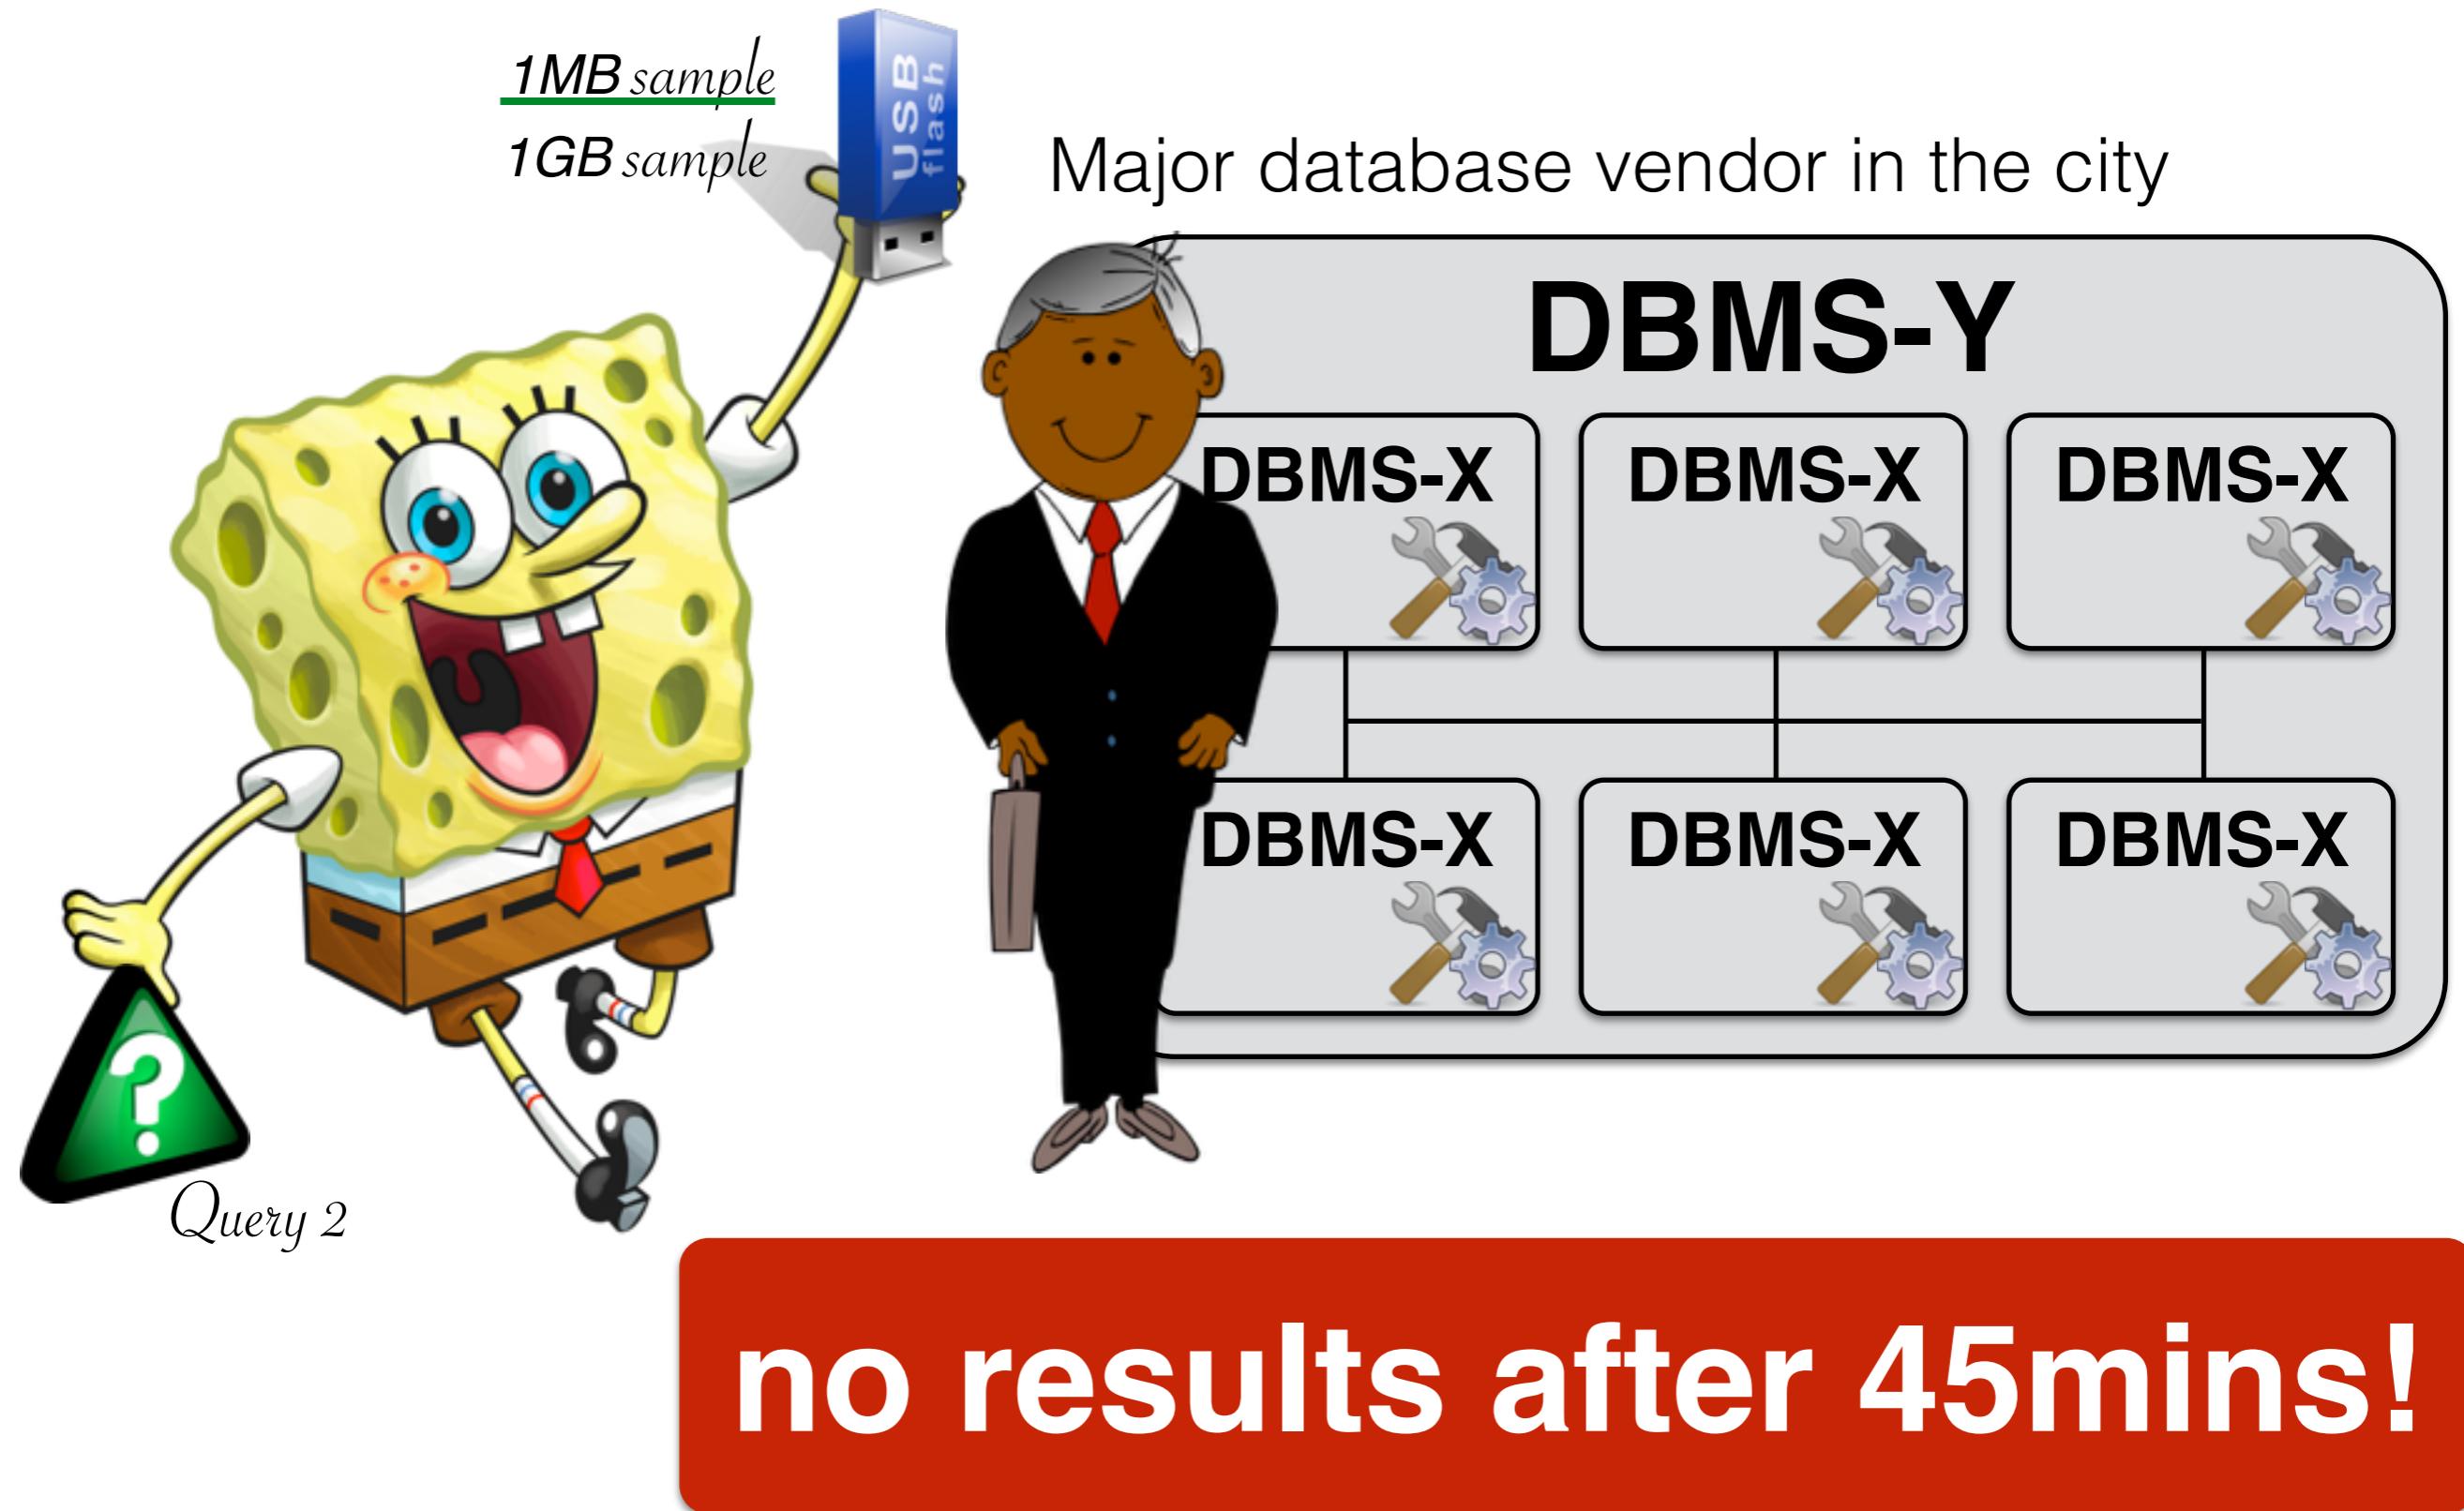
Bob@QCRI



Bob@QCRI



Bob@QCRI



Bob@QCRI

— cross product doesn't scale! —

Bob@QCRI

— cross product doesn't scale! —

1GB dataset x 1GB dataset

Bob@QCRI

— cross product doesn't scale! —

1GB dataset x 1GB dataset



100 trillion comparisons!
(0.0001ms per comparison)

Bob@QCRI

— cross product doesn't scale! —

1GB dataset x 1GB dataset



100 trillion comparisons!
(0.0001ms per comparison)



> 115 days!



> 300 thousands
years for 1TB!!!

Bob@QCRI

Can you handle
nonequi-joins?

DBMS-Z



Bob

Can you handle
nonequi-joins?

Cross product
+
selection

DBMS-Z



Bob@QCRI

Distributed frameworks:

- MapReduce
- Spark



Bob@QCRI

Distributed frameworks:

- MapReduce
- Spark



Bob@QCRI



**Distributed
frameworks:**

- MapReduce
- Spark

Joins:

- Sort-merge
- Band
- Interval

Bob@QCRI



**Distributed
frameworks:**

- MapReduce
- Spark

Joins:

- Sort-merge
- Range
- Interval

Bob@QCRI



Distributed frameworks:

- MapReduce
- Spark

Joins:

- Sort-merge
- Range
- Interval

But, can we avoid a cross product?

Bob@QCRI



Distributed frameworks:

- MapReduce
- Spark

Joins:

- Sort-merge
- Band
- Interval

Indexes:

- R-trees
- Space filling curves
- Bitmap
- Gist

But, can we avoid a cross product?

Bob@QCRI



Distributed frameworks:

- MapReduce
- Spark

Joins:

- Sort-merge
- Band
- Interval

Indexes:

- R-trees
- Space filling curves
- Bitmap
- Gist

But, can we avoid a cross product?

No existing technique
solves Bob's problem!

DataAnalyticsTeam@QCRI



DataAnalyticsTeam @ QCRI



More applications:

- Temporal DBs,
- Spatial DBs,
- Sensor networks,
- Social networks, ...



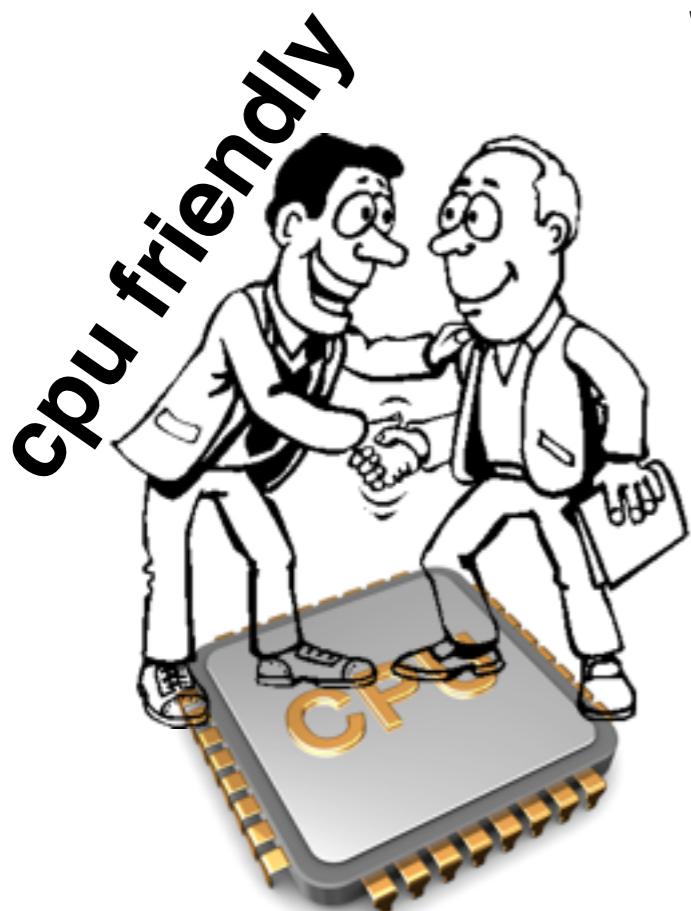
RAM locality is King!

— Jim Gray —

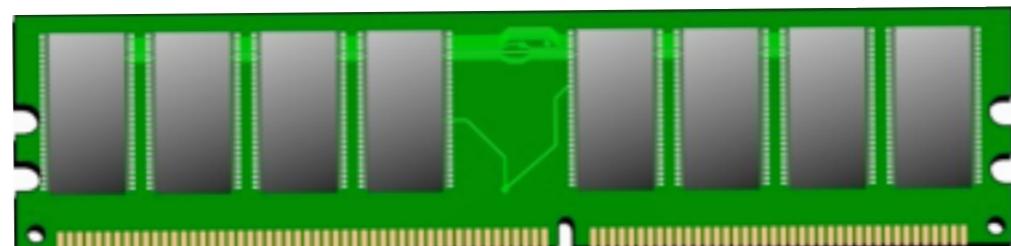


RAM locality is King!

— Jim Gray —

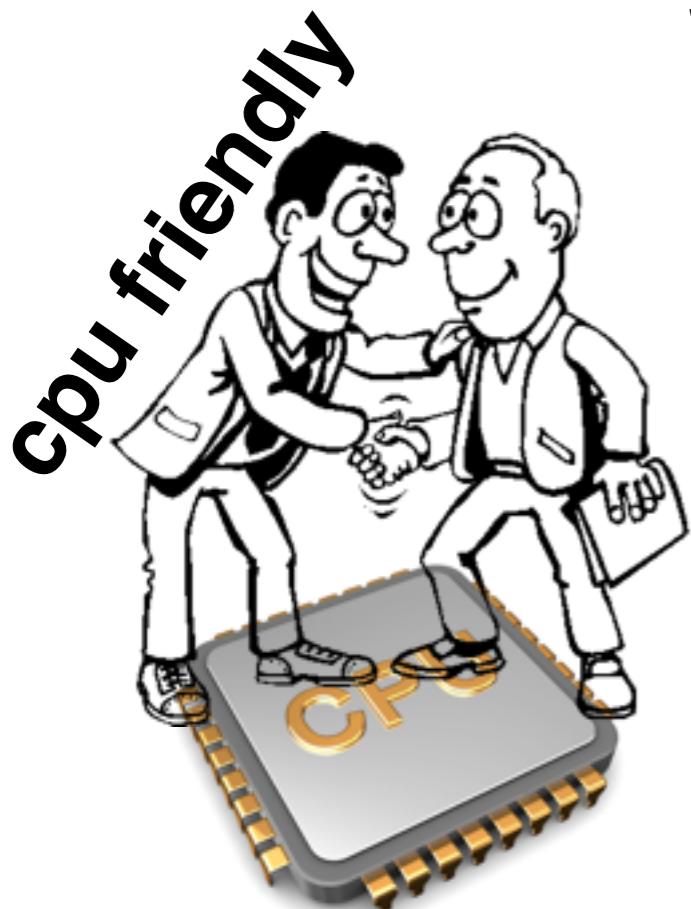


small memory footprint



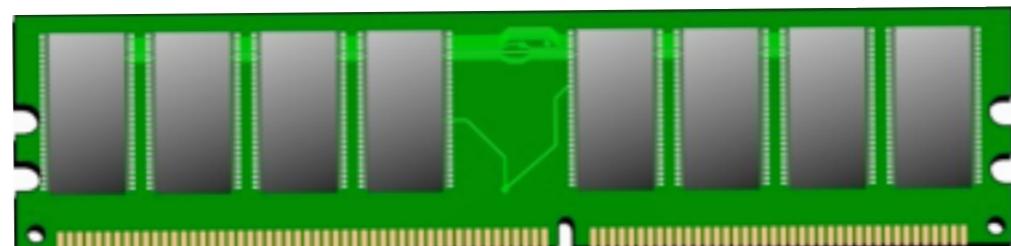
RAM locality is King!

— Jim Gray —



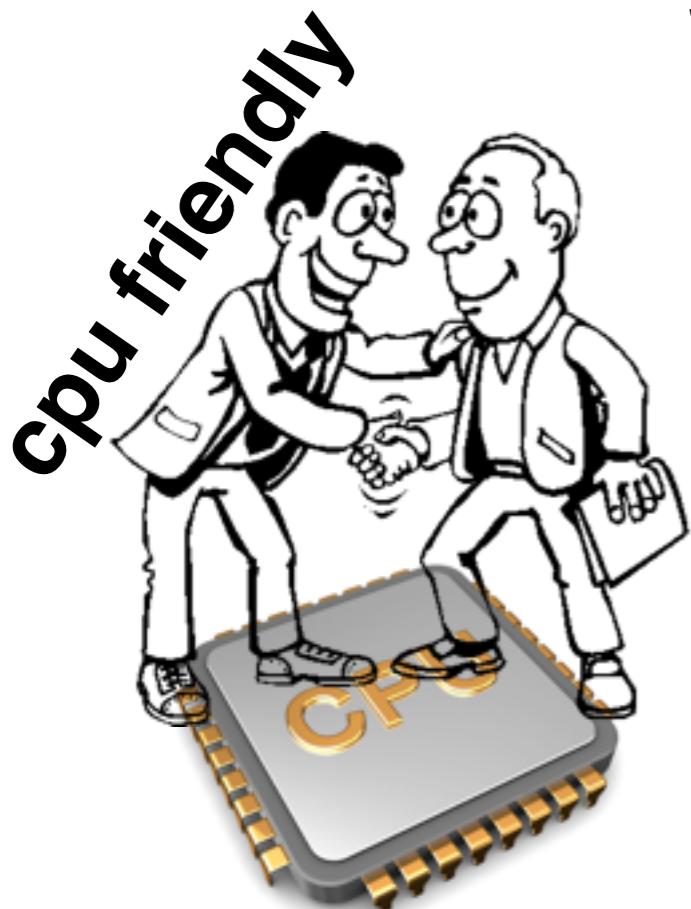
cpu friendly

small memory footprint



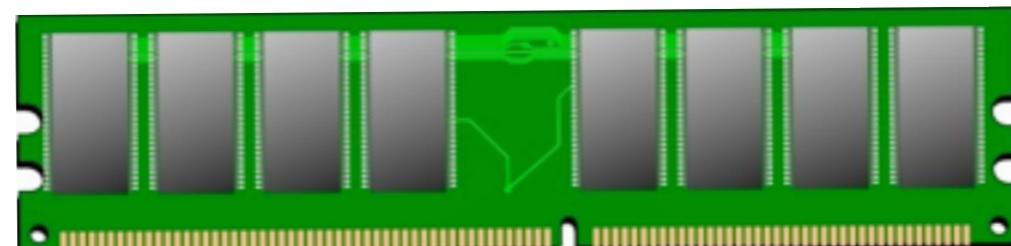
RAM locality is King!

— Jim Gray —



cpu friendly

small memory footprint



Fast!

Q2:

Data quality rule

```
SELECT e1.id, e2.id  
FROM employees e1, e2  
WHERE e1.salary > e2.salary  
AND e1.tax < e2.tax;
```

IJoin Algorithm

Employees

id	ssn	salary	tax
100	3456	100	10
101	9876	90	5
102	6790	160	14
103	4567	140	40
104	2095	150	15

Q2:

Data quality rule

```
SELECT e1.id, e2.id  
FROM employees e1, e2  
WHERE e1.salary > e2.salary
```

IJoin Algorithm

Employees

id	ssn	salary	tax
100	3456	100	10
101	9876	90	5
102	6790	160	14
103	4567	140	40
104	2095	150	15

Q2:

Data quality rule

SELECT e1.id, e2.id

FROM employees e1, e2

WHERE e1.salary > e2.salary

IEJoin Algorithm

Employees

id	ssn	salary	tax
100	3456	100	10
101	9876	90	5
102	6790	160	14
103	4567	140	40
104	2095	150	15

Sort on Salary

salary	tax	id
160	14	102
150	15	104
140	40	103
100	10	100
90	5	101



Q2:

Data quality rule

SELECT e1.id, e2.id

FROM employees e1, e2

WHERE e1.salary > e2.salary

AND e1.tax < e2.tax;

IEJoin Algorithm

Employees

id	ssn	salary	tax
100	3456	100	10
101	9876	90	5
102	6790	160	14
103	4567	140	40
104	2095	150	15

Sort on Salary

salary	tax	id
160	14	102
150	15	104
140	40	103
100	10	100
90	5	101



Q2:

Data quality rule

SELECT e1.id, e2.id

FROM employees e1, e2

WHERE e1.salary > e2.salary

AND e1.tax < e2.tax;

IEJoin Algorithm

Employees

id	ssn	salary	tax
100	3456	100	10
101	9876	90	5
102	6790	160	14
103	4567	140	40
104	2095	150	15

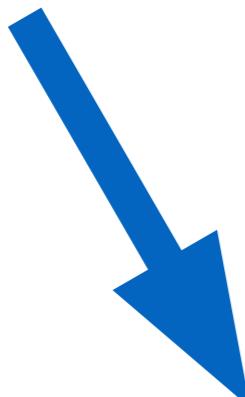
Sort on Salary

salary	tax	id
160	14	102
150	15	104
140	40	103
100	10	100
90	5	101



Sort on Tax

salary	tax	id
140	40	103
150	15	104
160	14	102
100	10	100
90	5	101



Q2:

Data quality rule

SELECT e1.id, e2.id

FROM employees e1, e2

WHERE e1.salary > e2.salary

AND e1.tax < e2.tax;

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

IEJoin Algorithm

Sort on Tax

id	salary	tax
103	140	40
104	150	15
102	160	14
100	100	10
101	90	5

Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

IEJoin Algorithm

Sort on Tax

id	salary	tax	order_sal
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

PA

Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

IEJoin Algorithm

Sort on Tax

id	salary	tax	order_sal
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array

0 | 0 | 0 | 0 | 0

Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

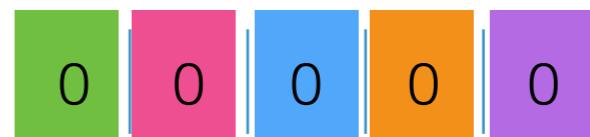
id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

IEJoin Algorithm

Sort on Tax

id	salary	tax	order_sal
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array



Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

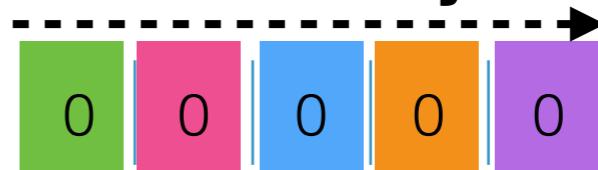
id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

IEJoin Algorithm

Sort on Tax

id	salary	tax	order_sal
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array



Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

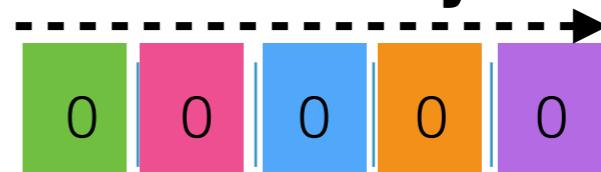
id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

IEJoin Algorithm

Sort on Tax

id	salary	tax	PA
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array



Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

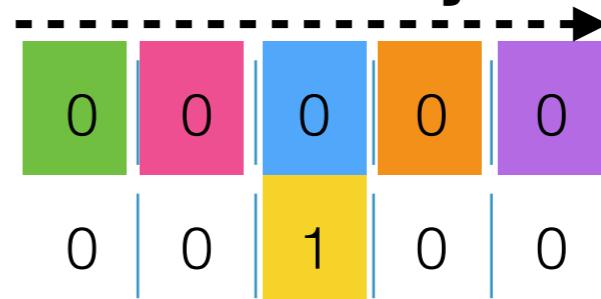
id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

IEJoin Algorithm

Sort on Tax

id	salary	tax	PA
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array



Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

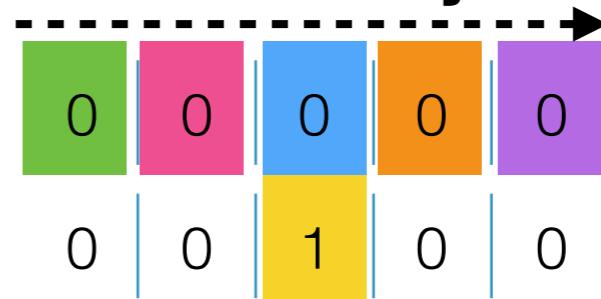
id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

IEJoin Algorithm

Sort on Tax

id	salary	tax	PA
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array



Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

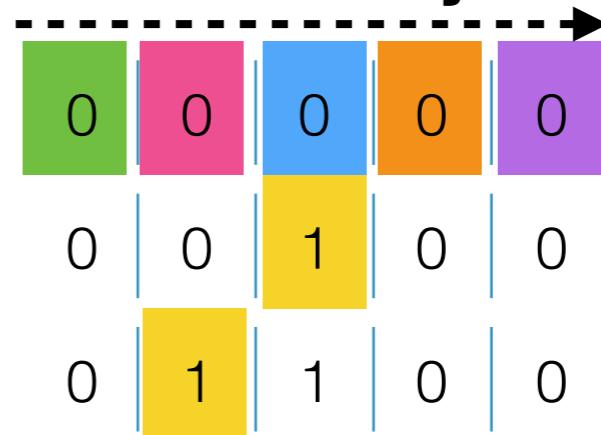
id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

IEJoin Algorithm

Sort on Tax

id	salary	tax	PA
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array



Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

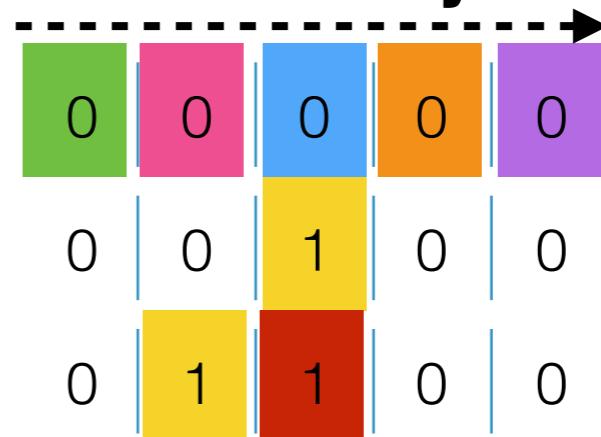
id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

IEJoin Algorithm

Sort on Tax

id	salary	tax	PA
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array



Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

Results

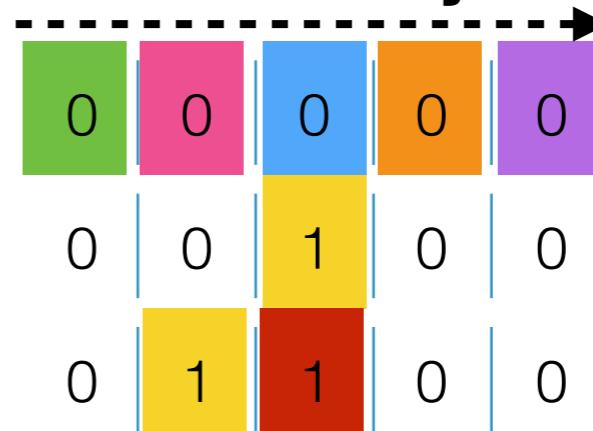
e1.id	e2.id
104	103

IEJoin Algorithm

Sort on Tax

id	salary	tax	PA
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array



Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

Results

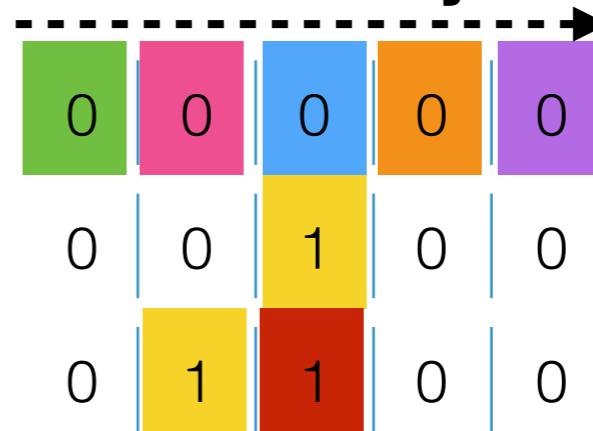
e1.id	e2.id
104	103

IEJoin Algorithm

Sort on Tax

id	salary	tax	PA	order_sal
103	140	40		2
104	150	15		1
102	160	14		0
100	100	10		3
101	90	5		4

Bit Array



Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

Results

e1.id	e2.id
104	103

IEJoin Algorithm

Sort on Tax

id	salary	tax	PA
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array

0	0	0	0	0
0	0	1	0	0
0	1	1	0	0
1	1	1	0	0

Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

Results

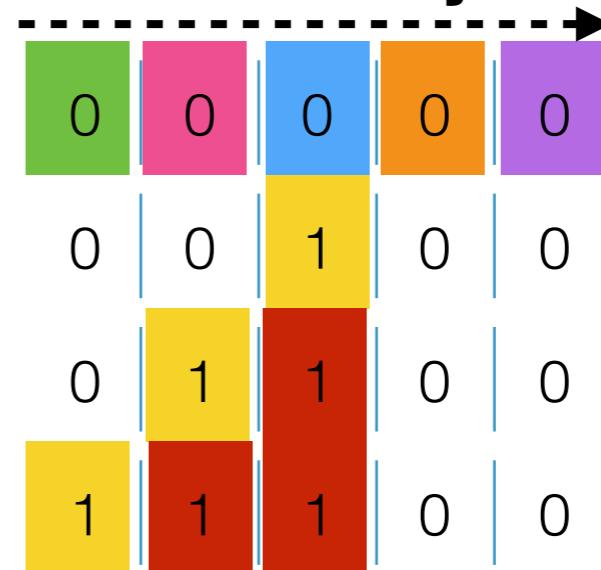
e1.id	e2.id
104	103
102	104
102	103

IEJoin Algorithm

Sort on Tax

id	salary	tax	order_sal
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array



Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

Results

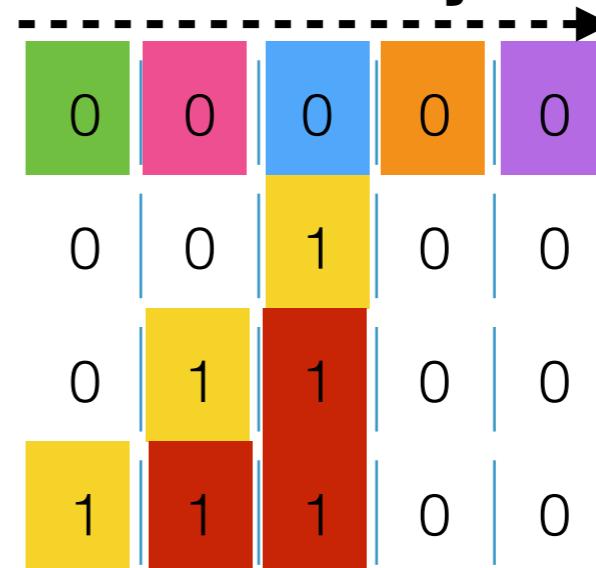
e1.id	e2.id
104	103
102	104
102	103

IEJoin Algorithm

Sort on Tax

id	salary	tax	order_sal
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array



Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

Results

e1.id	e2.id
104	103
102	104
102	103

IEJoin Algorithm

Sort on Tax

id	salary	tax	order_sal
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array

0	0	1	0	0
0	1	1	0	0
1	1	1	0	0
1	1	1	1	0

Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

Results

e1.id	e2.id
104	103
102	104
102	103

IEJoin Algorithm

Sort on Tax

id	salary	tax
103	140	40
104	150	15
102	160	14
100	100	10
101	90	5

PA

order_sal
2
1
0
3
4

Bit Array

0	0	1	0	0
0	1	1	0	0
1	1	1	0	0
1	1	1	1	0

Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

Results

e1.id	e2.id
104	103
102	104
102	103

IEJoin Algorithm

Sort on Tax

id	salary	tax
103	140	40
104	150	15
102	160	14
100	100	10
101	90	5

PA

order_sal
2
1
0
3
4

Bit Array

0	0	1	0	0
0	1	1	0	0
1	1	1	0	0
1	1	1	1	0
1	1	1	1	1

Q2:

Data quality rule

```
SELECT e1.id, e2.id
FROM employees e1, e2
WHERE e1.salary > e2.salary
      AND e1.tax < e2.tax;
```

Sort on Salary

id	salary	tax
102	160	14
104	150	15
103	140	40
100	100	10
101	90	5

Results

e1.id	e2.id
104	103
102	104
102	103

IEJoin Algorithm

Sort on Tax

id	salary	tax	order_sal
103	140	40	2
104	150	15	1
102	160	14	0
100	100	10	3
101	90	5	4

Bit Array

0	0	1	0	0
0	1	1	0	0
1	1	1	0	0
1	1	1	1	0
1	1	1	1	1

Voila!



Enhancing IEJoin Algorithm

Emp1

ssn	salary	tax
3456	100	10
9876	90	5
6790	160	14
4567	140	40
2095	150	15

Emp2

ssn	salary	tax
3456	100	10
9876	90	5
6790	160	14
4567	140	40
2095	150	15

Enhancing IEJoin Algorithm

Emp1

ssn	salary	tax
3456	100	10
9876	90	5
6790	160	14
4567	140	40
2095	150	15

Emp1: Sort on Salary

salary	tax
160	14
.	.
.	.
2	.
.	.

Emp1: Sort on Tax

salary	tax
140	40
.	.
.	.
2	.
.	.

Emp2

ssn	salary	tax
3456	100	10
9876	90	5
6790	160	14
4567	140	40
2095	150	15

Emp2: Sort on Salary

salary	tax
210	42
.	.
.	.
2	.
.	.

Emp2: Sort on Tax

salary	tax
99	23
.	.
.	.
2	.
.	.

Enhancing IEJoin Algorithm

Emp1

ssn	salary	tax
3456	100	10
9876	90	5
6790	160	14
4567	140	40
2095	150	15

Emp1: Sort on Salary

salary	tax
160	14
·	·
·	·
2	3
·	·

Emp1: Sort on Tax

salary	tax
140	40
·	·
·	·
5	·
·	·

Emp2

ssn	salary	tax
3456	100	10
9876	90	5
6790	160	14
4567	140	40
2095	150	15

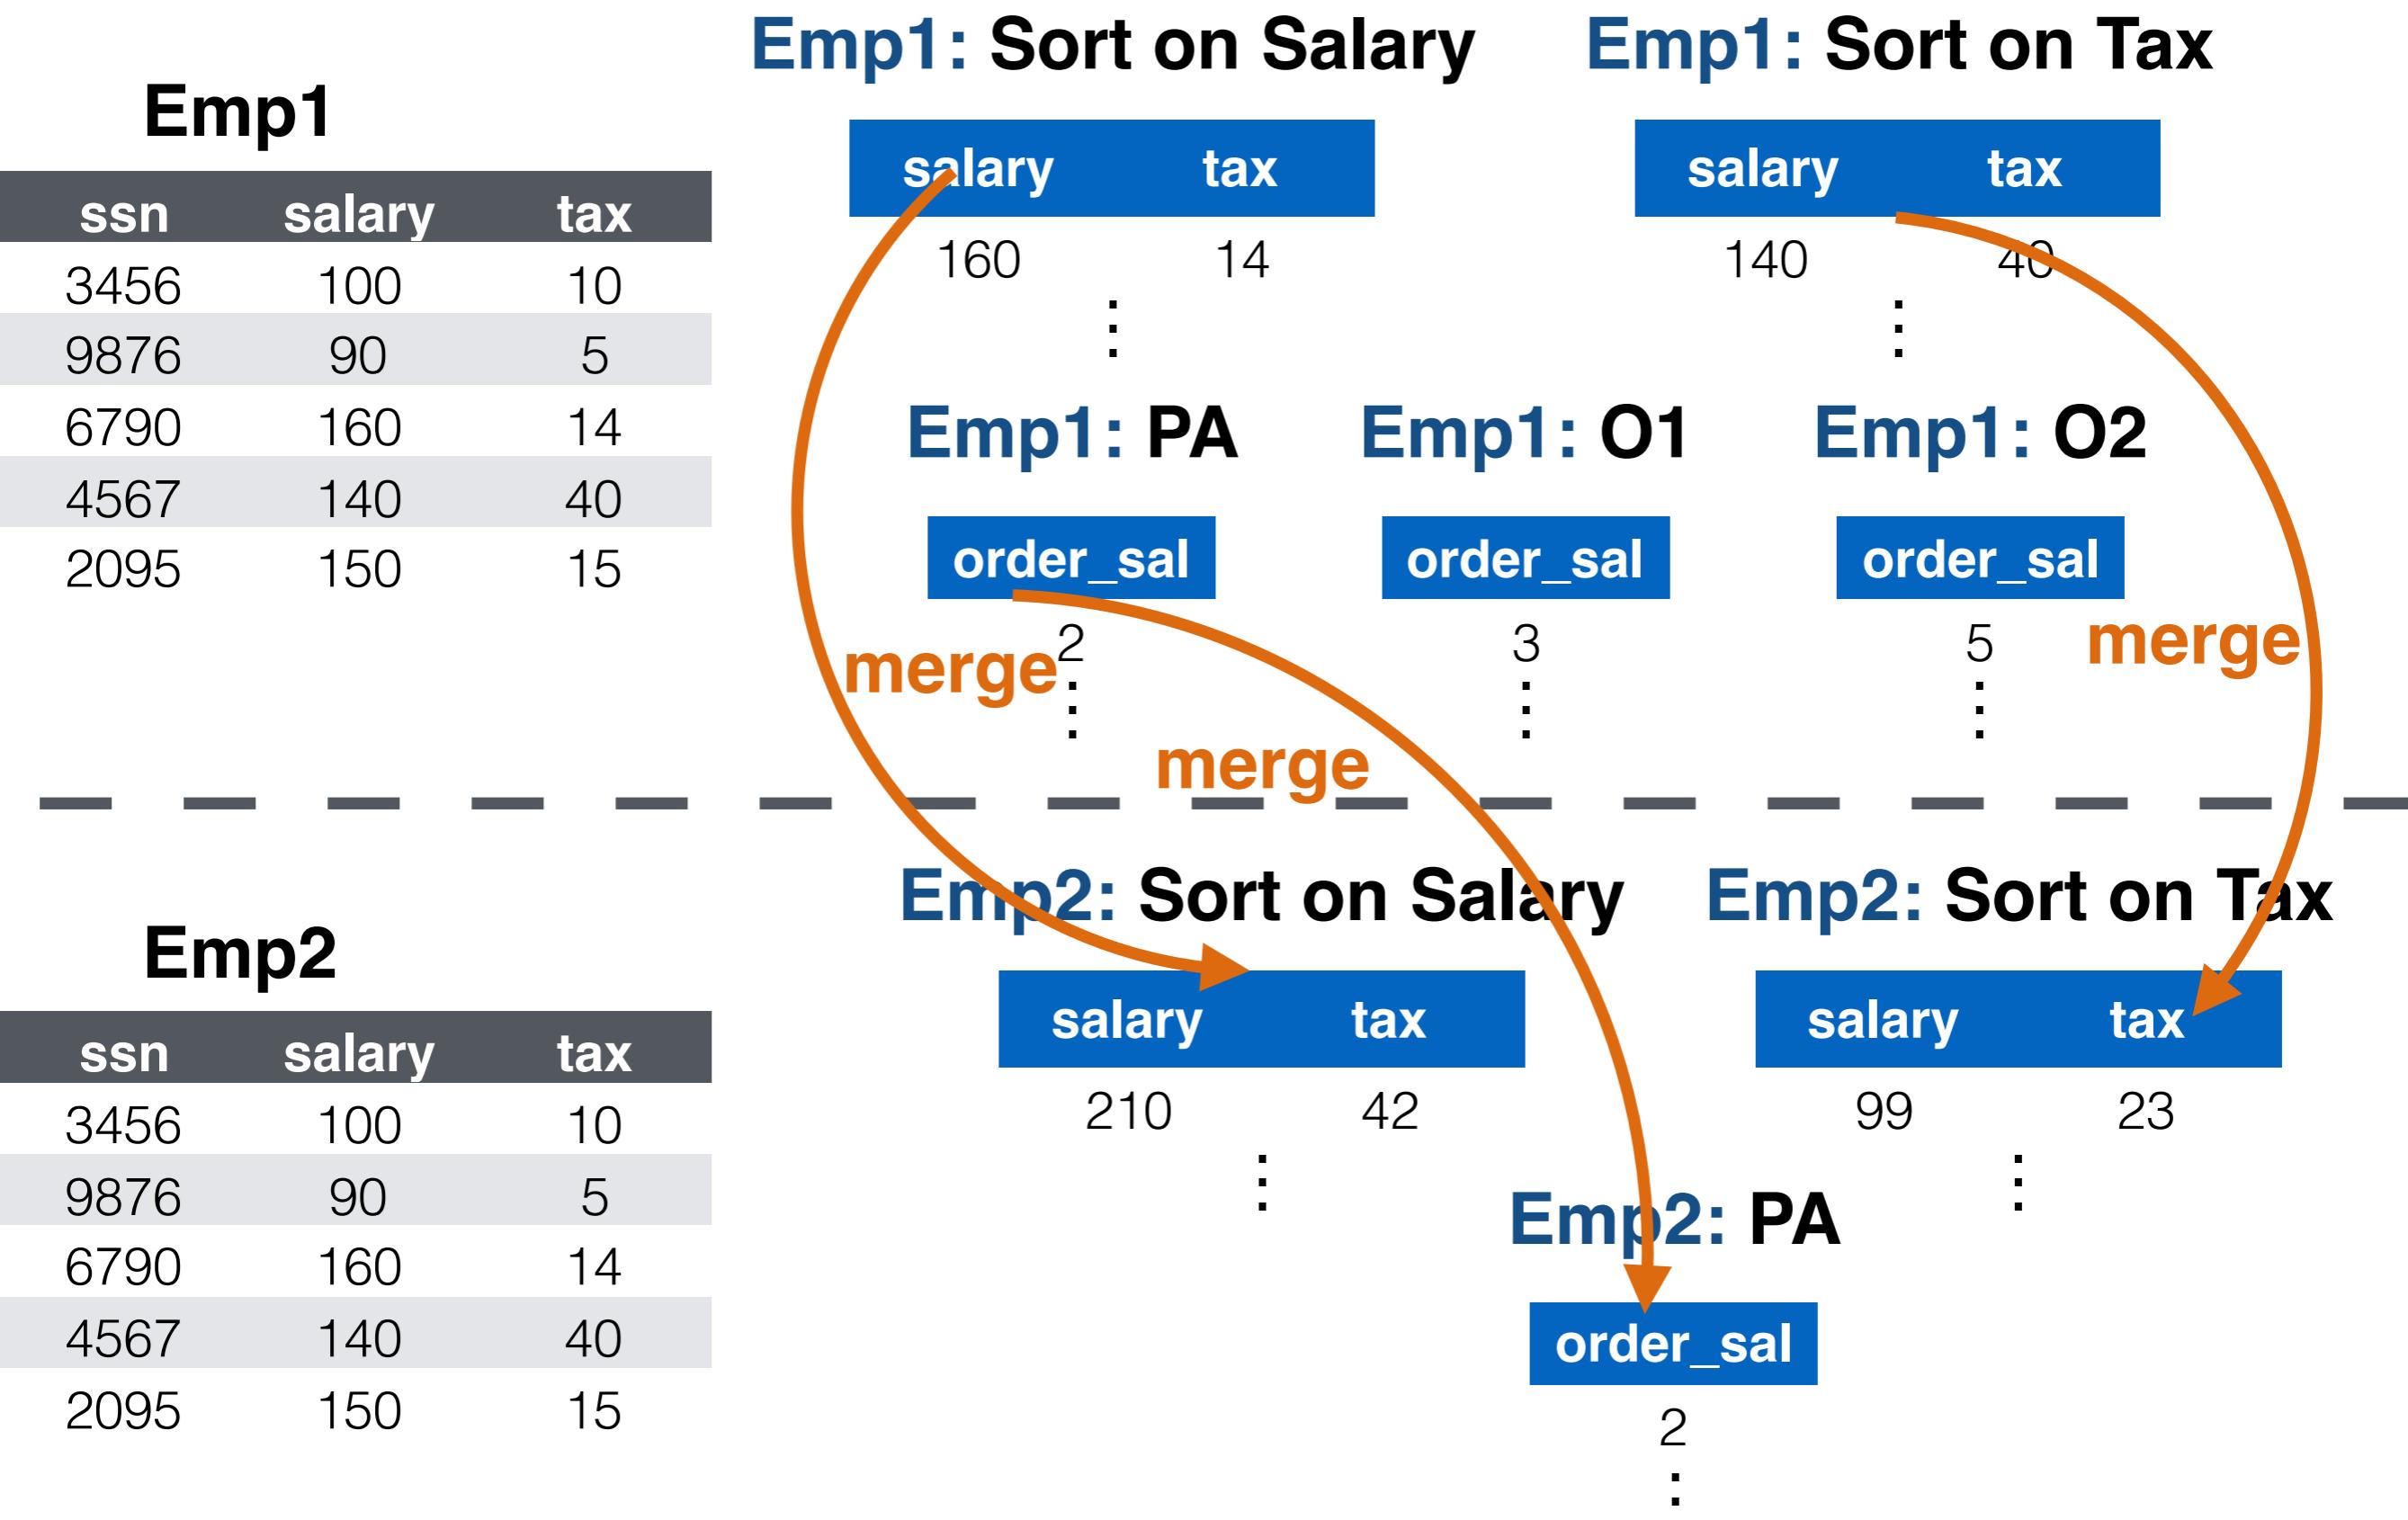
Emp2: Sort on Salary

salary	tax
210	42
·	·
·	·
2	·
·	·

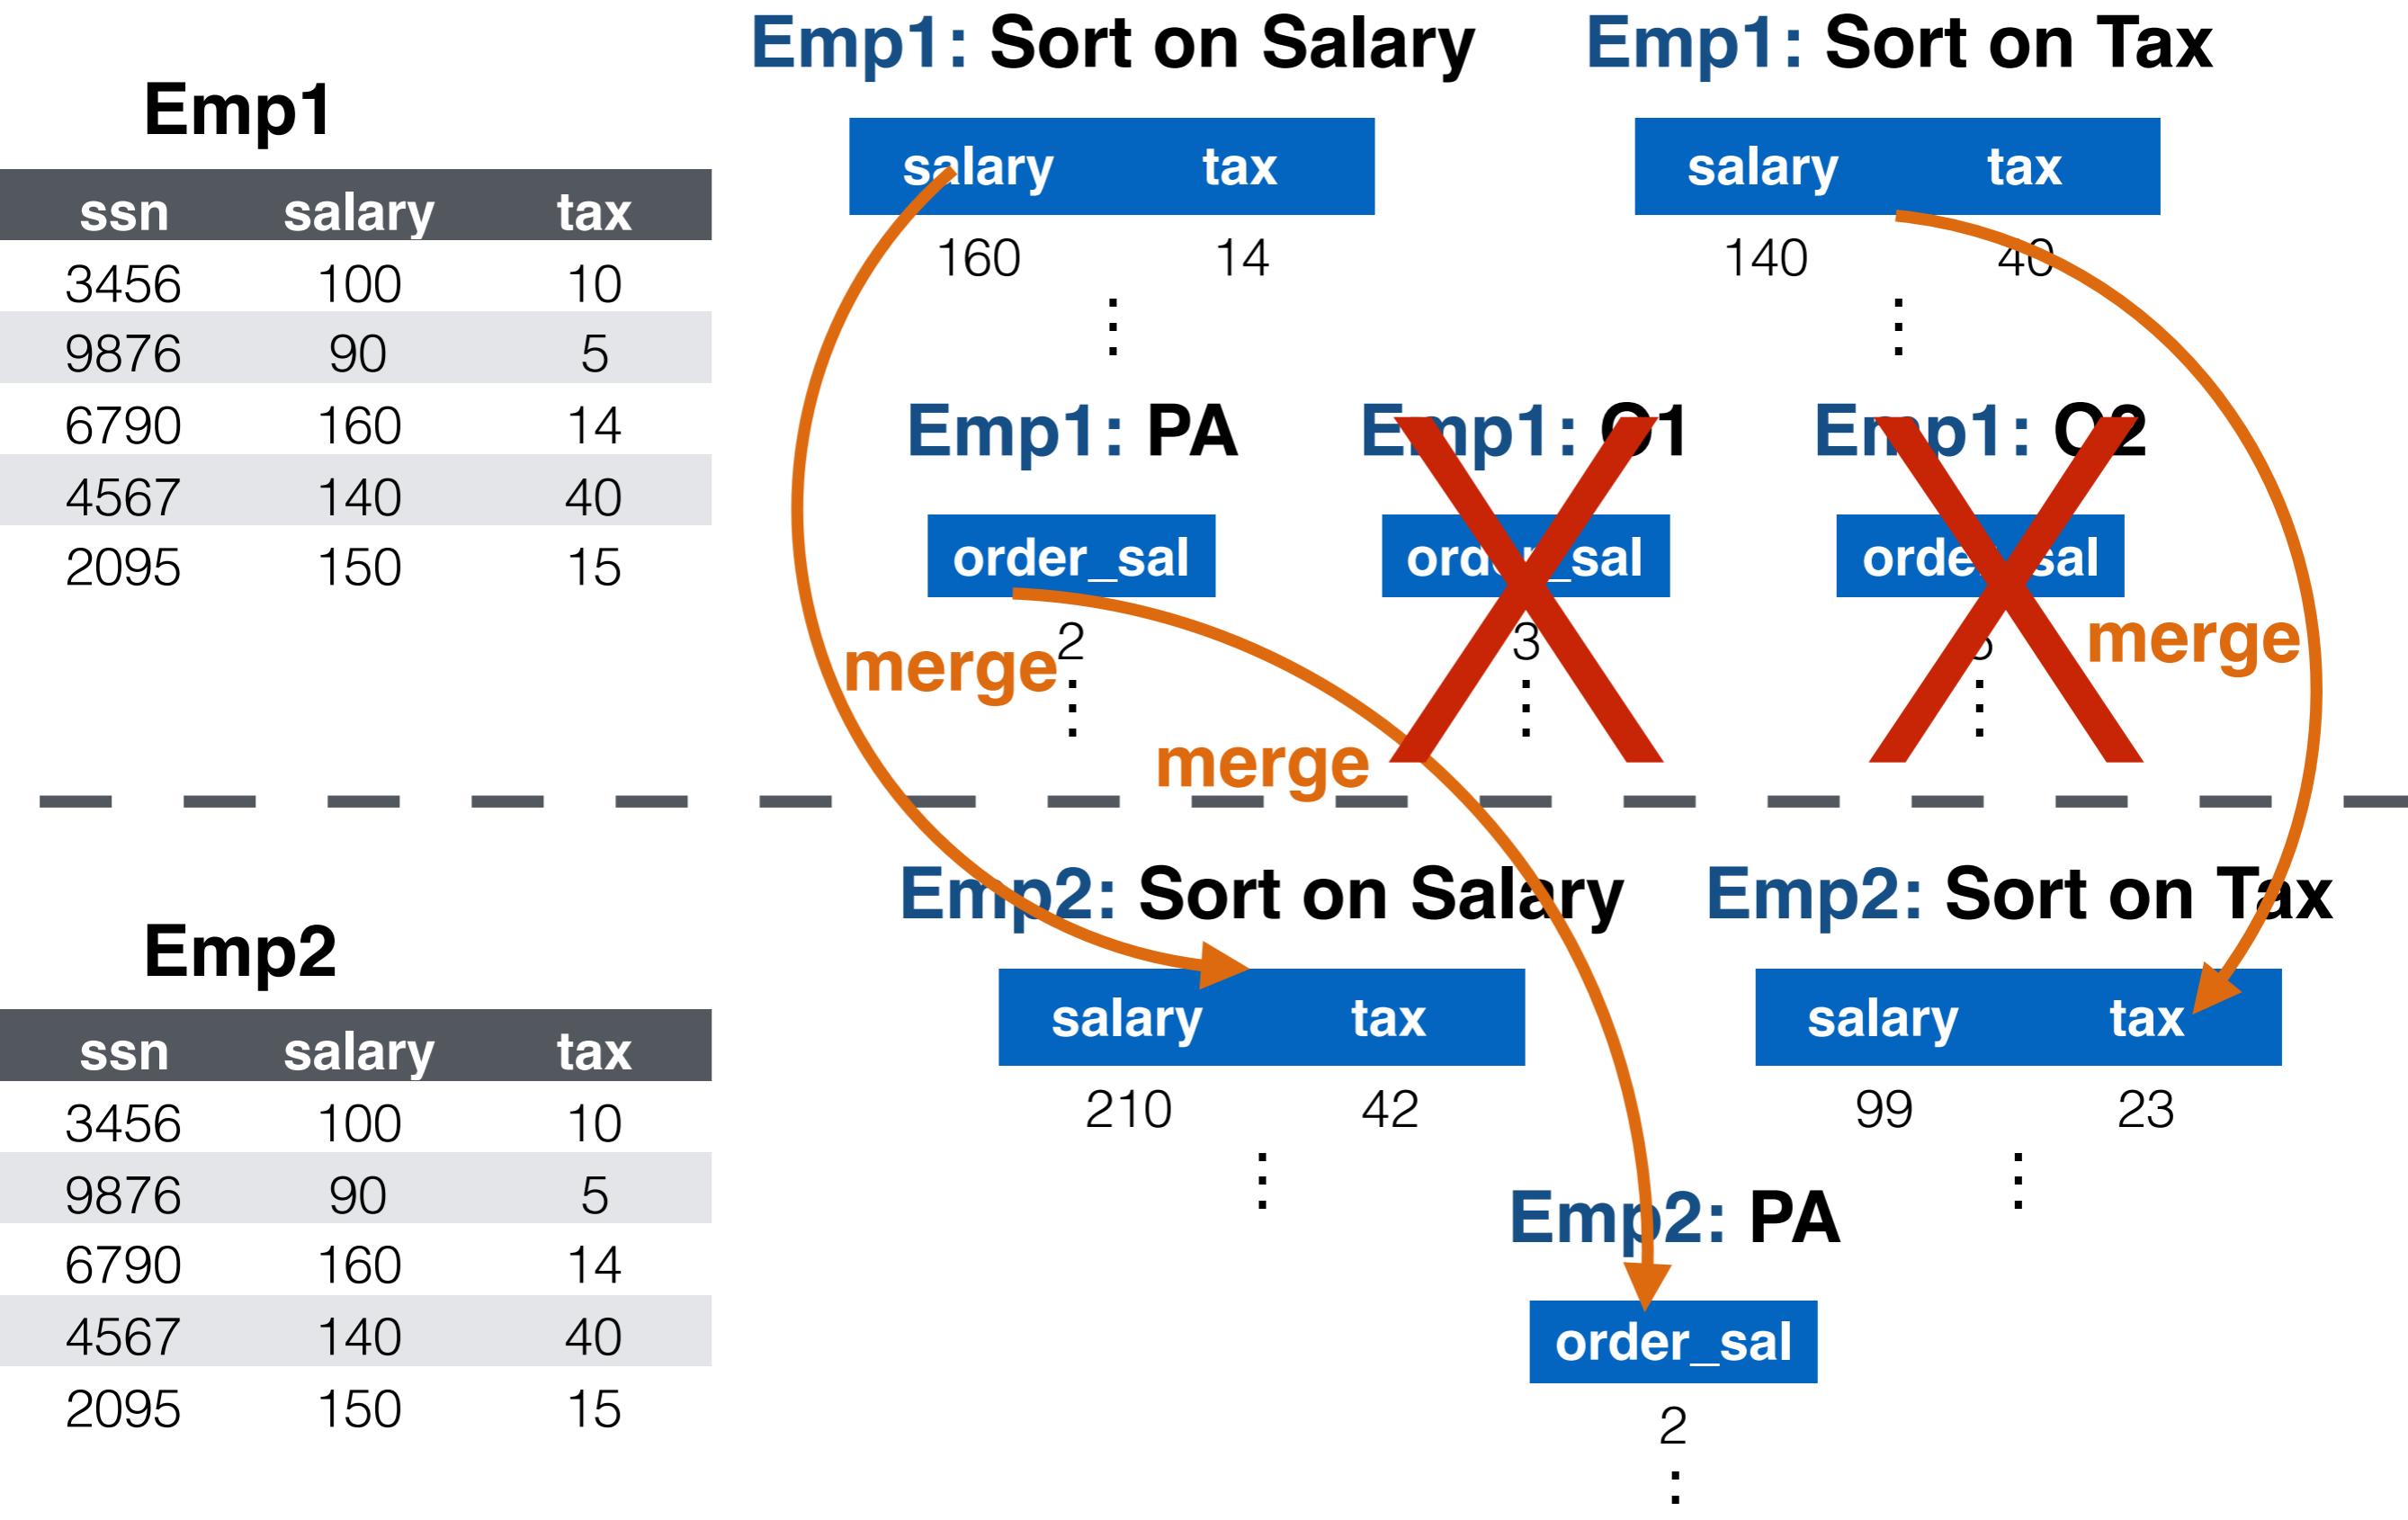
Emp2: Sort on Tax

salary	tax
99	23
·	·
·	·
order_sal	·
2	·
·	·

Enhancing IEJoin Algorithm



Enhancing IEJoin Algorithm



Enhancing IEJoin Algorithm

Sort on Salary

salary	tax
--------	-----

210 42

:

Sort on Tax

salary	tax
--------	-----

99 23

:

PA

order_sal

2
:

Bit Array

0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0

Enhancing IEJoin Algorithm

Sort on Salary

salary	tax
--------	-----

210 42

:

Sort on Tax

salary	tax
--------	-----

99 23

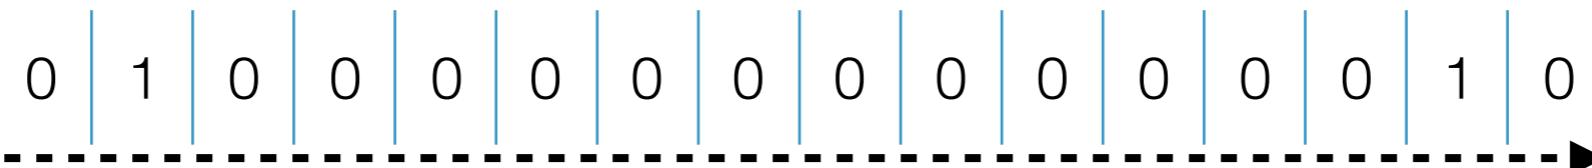
:

PA

order_sal

2
:
:

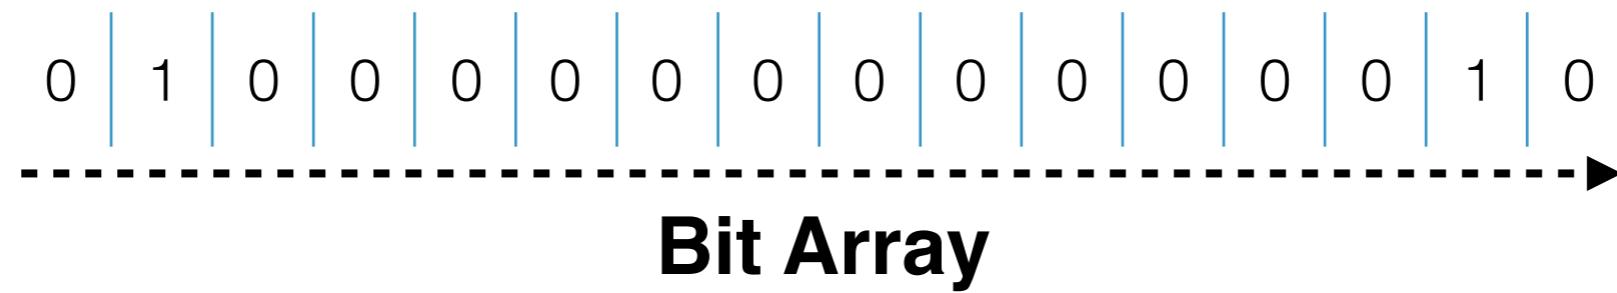
Bit Array



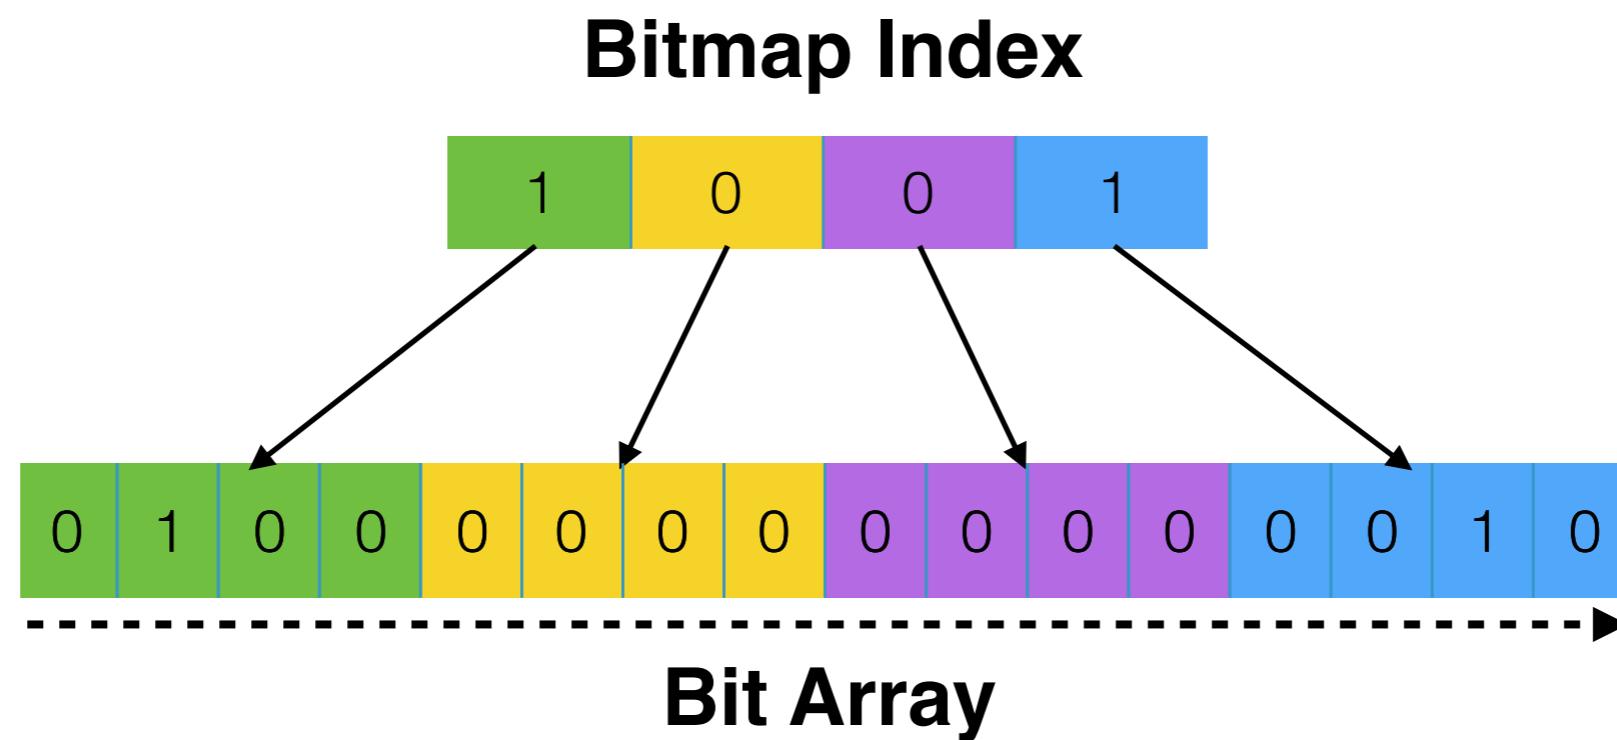
+

Postfiltering

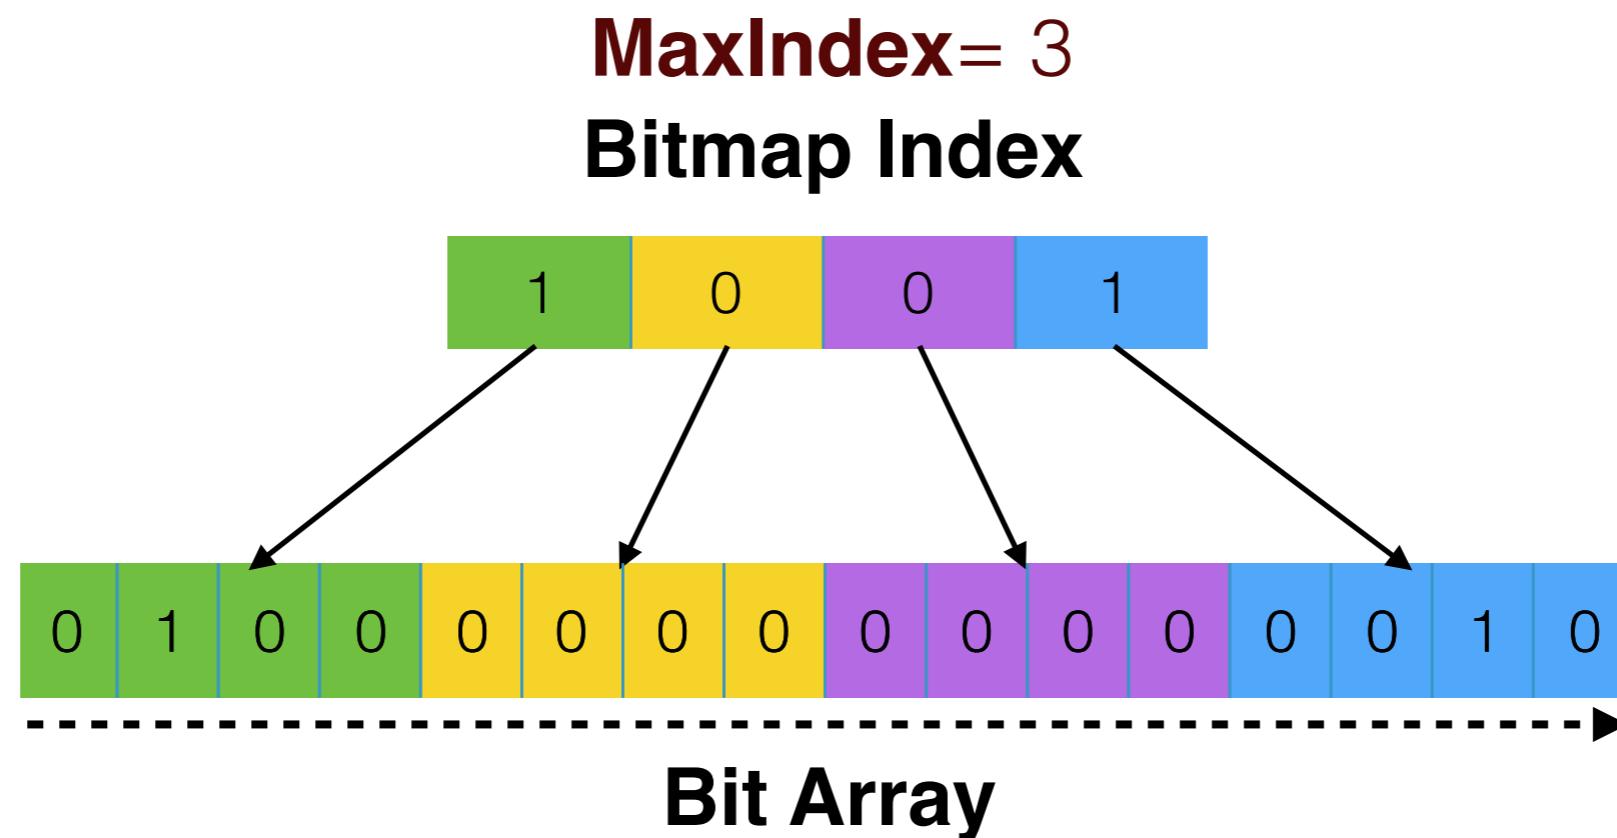
Enhancing IEJoin Algorithm



Enhancing IEJoin Algorithm



Enhancing IEJoin Algorithm



Scaling IEJoin Up

R

S

Phase 1: Pre-processing

Phase 2: IEJoin

Phase 3: Post-processing

Scaling IEJoin Up

R

S

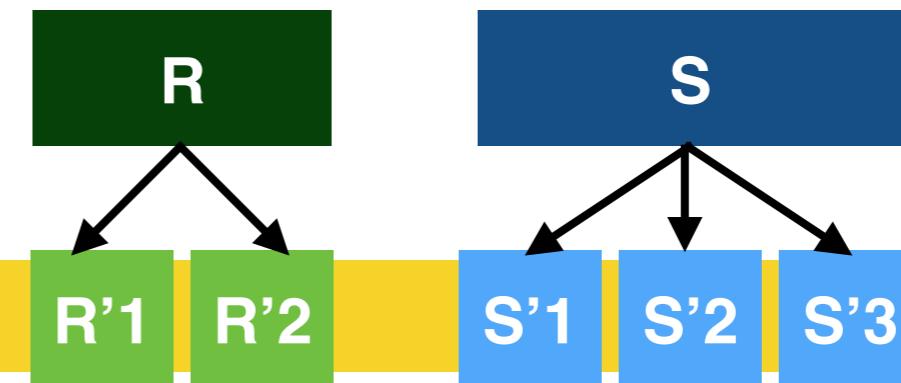
Phase 1: Pre-processing

Phase 2: IEJoin

Phase 3: Post-processing

sort on 1 att. + projection + partition

Scaling IEJoin Up



Phase 1: Pre-processing

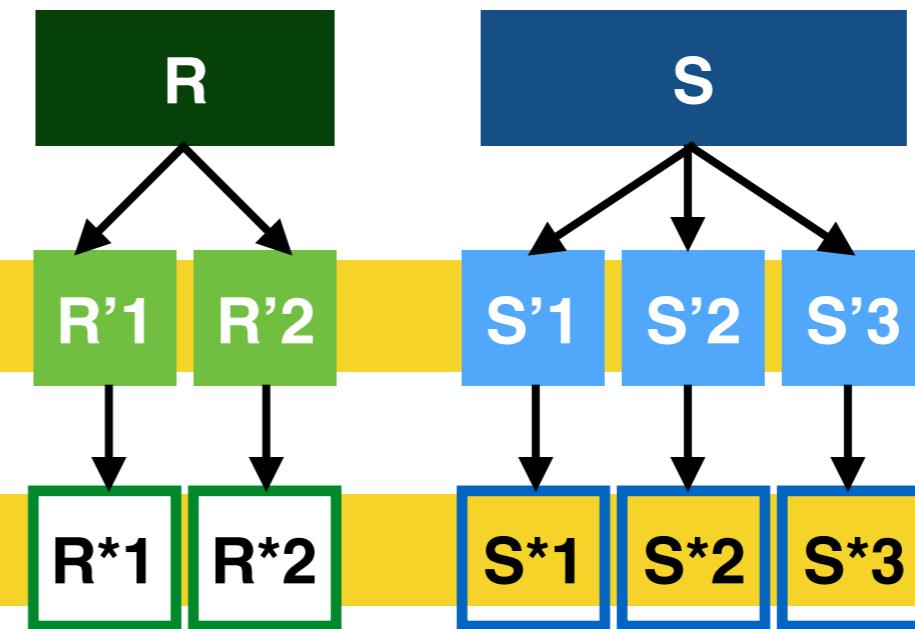
Phase 2: IEJoin

Phase 3: Post-processing

sort on 1 att. + projection + partition

metadata (min-max values)

Scaling IEJoin Up



Scaling IEJoin Up

Phase 1: Pre-processing

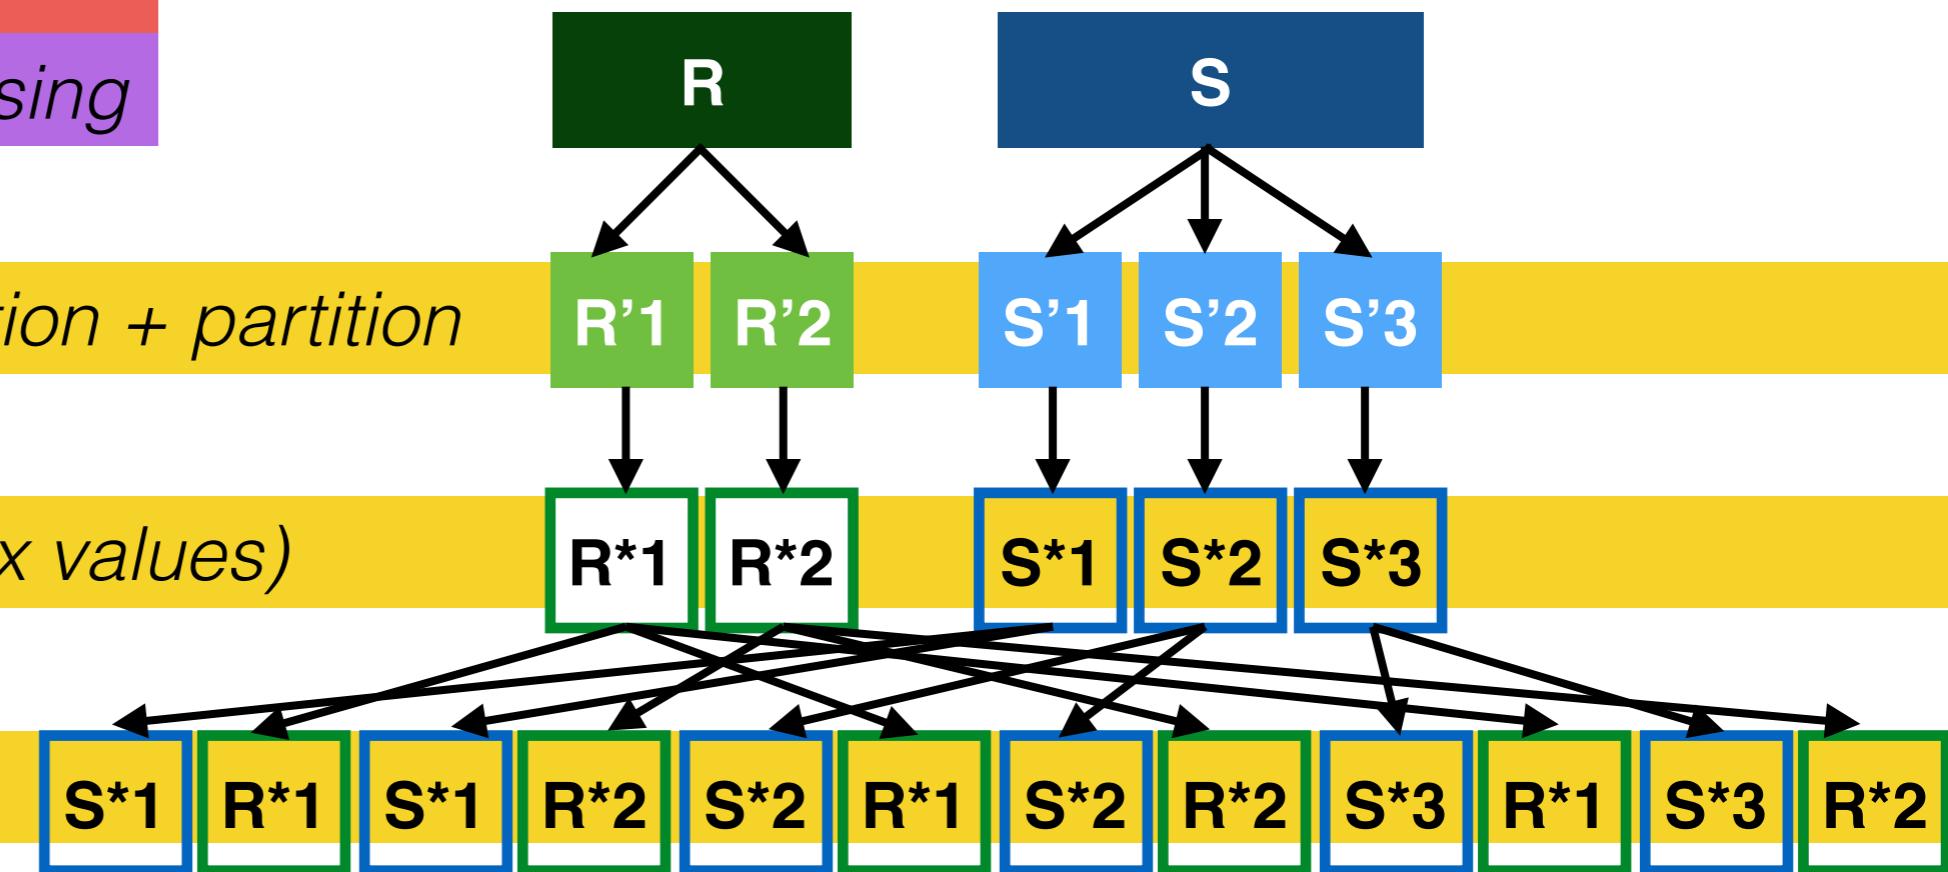
Phase 2: IEJoin

Phase 3: Post-processing

sort on 1 att. + projection + partition

metadata (min-max values)

cross product



Scaling IEJoin Up

Phase 1: Pre-processing

Phase 2: IEJoin

Phase 3: Post-processing

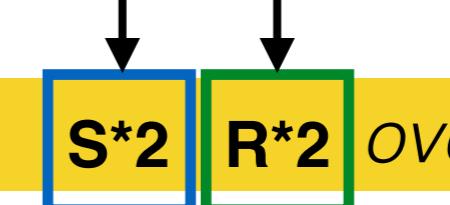
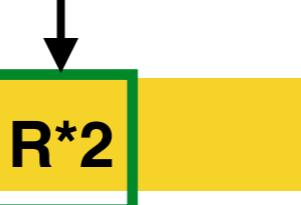
sort on 1 att. + projection + partition

metadata (min-max values)

cross product

S

R



overlapping blocks



Scaling IEJoin Up

Phase 1: Pre-processing

Phase 2: IEJoin

Phase 3: Post-processing

sort on 1 att. + projection + partition



metadata (min-max values)



cross product



overlapping blocks



recover projection

Scaling IEJoin Up

Phase 1: Pre-processing

Phase 2: IEJoin

Phase 3: Post-processing

sort on 1 att. + projection + partition



metadata (min-max values)



cross product



overlapping blocks

recover projection
perform IEJoin

Scaling IEJoin Up

Phase 1: Pre-processing

Phase 2: IEJoin

Phase 3: Post-processing

sort on 1 att. + projection + partition

metadata (min-max values)

cross product



More Optimizations...

Query Optimization & Incremental IJoin

— Fast and Scalable Inequality Joins —
VLDBJ, special issue on best VLDB'16 papers



Implemented on...

IJoin



PostgreSQL



Compared with...

BTree

GiST

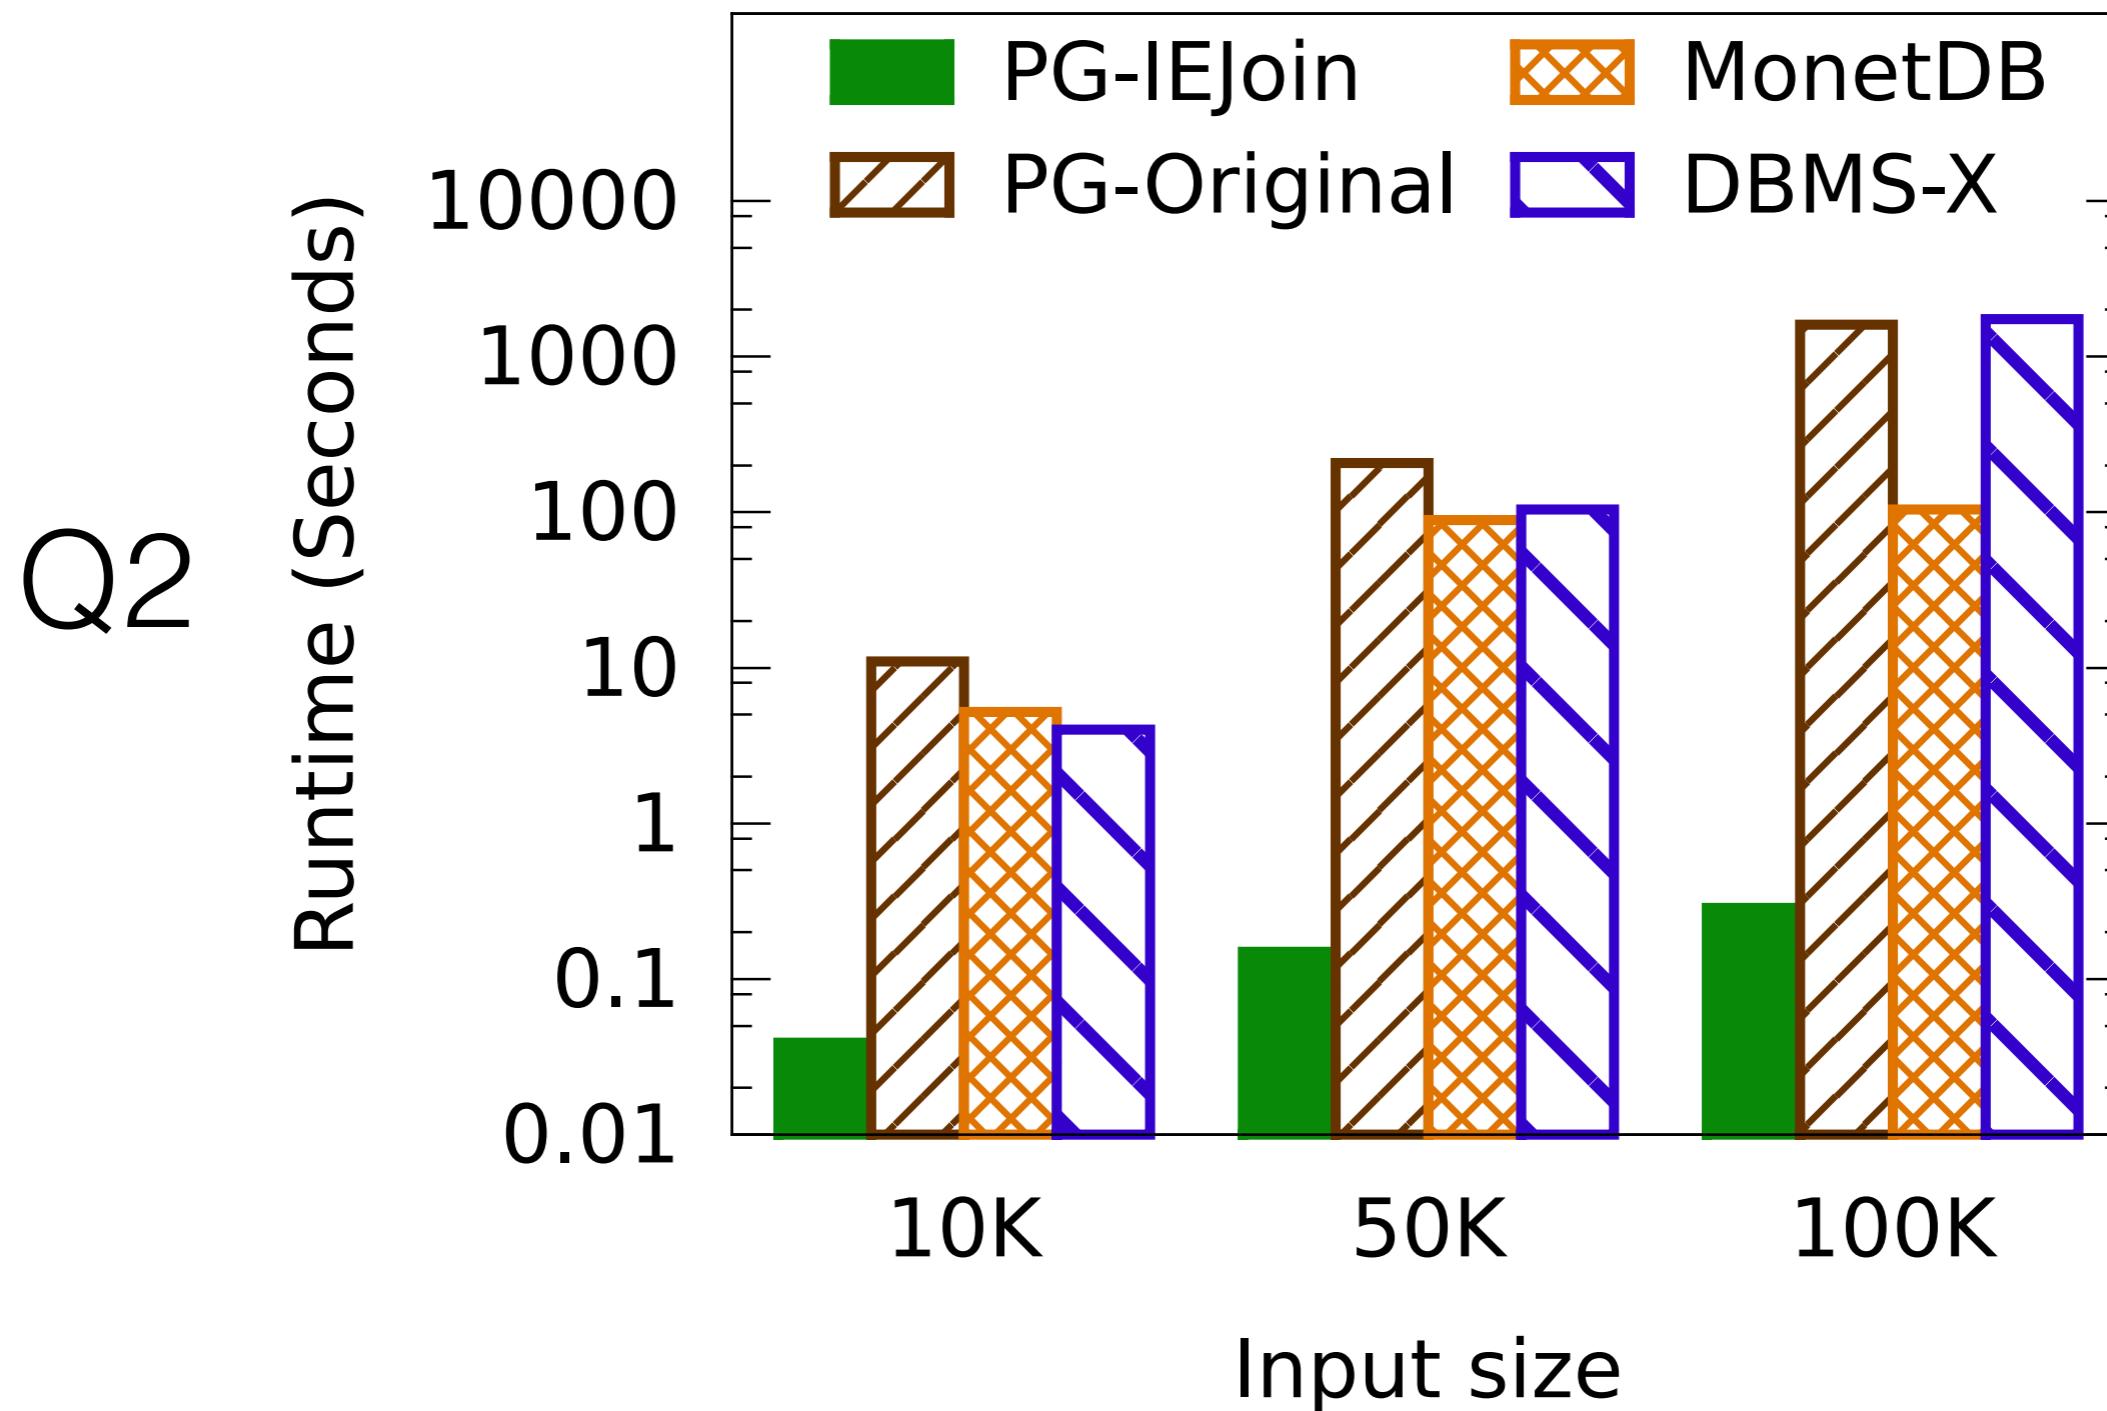


PostgreSQL

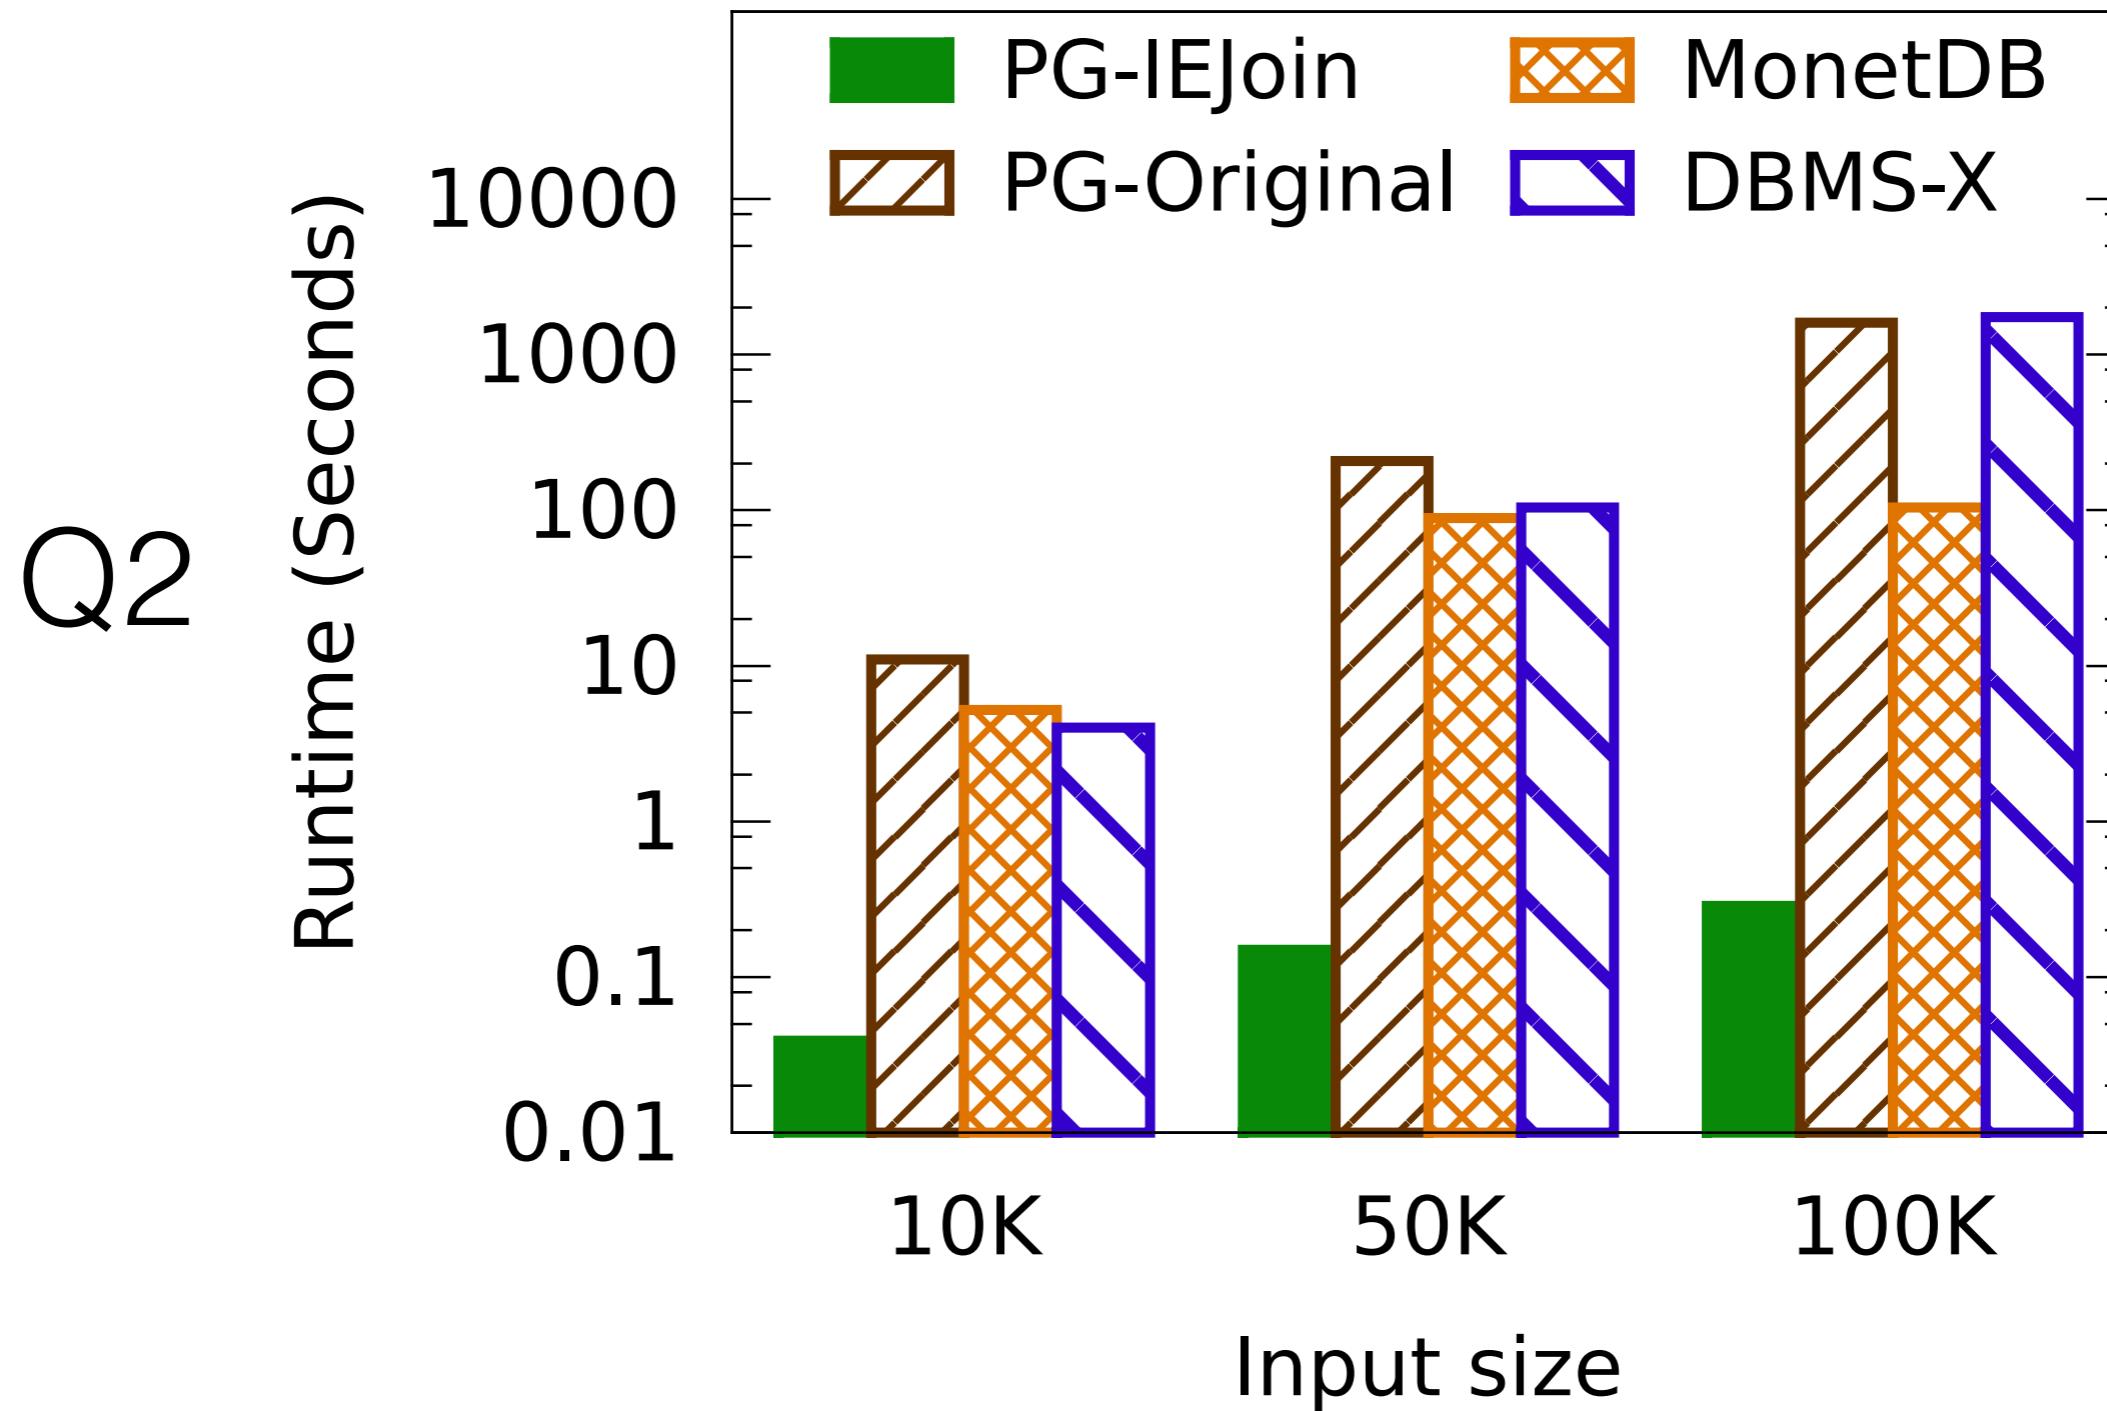
DBMS-X



Single-Node Performance

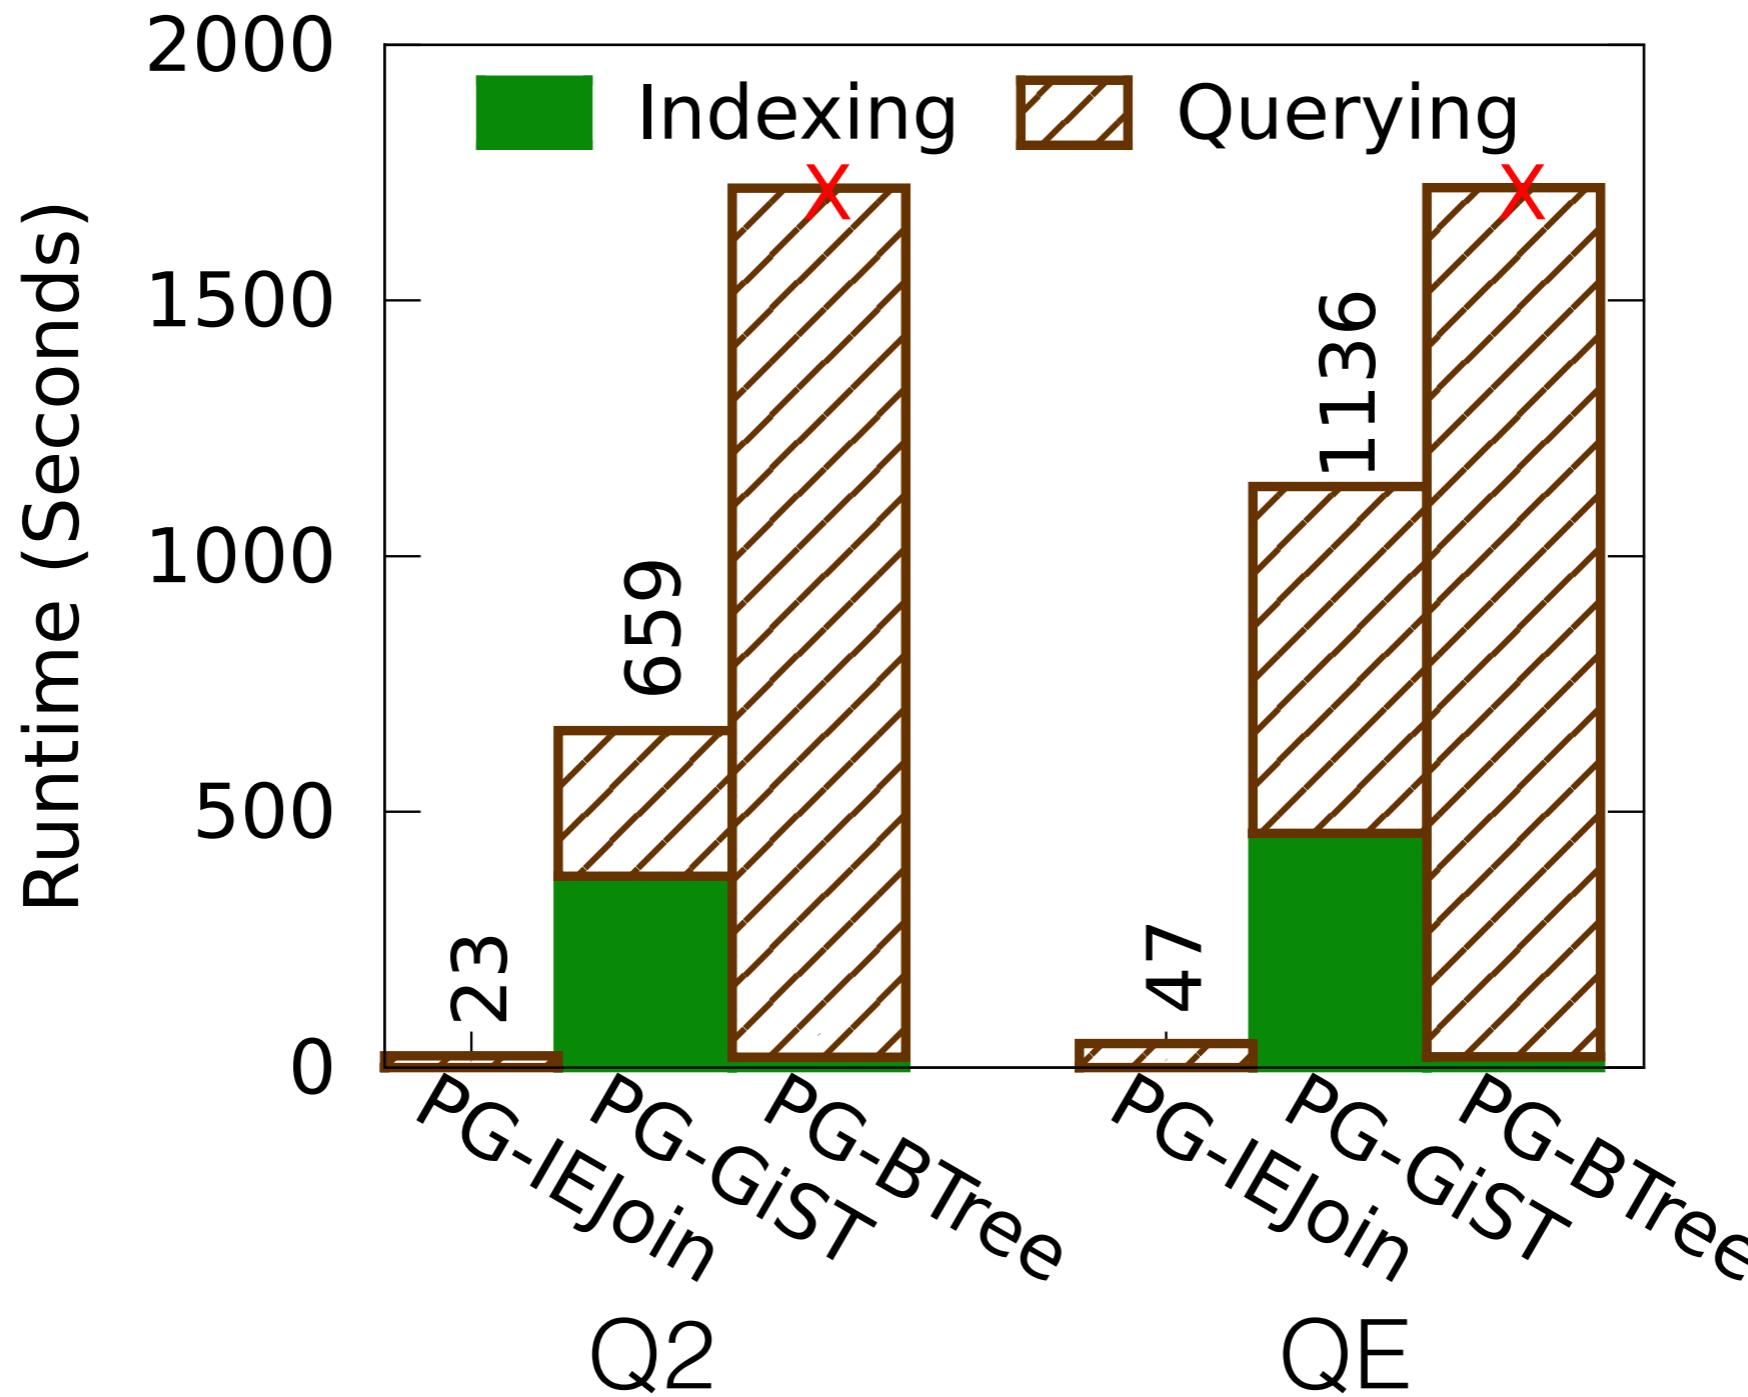


Single-Node Performance

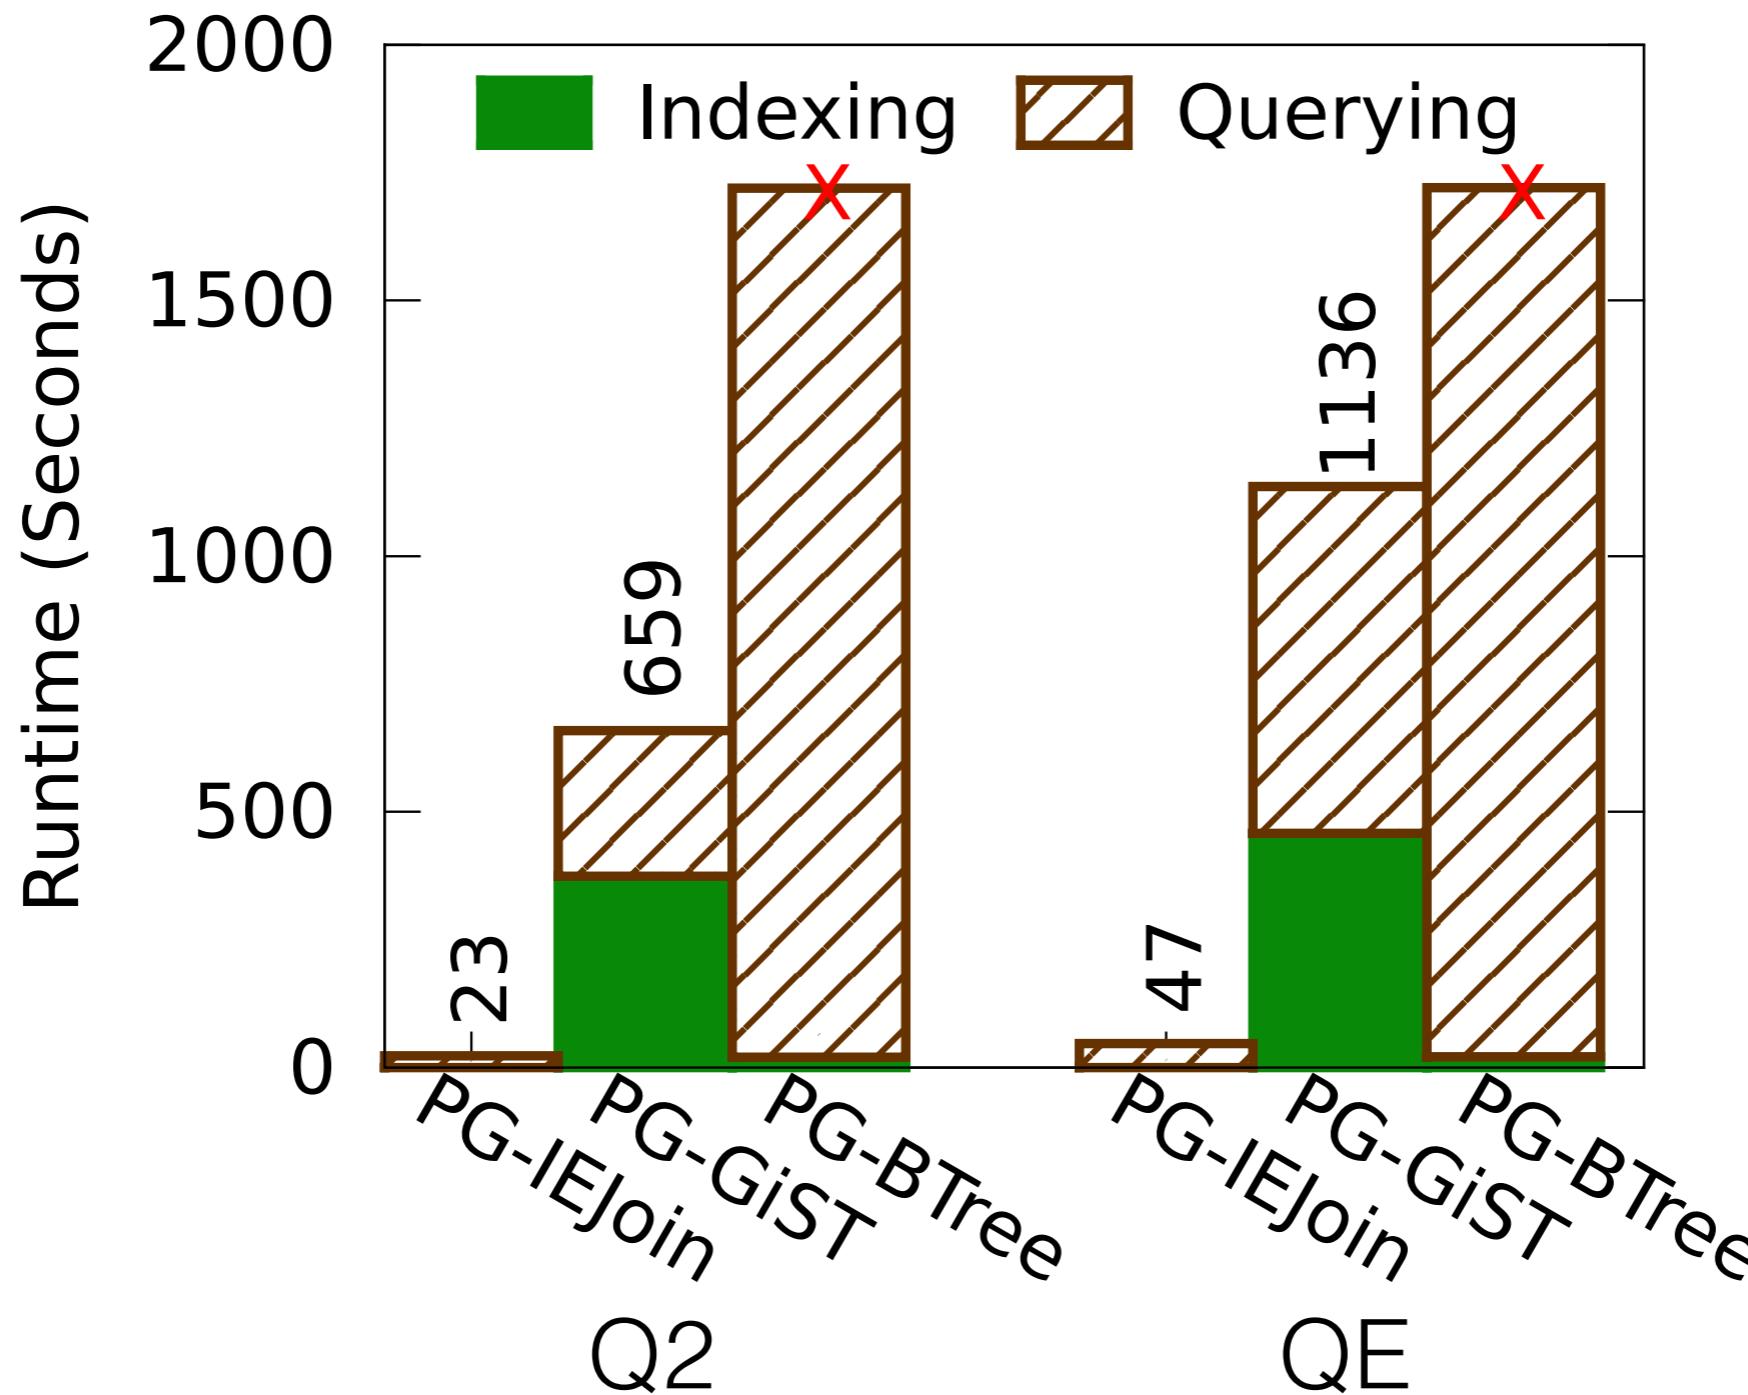


up to more than 3 orders of magnitude!

IEJoin vs Indexing

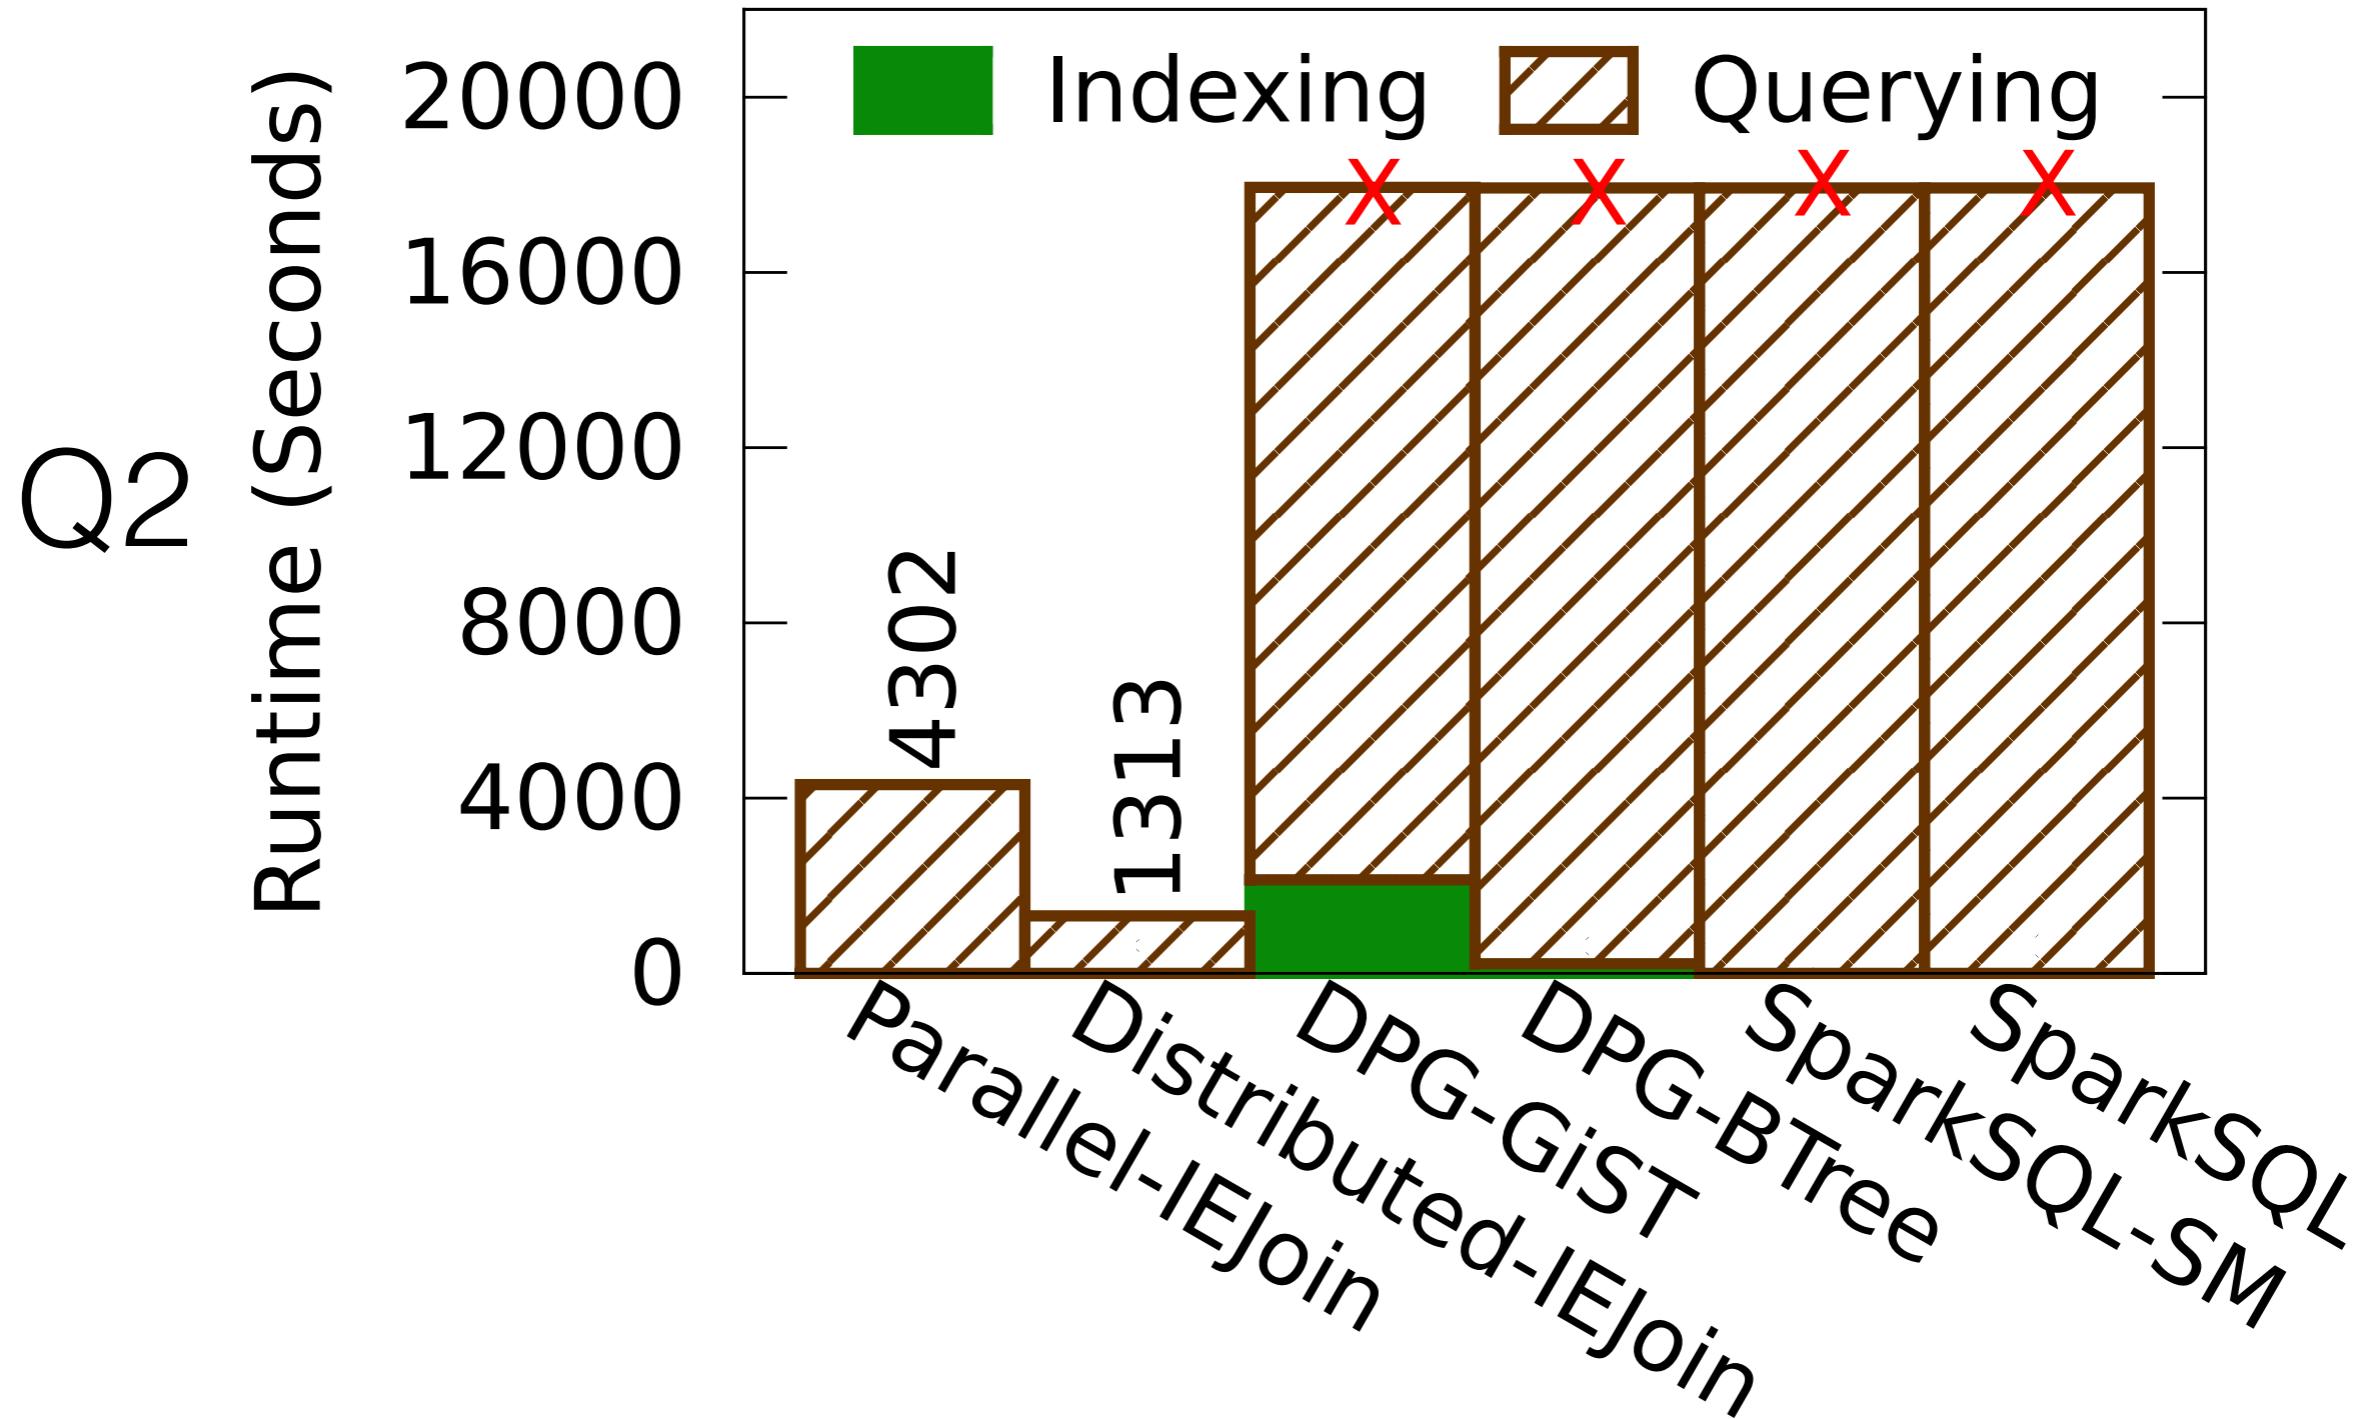


IEJoin vs Indexing

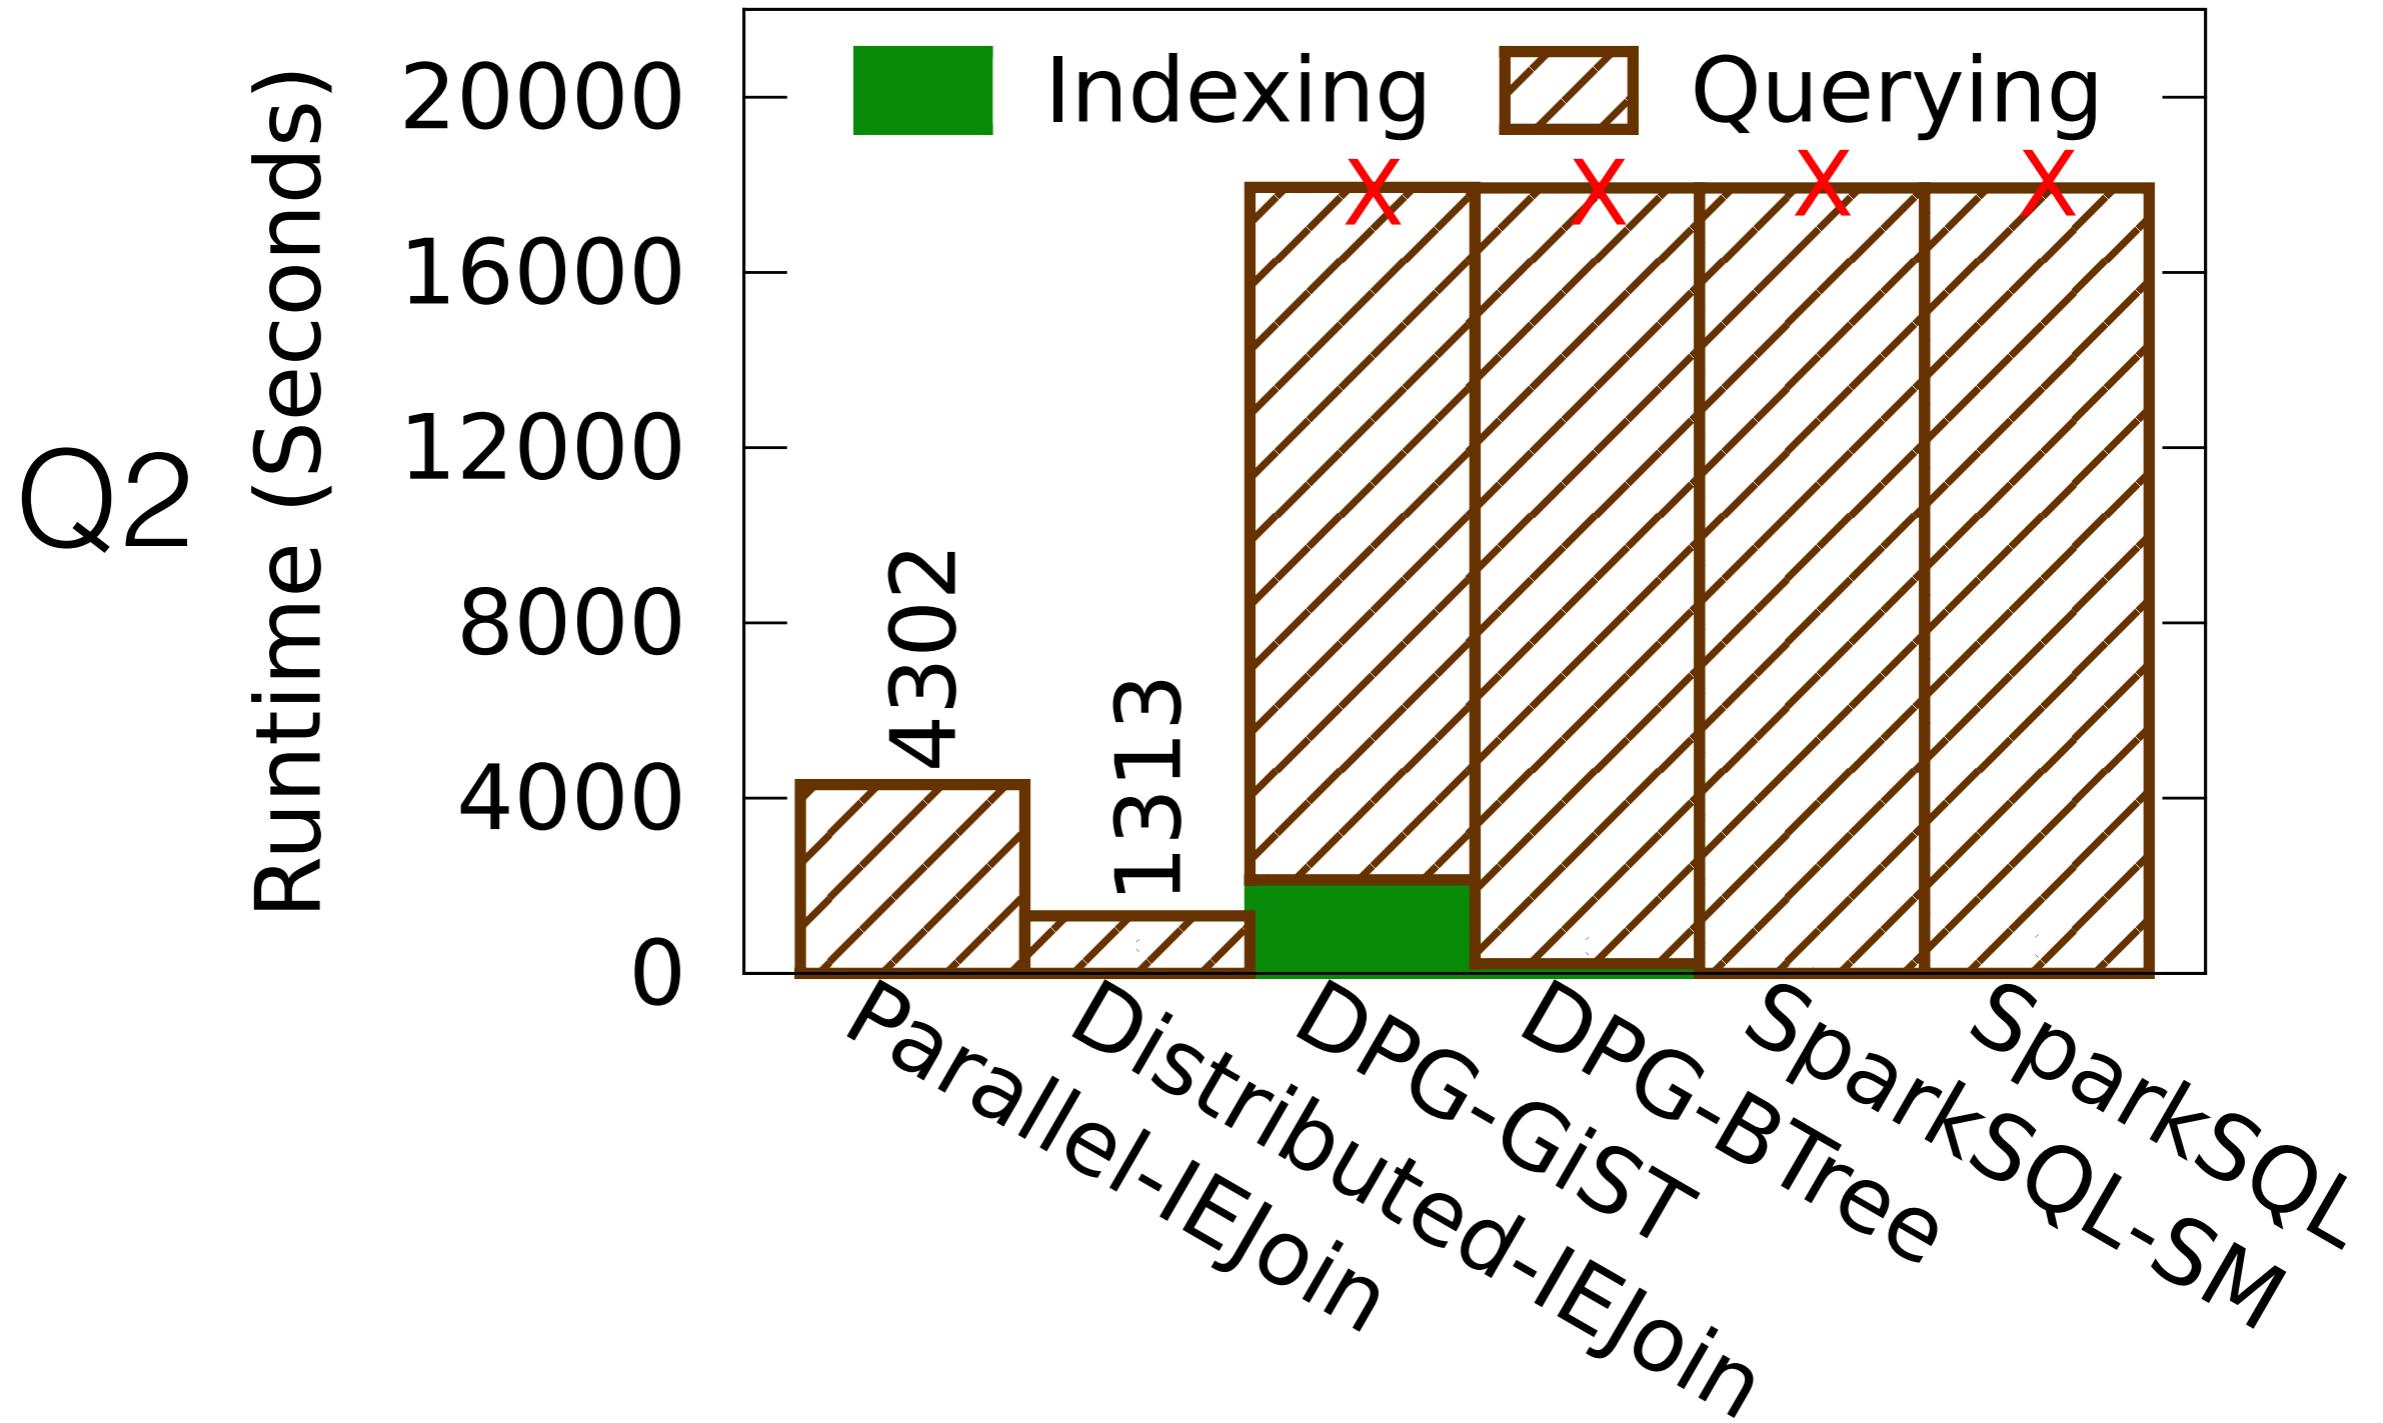


more than 1 order of magnitude!

Multi-Node Performance

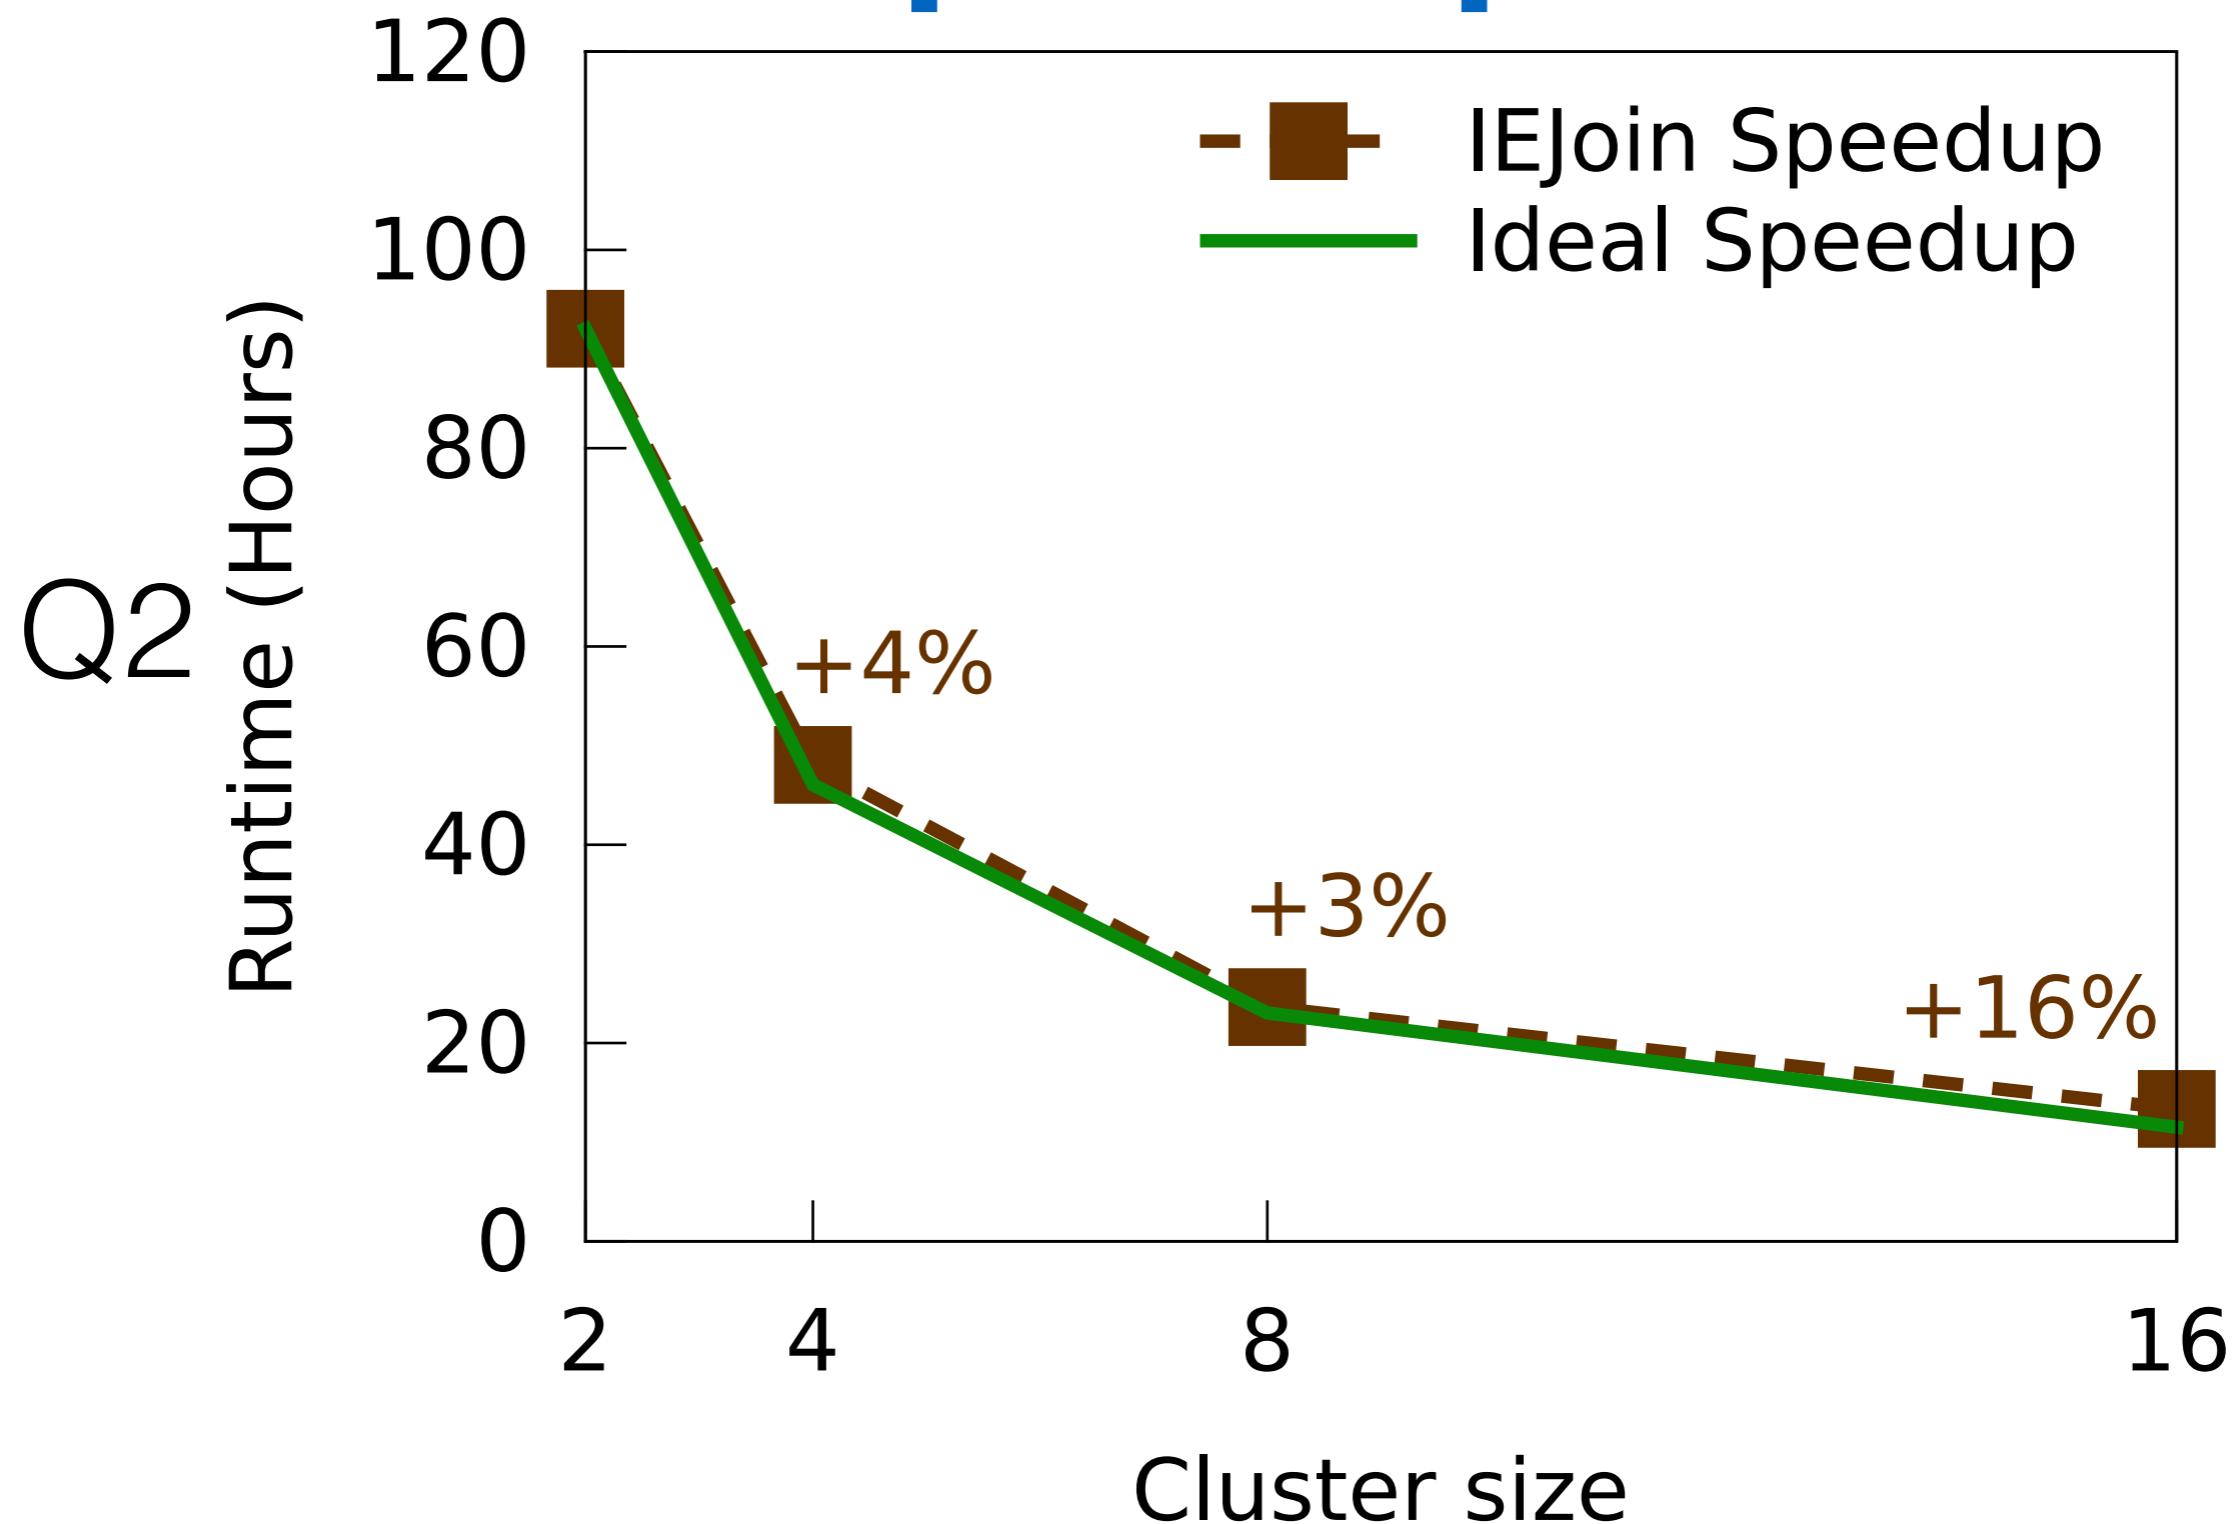


Multi-Node Performance

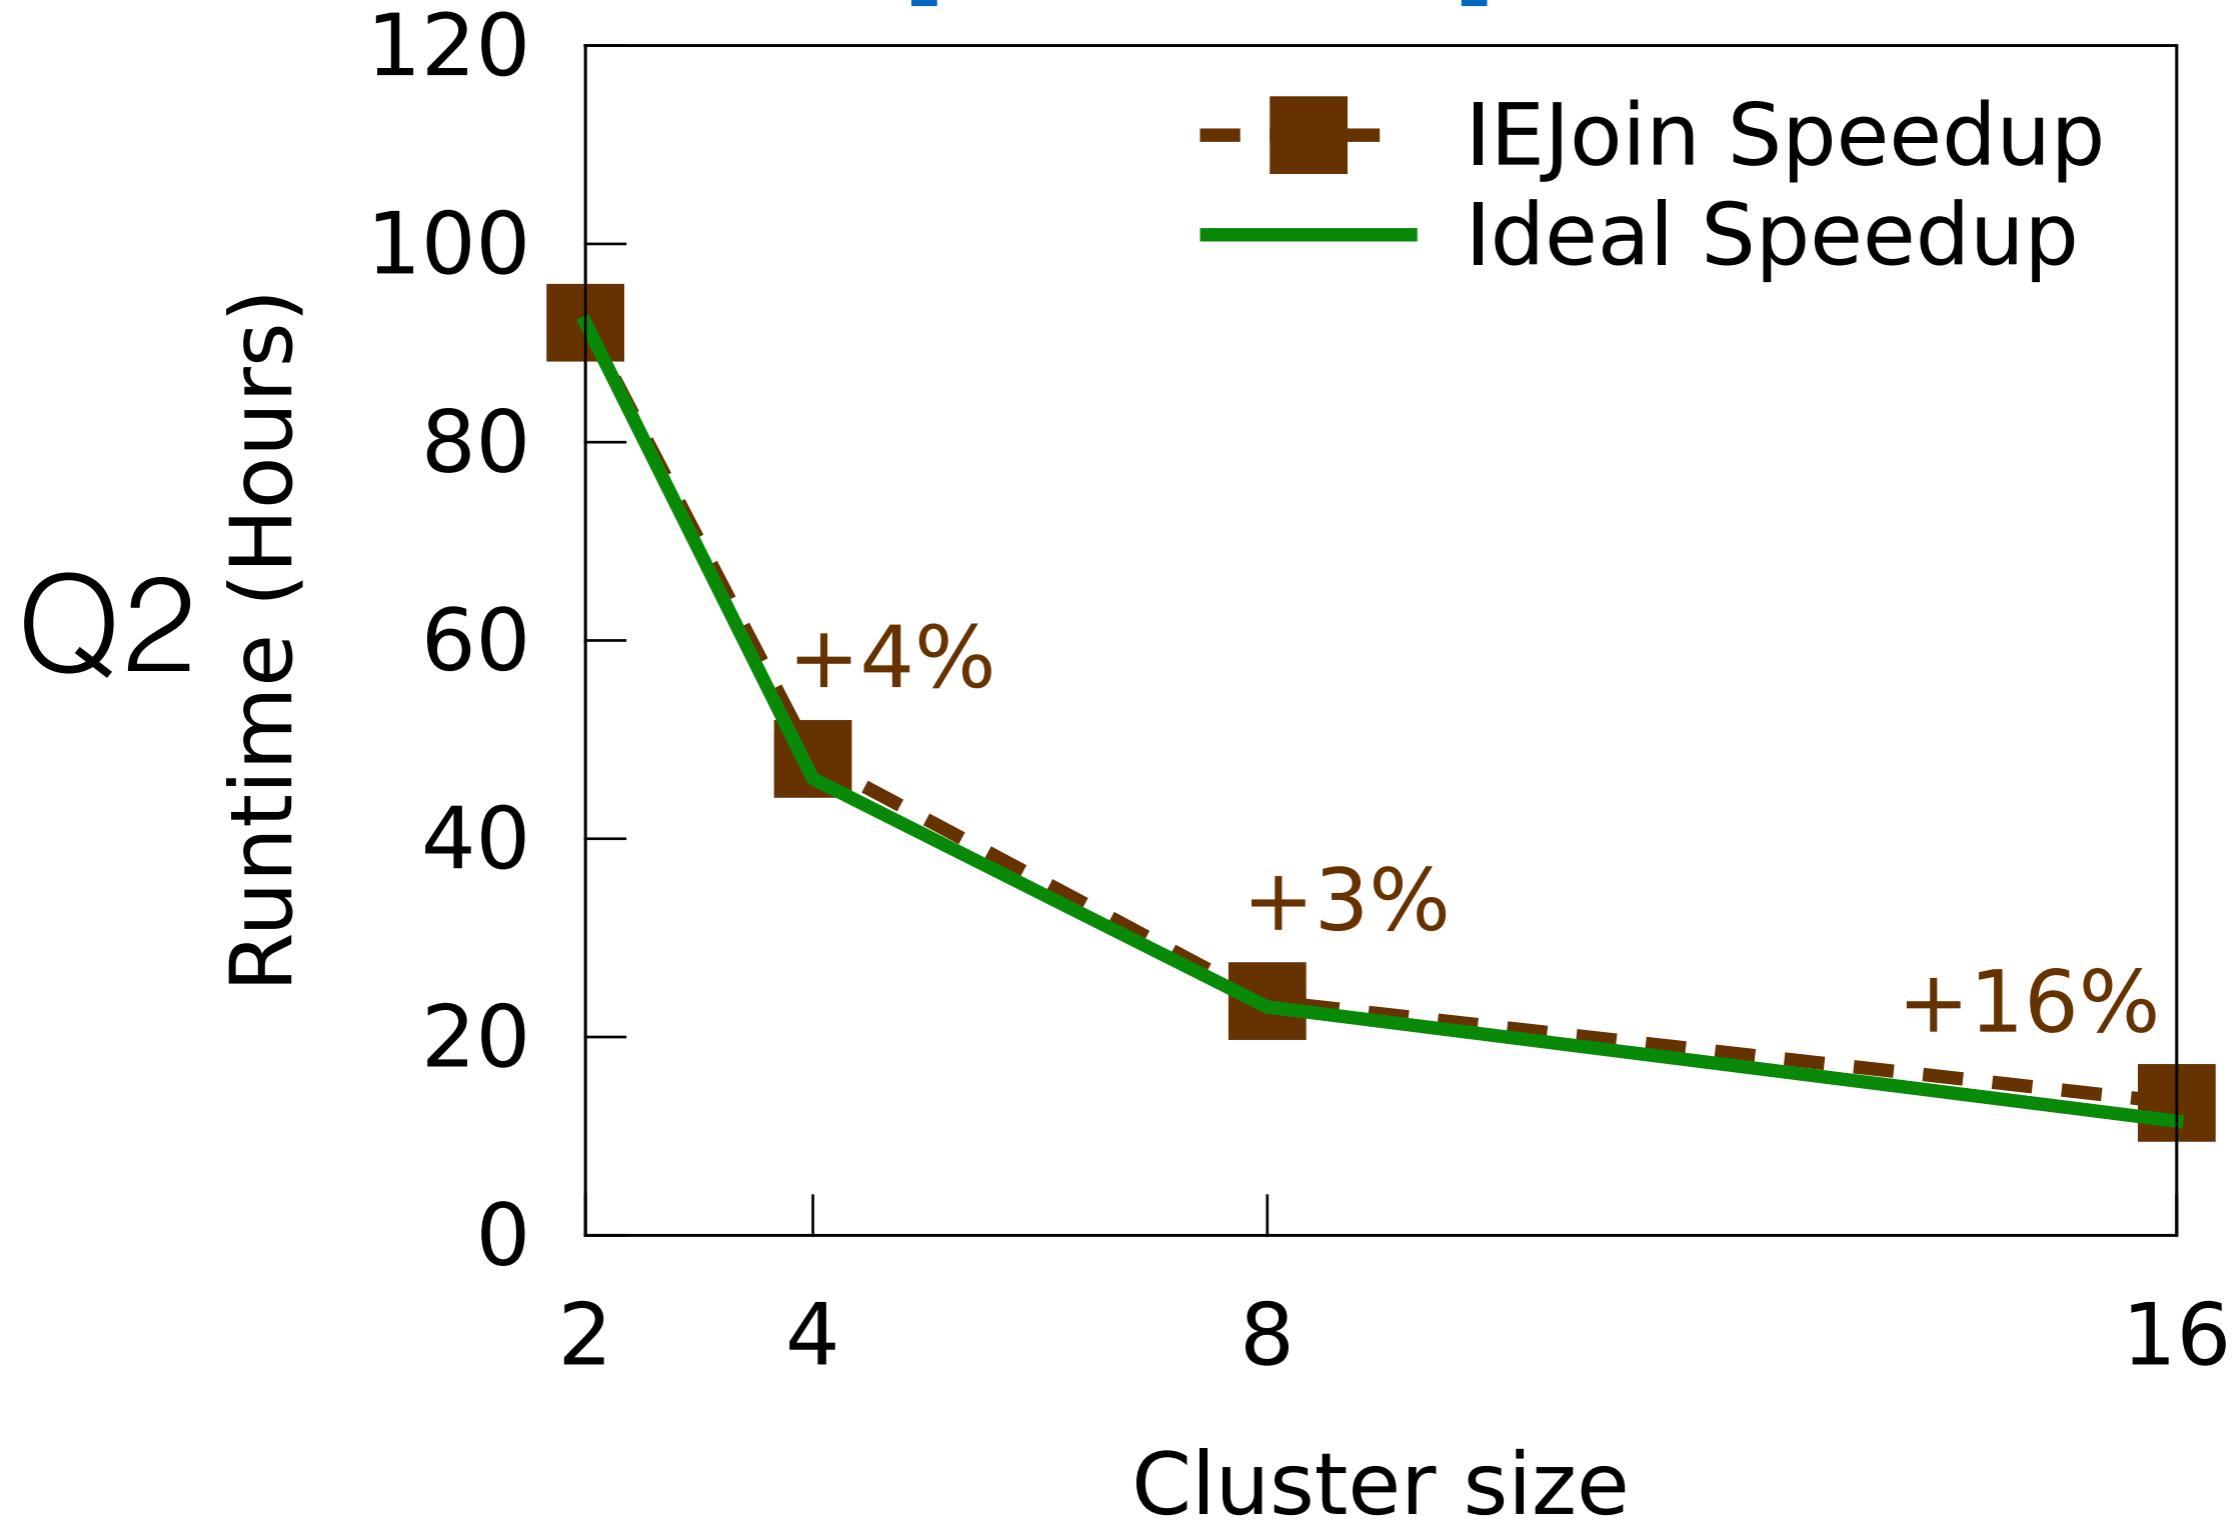


more than 2 orders of magnitude!

Speedup



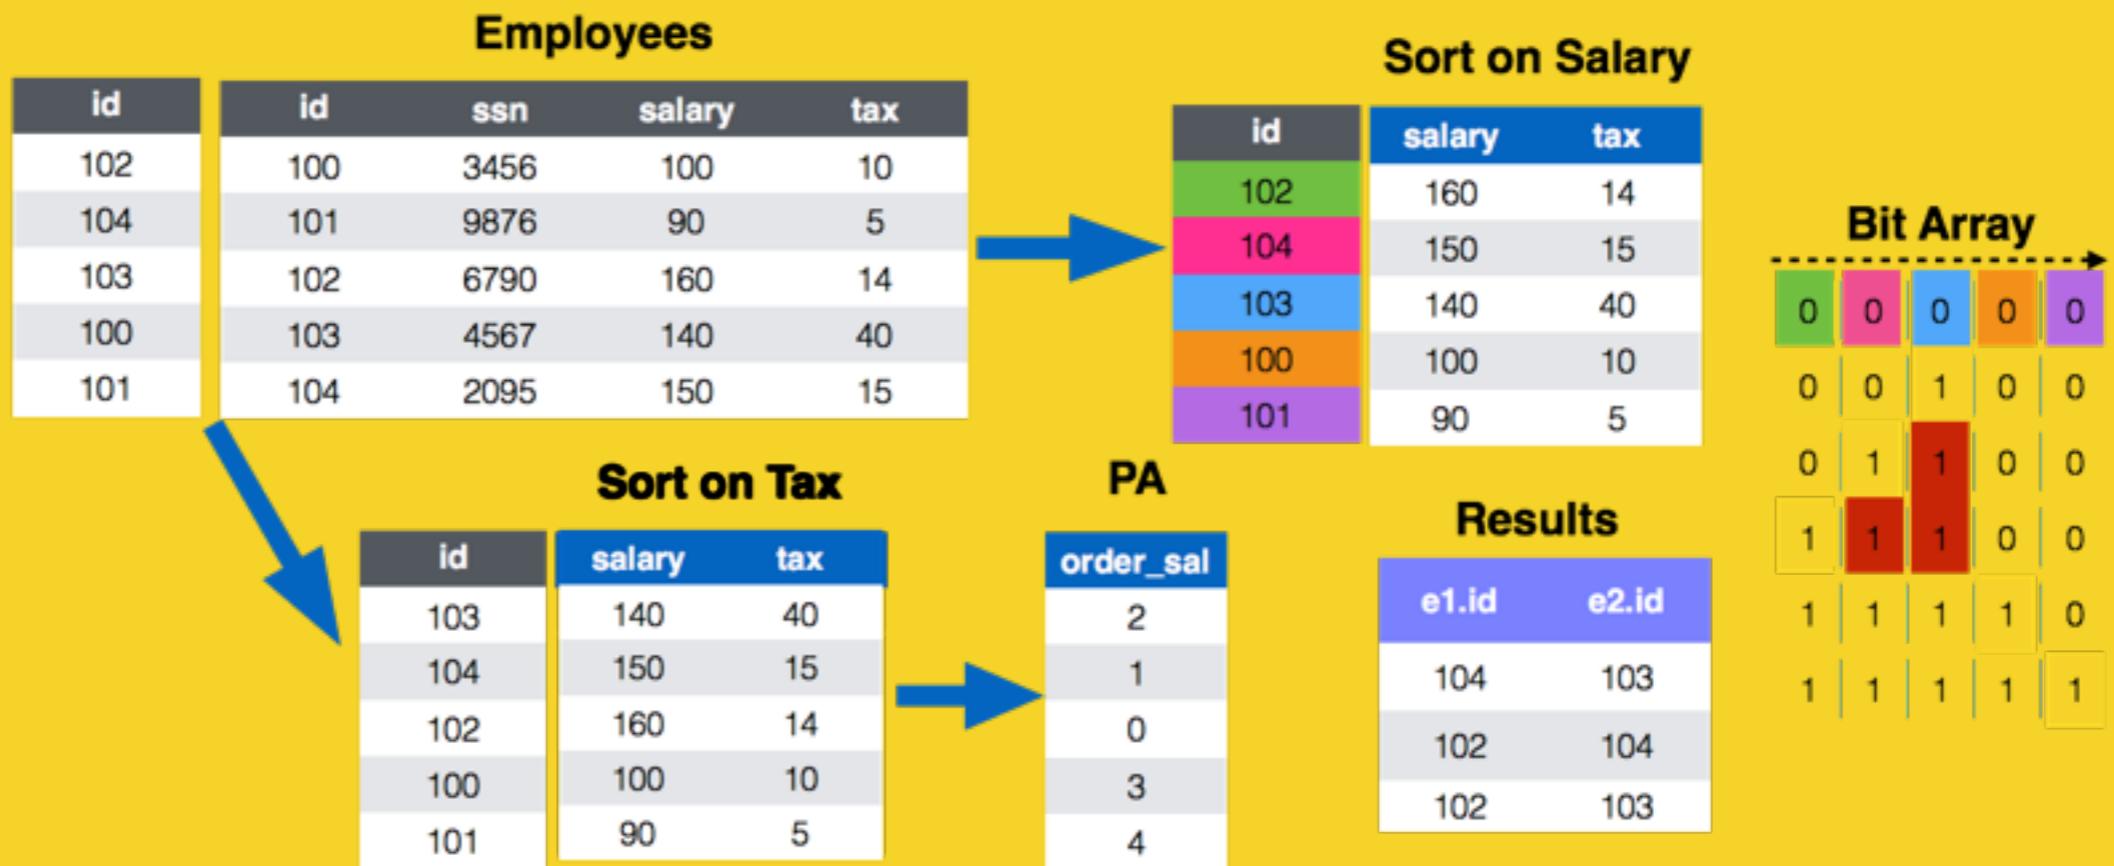
Speedup



close to **ideal** speedup!

Summary

Simple



Fast

